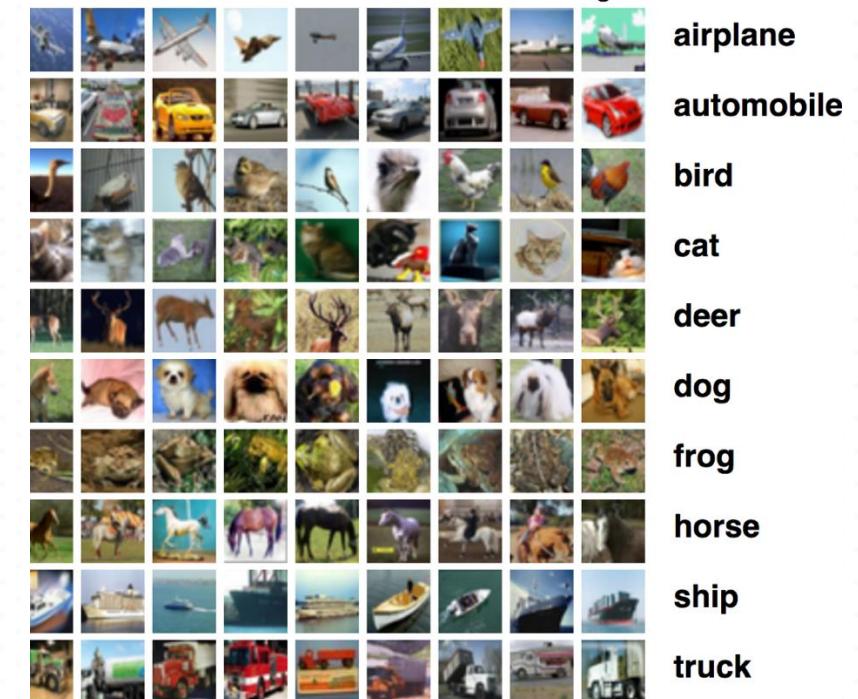


# ECE 4252/8803: Fundamentals of Machine Learning (FunML)

## Fall 2024

### Lecture 5: Classifiers (Introduction to Neural Networks)



# Reminders

- Reminders
  - HW 1 is due this Friday, 6-Sept
  - HW 2 was posted and due next Friday, 13-Sept
- Lecture 2 Scribe notes were posted
  - Will be posted by weekends going forward
  - Please see the posted notes to get an idea of expectation
  - Q/A at the end added by GTA, Shiva
  - Focus must be on:
    - Visuals
    - Numerical examples
    - Concepts

# Classifier Comparison

Methods	Assumptions on Feat. Dist.	Feat. Normalization	Cost Function	Regularization	Linear Classifier	Prob. View of Prediction	Generative/Discriminative	Parametric/Non-parametric	Overfitting
Logistic Regression	No	Required	BCE (convex)	Additional term	Linear	Yes	Discriminative	Parametric	Not often
K Nearest Neighbors	No	Required	N/A	N/A	Non-linear	N/A	Discriminative	Non-parametric	when k is too small
Decision Trees	No	Not Required	N/A	N/A	Non-linear	N/A	Discriminative	Non-parametric	with large depth
Support Vector Machines	No	Required	Hinge (convex)	C (control robustness)	Linear/ Non-linear(kernel)	N/A	Discriminative	Parametric	Not often
Naïve Bayes	Conditional independent	Not Required	N/A	N/A	Non-linear /Linear (Gaussian)	Yes	Generative	Parametric	Not often
Artificial Neural Networks	No	Required	Non-convex	Additional term	Non-linear	Yes	Discriminative/ Generative	Parametric	with many layers

# Overview

In this Lecture..

Nearest Neighbor

Naïve Bayes

Logistic Regression

Decision Trees

Support Vector Machines

Artificial Neural Networks

- Overview
- Activation Function
- Perceptron Network
- Multi-layer ANN
- Feedforward and Backward Error Propagation
- Learning Algorithm
- Image Classification using ANNs

# Artificial Neural Networks

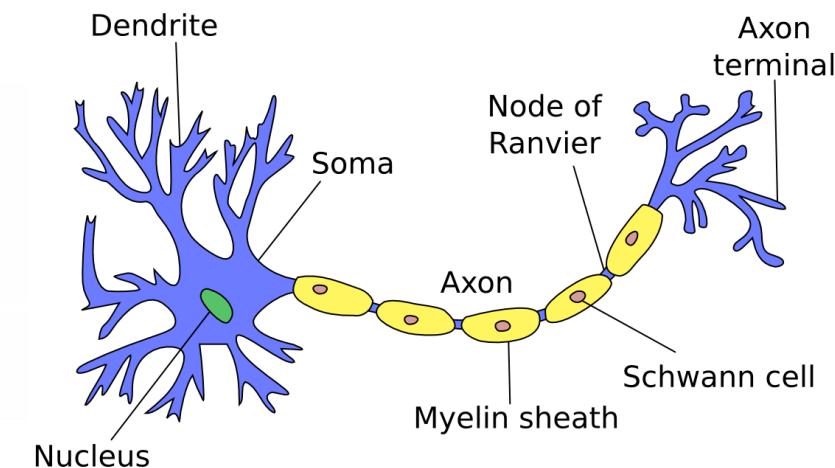
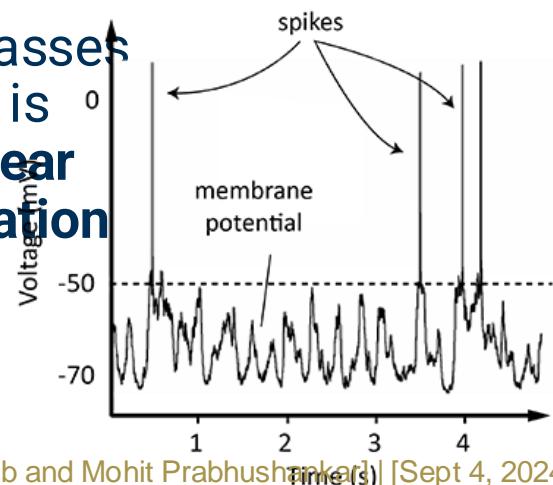
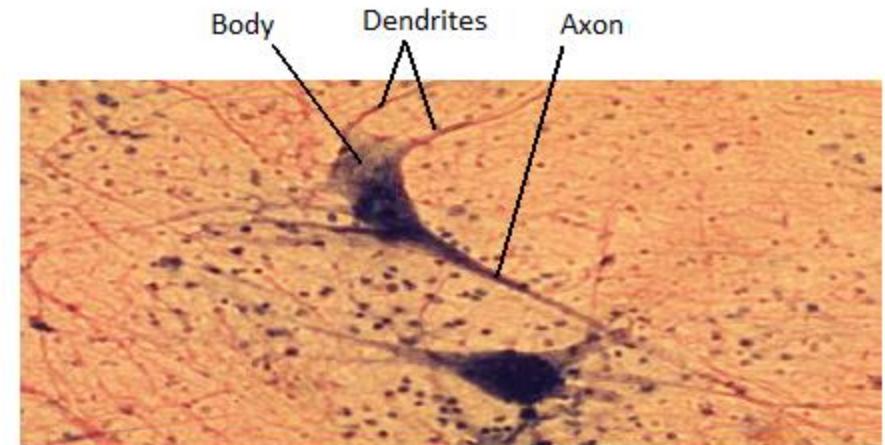
## Overview

- Overview
- Activation Function
- Perceptron Network
- Multi-layer ANN
- Feedforward and Backward Error Propagation
- Learning Algorithm
- Image Classification using ANNs

# Artificial Neural Networks

## Biological Neurons

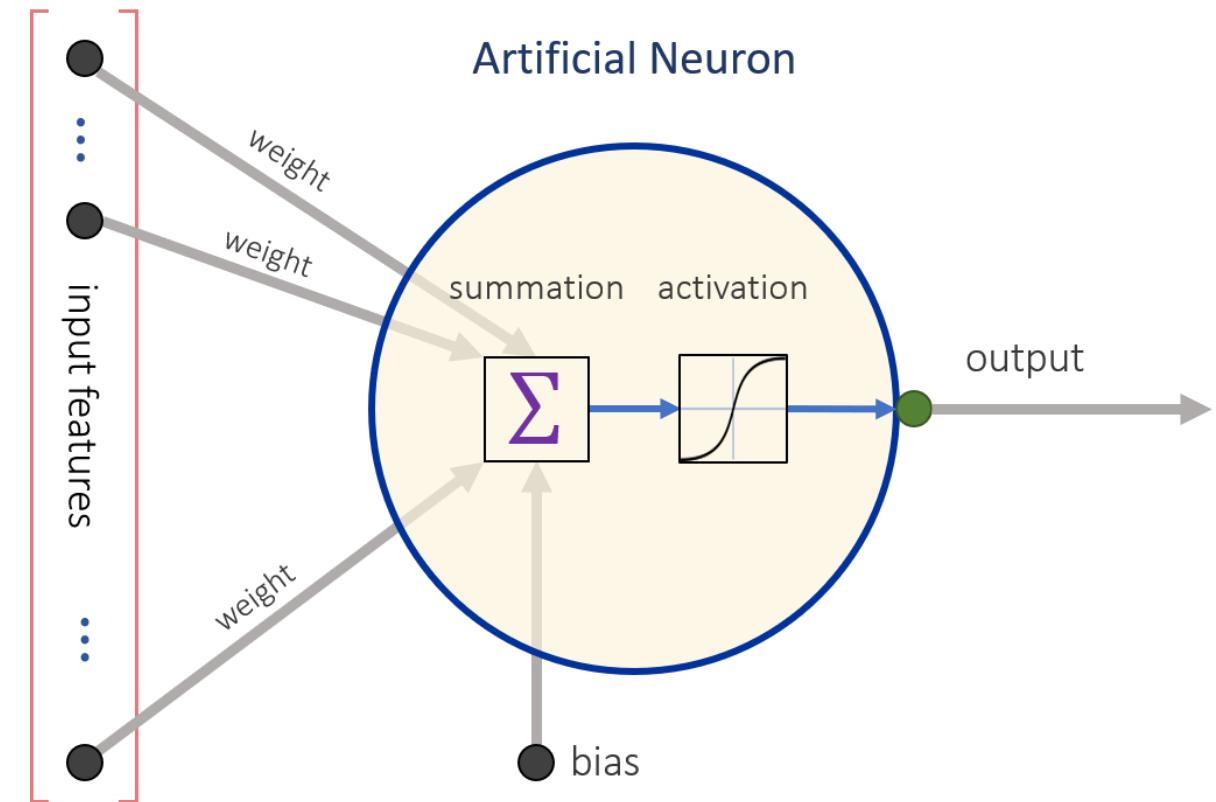
- The brain has approximately 100 billion neurons
- Neurons communicate through electro-chemical signals. Neurons are connected through junctions called synapses.
- Each neuron receives thousands of connections with other neurons, constantly receiving incoming signals.
- If the resulting sum of the signals surpasses a certain voltage threshold, a response is sent through the axon. This is a **non linear** relation and motivates the use of **activation**



# Artificial Neural Networks

## Artificial Neuron

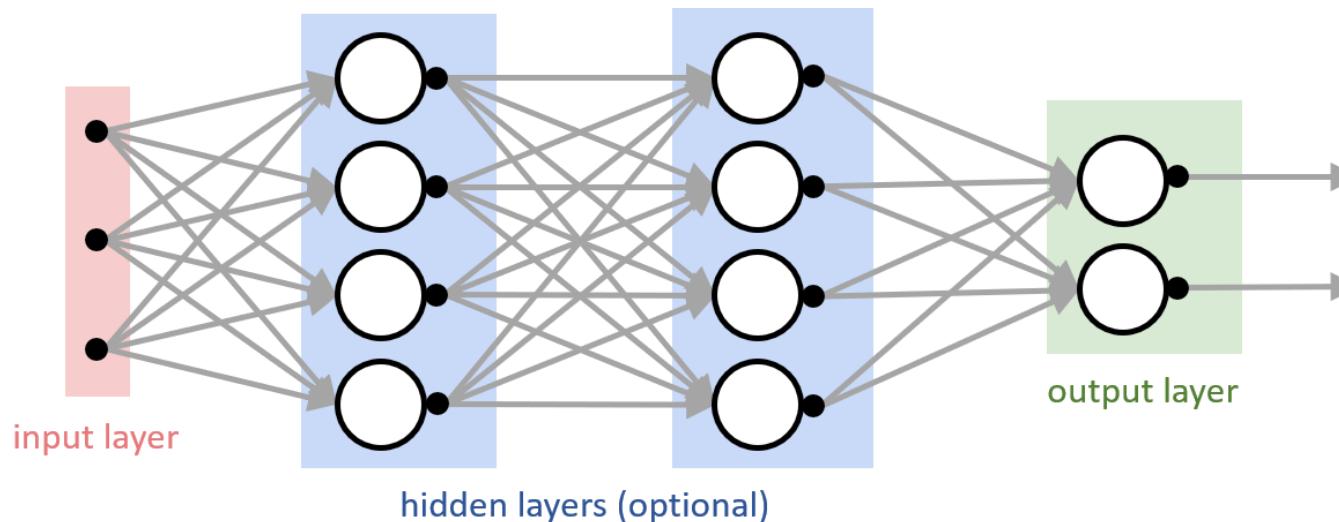
- A computational unit consisting of:
  - A single output
  - Multiple inputs
  - Input weights
  - A bias input
  - An activation function



# Artificial Neural Networks

## Artificial Neurons within a Network

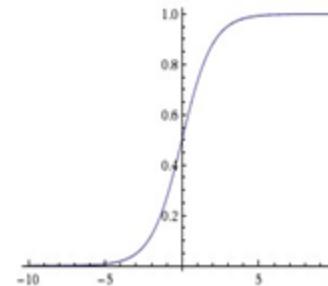
- Typically, artificial neurons in ANNs are connected in layers:
  - An **input layer** (Layer 0)
  - An **output layer** (Layer  $K$ )
  - Zero or more **hidden** (middle) layers (Layers 1 ...  $K - 1$ )



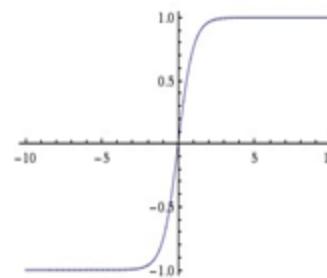
# Artificial Neural Networks

## Common Activation Functions

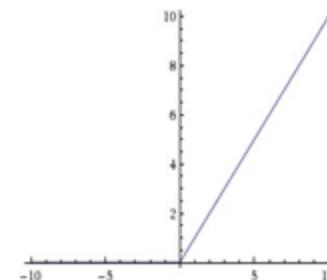
**Sigmoid**  $\sigma(x) = \frac{1}{(1+e^{-x})}$



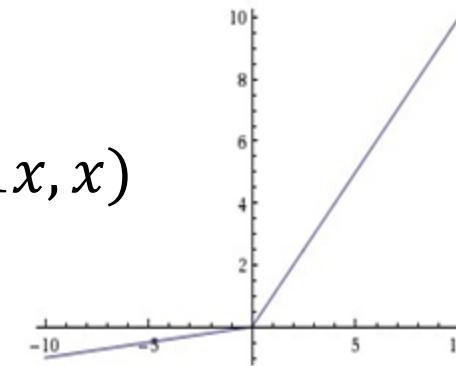
**Tanh**  $\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$



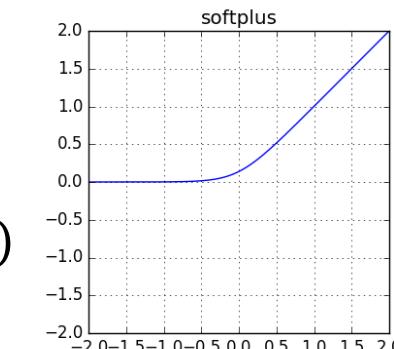
**ReLU**  $\max(0, x)$



**Leaky ReLU**  $\max(0.1x, x)$



**Maxout**  $\max(w_1^T x + b_1, w_2^T x + b_2)$

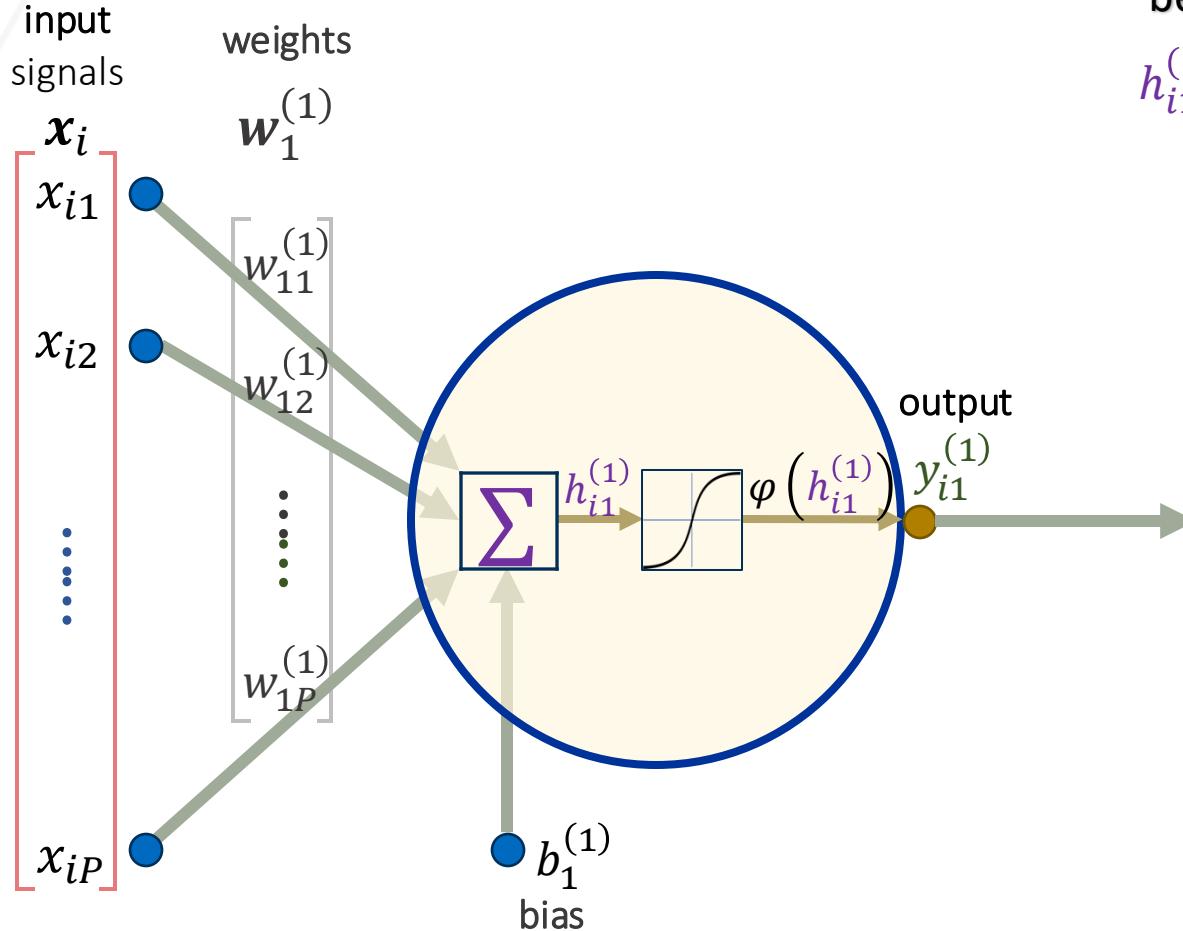


**SoftPlus**  $f(x) = \ln(1 + e^x)$

# Artificial Neural Networks

## The Perceptron

- Single-layer Perceptron (SLP)



weighted input to the neuron  
before activation  $\varphi$

$$h_{i1}^{(1)} = (\mathbf{w}_j^{(1)})^T \mathbf{x}_i + b_1^{(1)}$$

output of the neuron  
after activation  $\varphi$

$$\varphi(h_{i1}^{(1)})$$

The simplest form of the perceptron uses linear activation  $\varphi(h_{i1}^{(1)}) = h_{i1}^{(1)}$ , the outputs are binary (i.e.  $y_{i1}^{(1)} \in \{1, -1\}$ ):

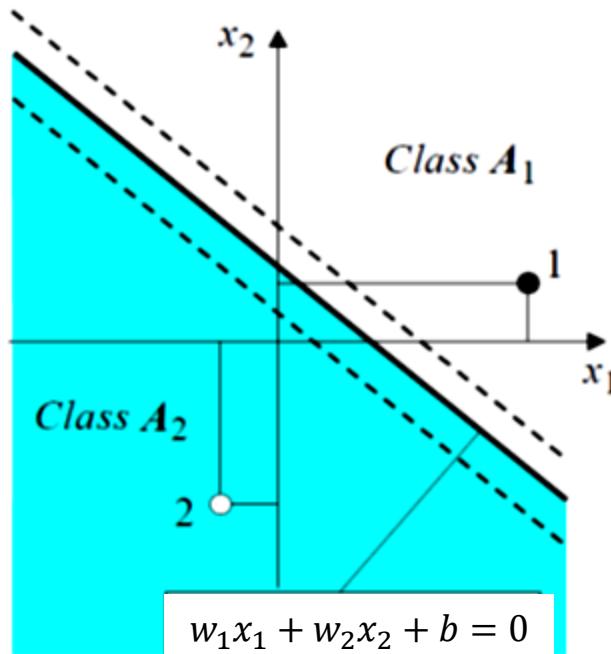
$$y_{i1}^{(1)} = \begin{cases} +1, & \text{if } (\mathbf{w}_j^{(1)})^T \mathbf{x}_i + b_1^{(1)} \geq 0 \\ -1, & \text{if } (\mathbf{w}_j^{(1)})^T \mathbf{x}_i + b_1^{(1)} < 0 \end{cases}$$

The above model applies to linearly separable data

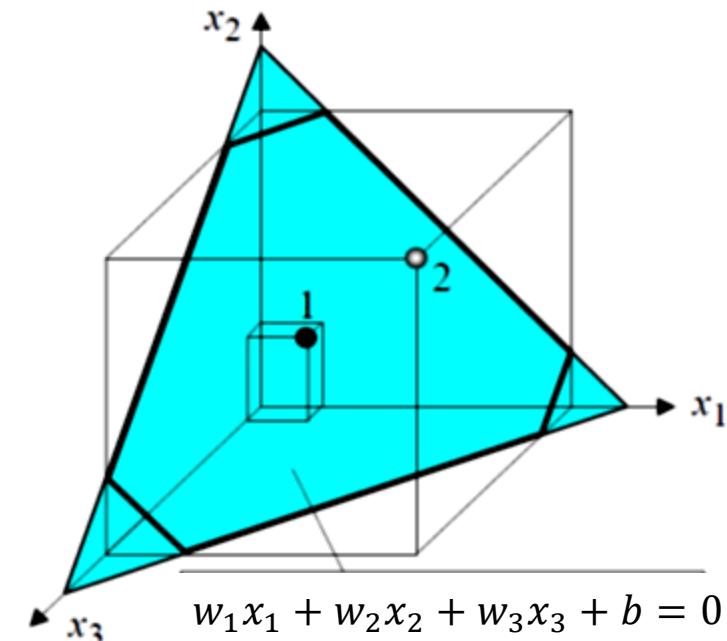
# Artificial Neural Networks

## The SLP Perceptron Model

- Linear separability in the perceptron



(a) Two-input perceptron.



(b) Three-input perceptron.

# Artificial Neural Networks

## The SLP Perceptron Learning Algorithm

### 1. Weight Initialization

Initial weights  $\mathbf{w} = [w_1, w_2, \dots, w_P]$  are set to random values

### 2. Neuron Activation

Calculate perceptron outputs with activation function such as sigmoid:

$$y = \sigma(\mathbf{w}^T \mathbf{x} + b)$$

### 3. Weight Update

Weights are updated based on the learning rule:

$$w_i^{t+1} = w_i^t + \alpha x_i^t e^t$$

Where

- $e$ : difference between the calculated output and desired output
- $\alpha$ : learning rate which is a positive constant less than 1

### 4. Iteration

Input next training sample, and the algorithm keeps iterating between steps 2 and 3 until convergence

# Artificial Neural Networks

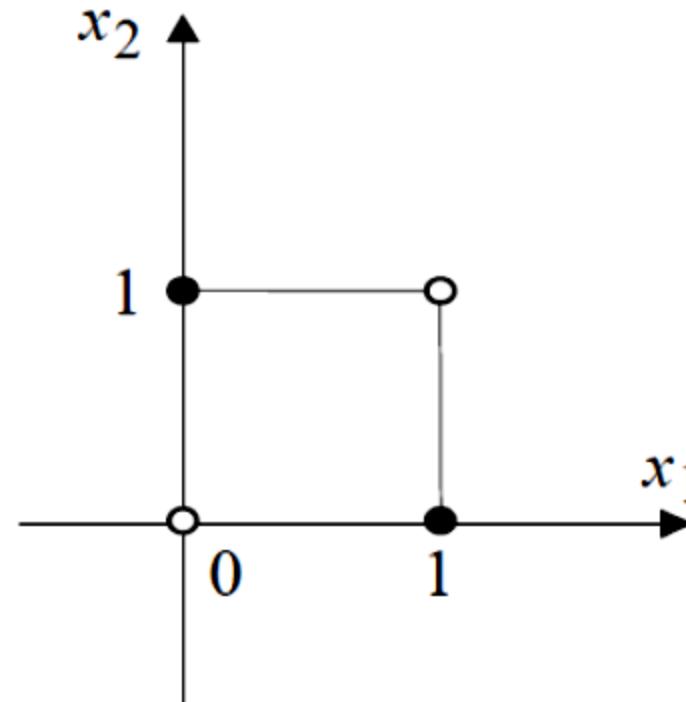
## The XOR Problem

The XOR, or “exclusive or”, problem is a classic problem in ANN research. It is the problem of using a neural network to predict the outputs of XOR logic gates given two binary inputs.

- is the output=0
- is the output=1

These are not linearly separable using one neuron (a line hyperplane).

We need a decision plane! → hidden layer (MLP)



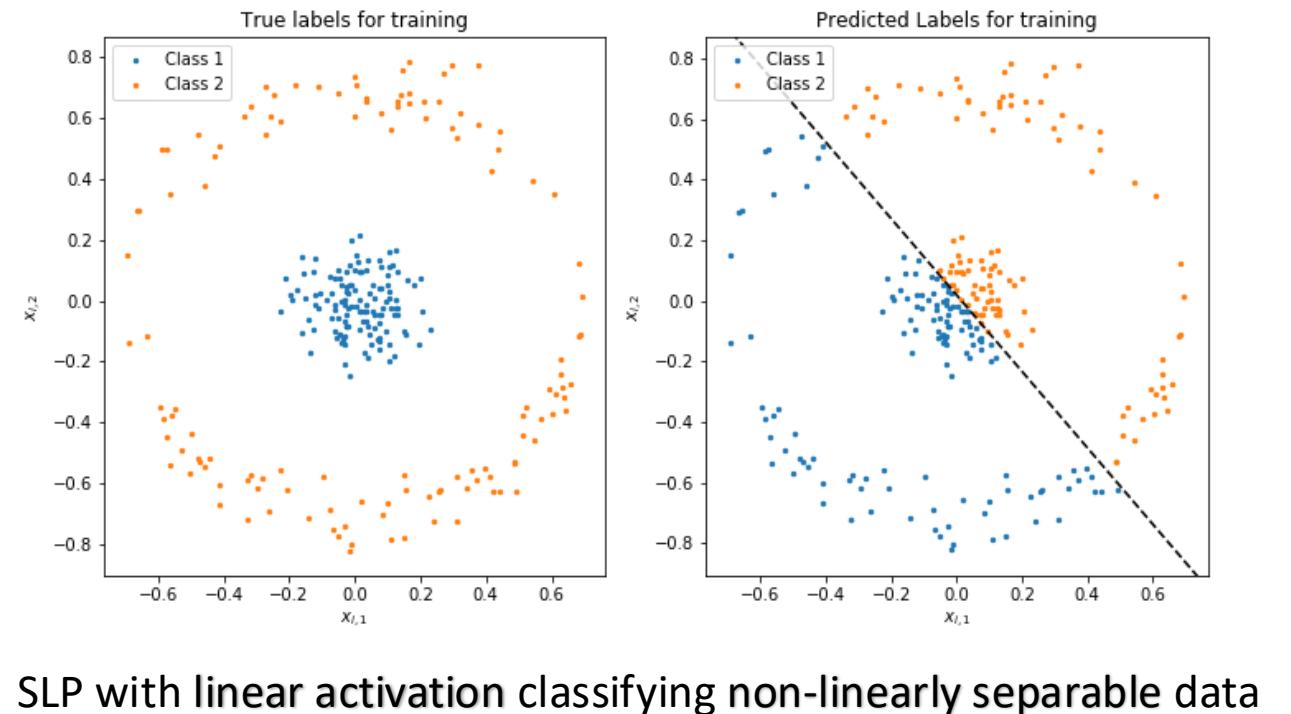
(c) *Exclusive-OR*  
 $(x_1 \oplus x_2)$

X1	X2	Y
0	0	0
0	1	1
1	0	1
1	1	0

# Artificial Neural Networks

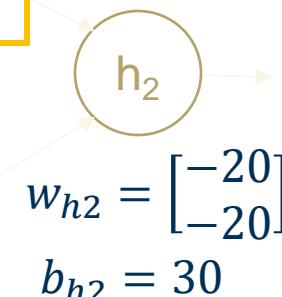
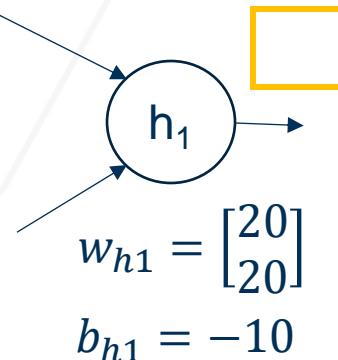
## Hidden Layers

- Single neurons with linear activation fail to classify non-linearly separable dataset
- Non-linearly separable data requires
  - Non-linear Activation
  - Multi-layer networks of neurons (Multi-layer perceptron)

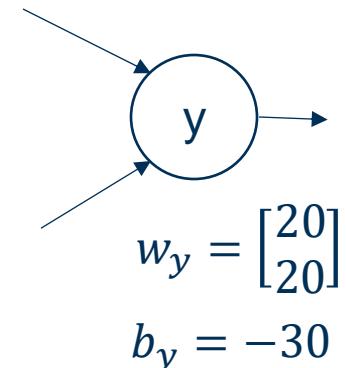


# Artificial Neural Networks

## Solving XOR with Hidden Layers



### Output Neuron

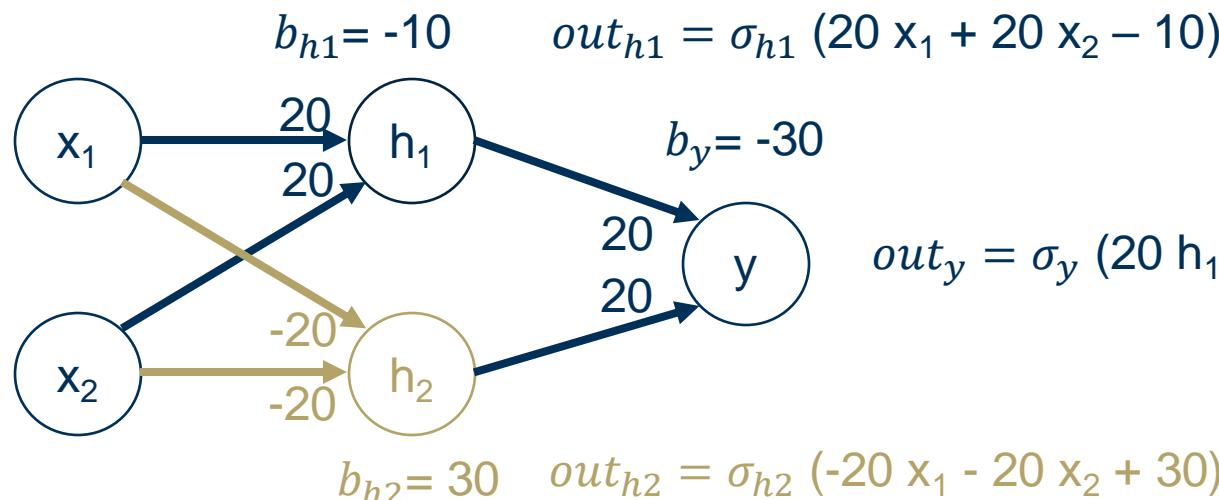


$$out_{h1} = \sigma_{h1}(w_{h1}[x_1 \ x_2] + b_{h1})$$

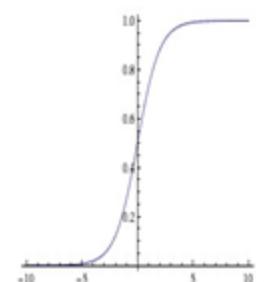
$$out_{h2} = \sigma_{h2}(w_{h2}[x_1 \ x_2] + b_{h2})$$

$$out_y = \sigma_y(w_y[out_{h1} \ out_{h2}] + b_y)$$

x1	x2	Y
0	0	0
0	1	1
1	0	1
1	1	0

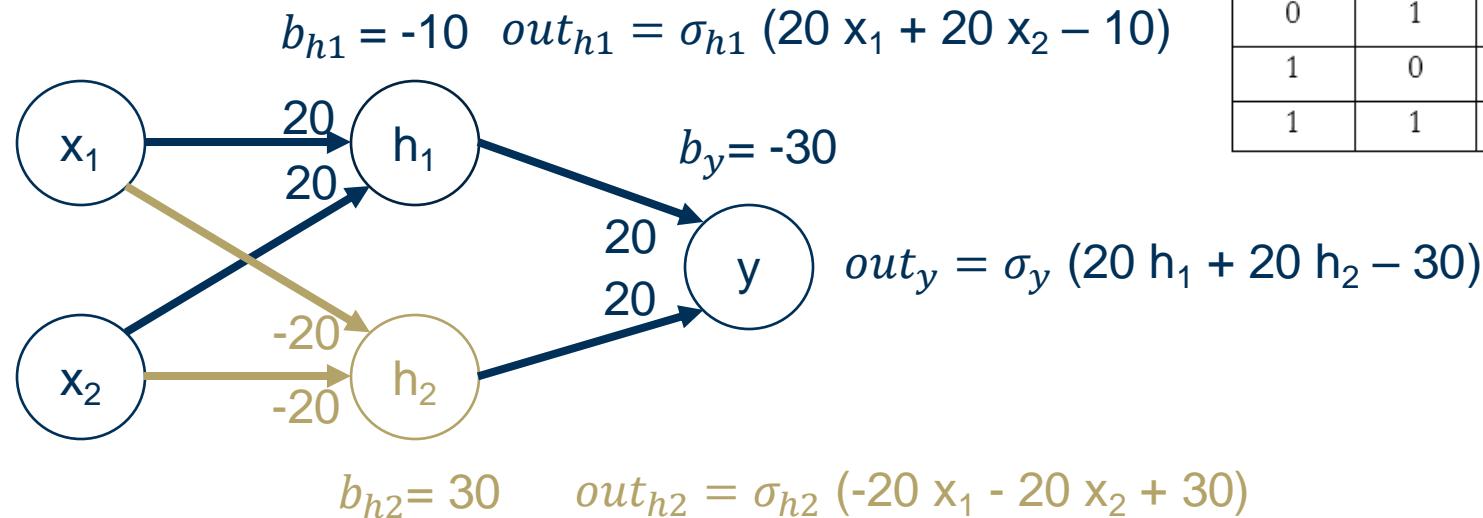
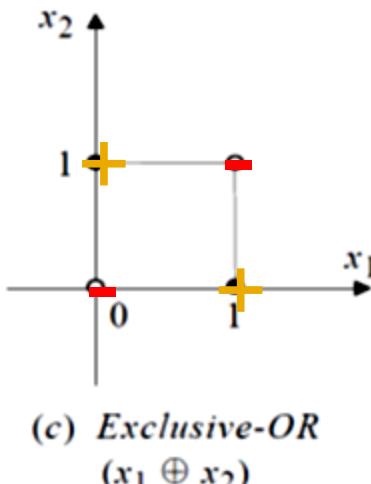


**Sigmoid**  $\sigma(x) = \frac{1}{(1+e^{-x})}$

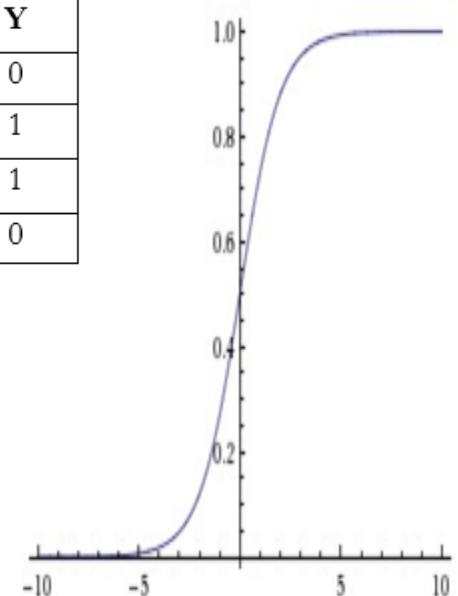


# Artificial Neural Networks

## Solving XOR with Hidden Layers



X1	X2	Y
0	0	0
0	1	1
1	0	1
1	1	0



$$\text{Input } (0,0) \rightarrow \sigma(20 \times 0 + 20 \times 0 - 10) \approx 0$$

$$\sigma(-20 \times 0 - 20 \times 0 + 30) \approx 1 \quad \sigma(20 \times 0 + 20 \times 1 - 30) \approx 0$$

$$\text{Input } (1,1) \rightarrow \sigma(20 \times 1 + 20 \times 1 - 10) \approx 1$$

$$\sigma(-20 \times 1 - 20 \times 1 + 30) \approx 0 \quad \sigma(20 \times 1 + 20 \times 0 - 30) \approx 0$$

$$\text{Input } (0,1) \rightarrow \sigma(20 \times 0 + 20 \times 1 - 10) \approx 1$$

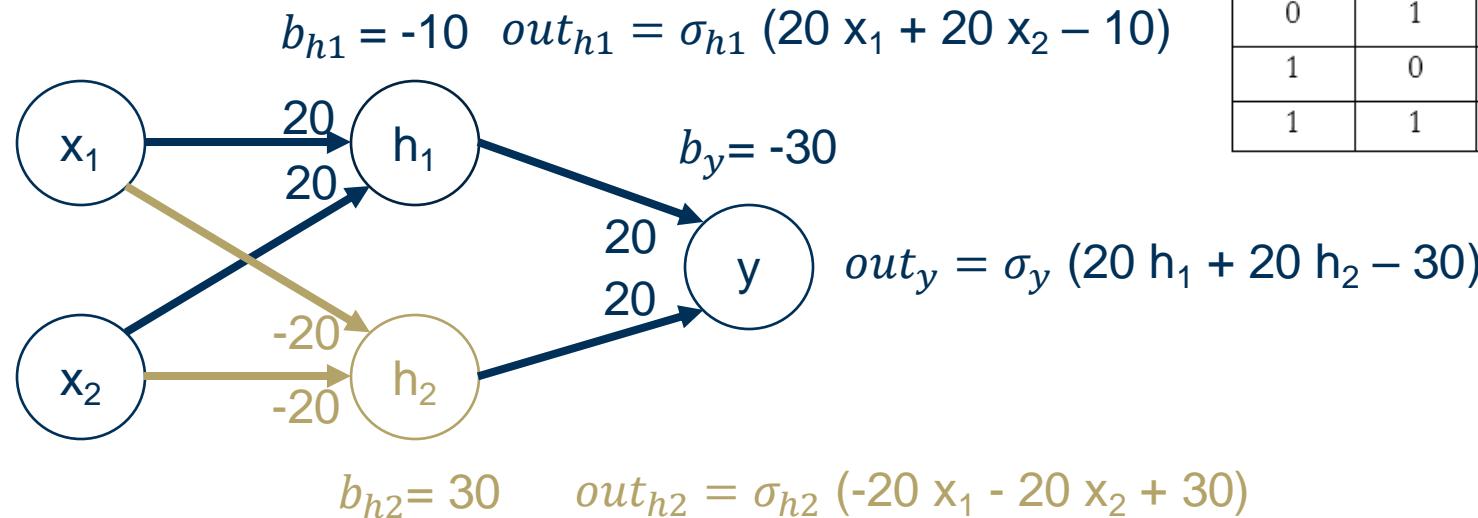
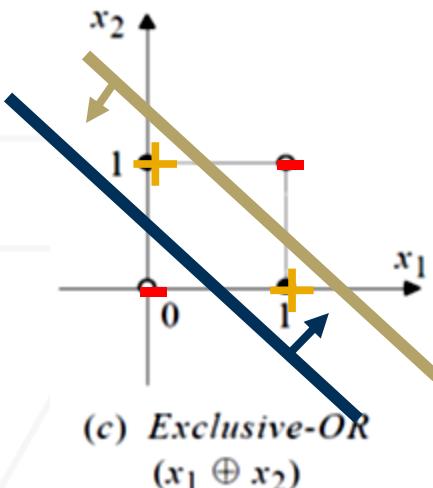
$$\sigma(-20 \times 0 - 20 \times 1 + 30) \approx 1 \quad \sigma(20 \times 1 + 20 \times 1 - 30) \approx 1$$

$$\text{Input } (1,0) \rightarrow \sigma(20 \times 1 + 20 \times 0 - 10) \approx 1$$

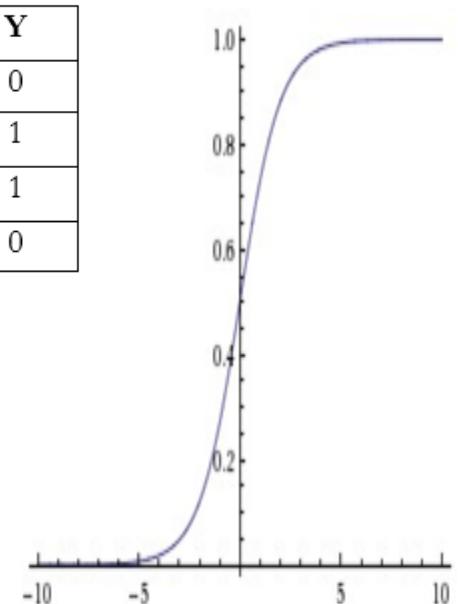
$$\sigma(-20 \times 1 - 20 \times 0 + 30) \approx 1 \quad \sigma(20 \times 1 + 20 \times 1 - 30) \approx 1$$

# Artificial Neural Networks

## Solving XOR with Hidden Layers



X1	X2	Y
0	0	0
0	1	1
1	0	1
1	1	0



Input (0,0)  $\rightarrow \sigma (20 \times 0 + 20 \times 0 - 10) \approx 0$

$\sigma (-20 \times 0 - 20 \times 0 + 30) \approx 1$  1  $\sigma (20 \times 0 + 20 \times 1 1 - 30) \approx 0$

Input (1,1)  $\rightarrow \sigma (20 \times 1 + 20 \times 1 - 10) \approx 1$

$\sigma (-20 \times 1 - 20 \times 1 + 30) \approx 0$  0.  $\sigma (20 \times 1 + 20 \times 0 0 - 30) \approx 0$

Input (0,1)  $\rightarrow \sigma (20 \times 0 + 20 \times 1 - 10) \approx 1$

$\sigma (-20 \times 0 - 20 \times 1 + 30) \approx 1$  1  $\sigma (20 \times 1 + 20 \times 1 1 - 30) \approx 1$

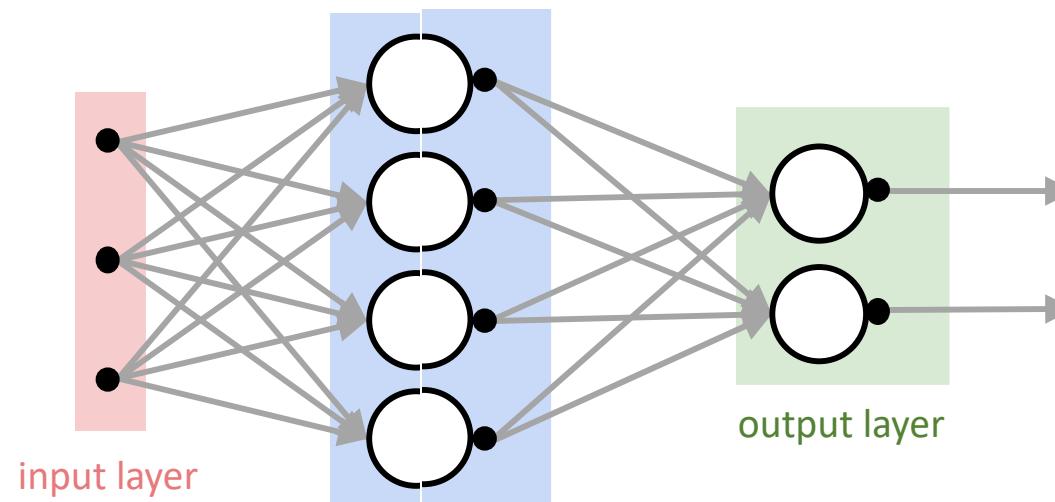
Input (1,0)  $\rightarrow \sigma (20 \times 1 + 20 \times 0 - 10) \approx 1$

$\sigma (-20 \times 1 - 20 \times 0 + 30) \approx 1$  1  $\sigma (20 \times 1 + 20 \times 1 1 - 30) \approx 1$

# Artificial Neural Networks

## Hidden Layers

- A NN with one hidden layer can represent:
  - any bounded continuous function (to some arbitrary  $\epsilon$ ) [Universal Approximation Theorem, Cybenko 1989]
  - any Boolean Function, but it requires  $2^k$  hidden units for  $1K$  inputs



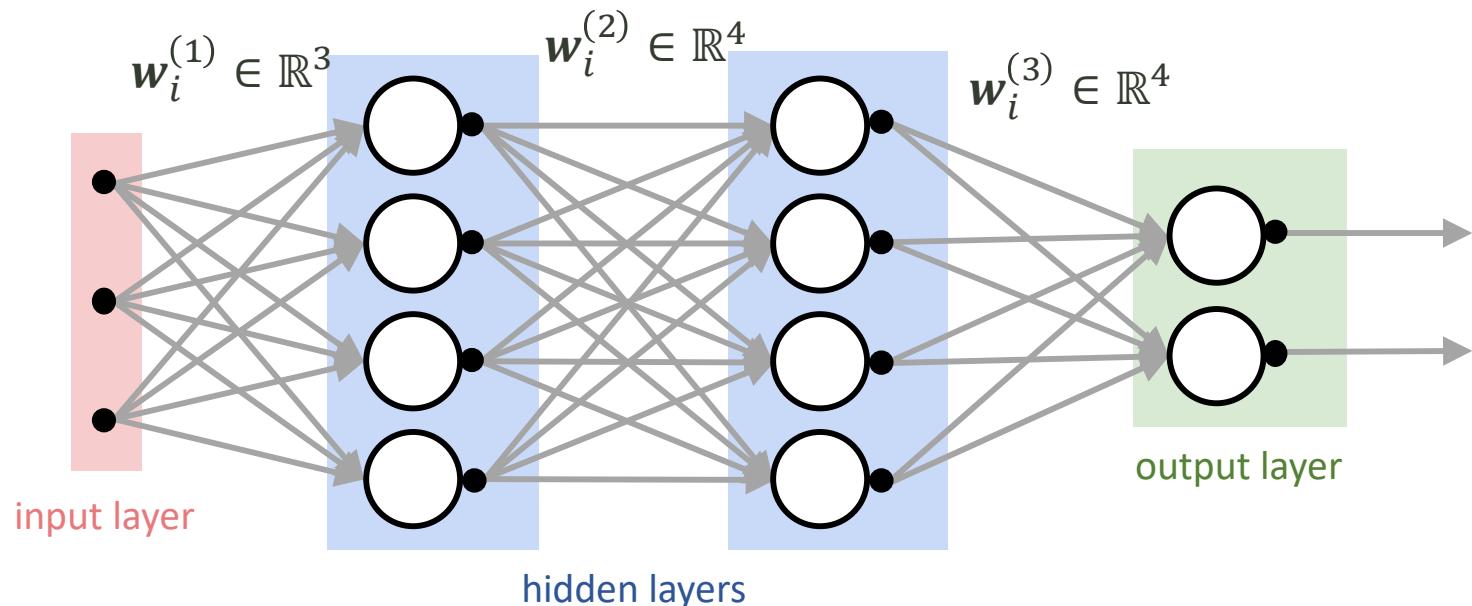
# Artificial Neural Networks

## Multiple Hidden Layers

**Layer-wise organization:** MLP consists of *fully-connected* layers in which neurons between two adjacent layers are fully pairwise connected, while neurons within a single layer share no connections.

**Sizing MLP:** the number of neurons or the number of parameters (more commonly)

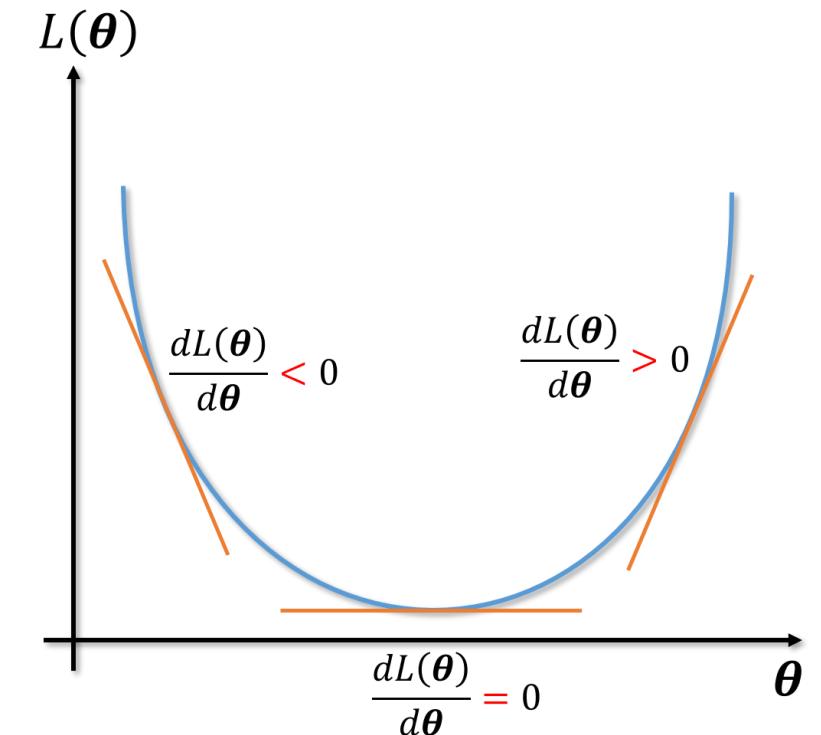
- The network has  $4 + 4 + 2 = 10$  neurons,  $[3 \times 4] + [4 \times 4] + [4 \times 2] = 12 + 16 + 8 = 36$  weights and  $4 + 4 + 2 = 10$  biases, for a total of 46 learnable parameters.



# Artificial Neural Networks

## Learning the Optimal Parameters

- ANN is essentially a function  $f_{\theta}: \mathcal{X} \rightarrow \mathcal{Y}$  mapping feature vectors  $x_i \in \mathcal{X}$  to predicted output vectors  $\hat{y}_i \in \mathcal{Y}$  based on a set of parameters  $\theta = \{W, b\}$
- Let  $L(\theta)$  be the loss function that measures the difference between predictions  $\hat{y}$  and desired outputs  $y$
- We want to find the **optimum**  $\theta^*$  such that:  
$$\theta^* = \operatorname{argmin}_{\theta} L(\theta)$$
- Considering the change in the loss function with respect to  $\theta$ , the loss function is minimum when  $\frac{dL(\theta)}{d\theta} = 0$ .
- Finding  $\theta^*$  can be achieved by backpropagation and gradient descent



# Artificial Neural Networks

## Quick Review of Derivatives

Derivative of a function,  $f(\theta)$ :

- how fast  $f$  changes around  $\theta$
- will  $f$  increase or decrease if we increase  $\theta$
- is  $\theta$  higher or lower than it should be ( $\theta^*$ ); shall we make it bigger or smaller to be where it should be

$$\frac{\partial}{\partial \theta} f(g(\theta)) = \frac{\partial}{\partial g} f(g(\theta)) \cdot \frac{\partial}{\partial \theta} g(\theta) \text{ chain rule}$$

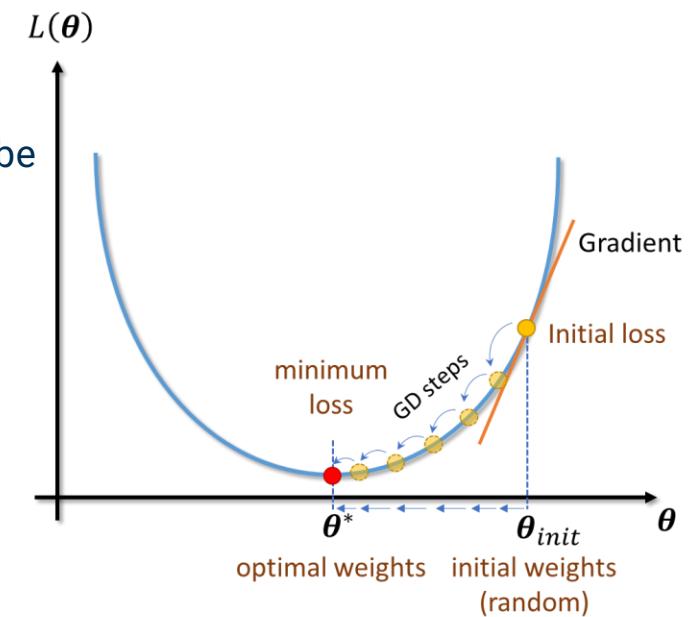
$$\frac{\partial}{\partial \theta} \sum_i f_i(\theta) = \sum_i \frac{\partial}{\partial \theta} f_i(\theta) \text{ derivative of the sum is the sum of the derivatives}$$

$$\frac{\partial}{\partial \theta_k} \sum_i a_i f(\theta_i) = a_k \frac{\partial}{\partial \theta_k} f(\theta_k) \text{ derivative wrt one element of the sum collapses the sum}$$

The sigmoid has a special property:

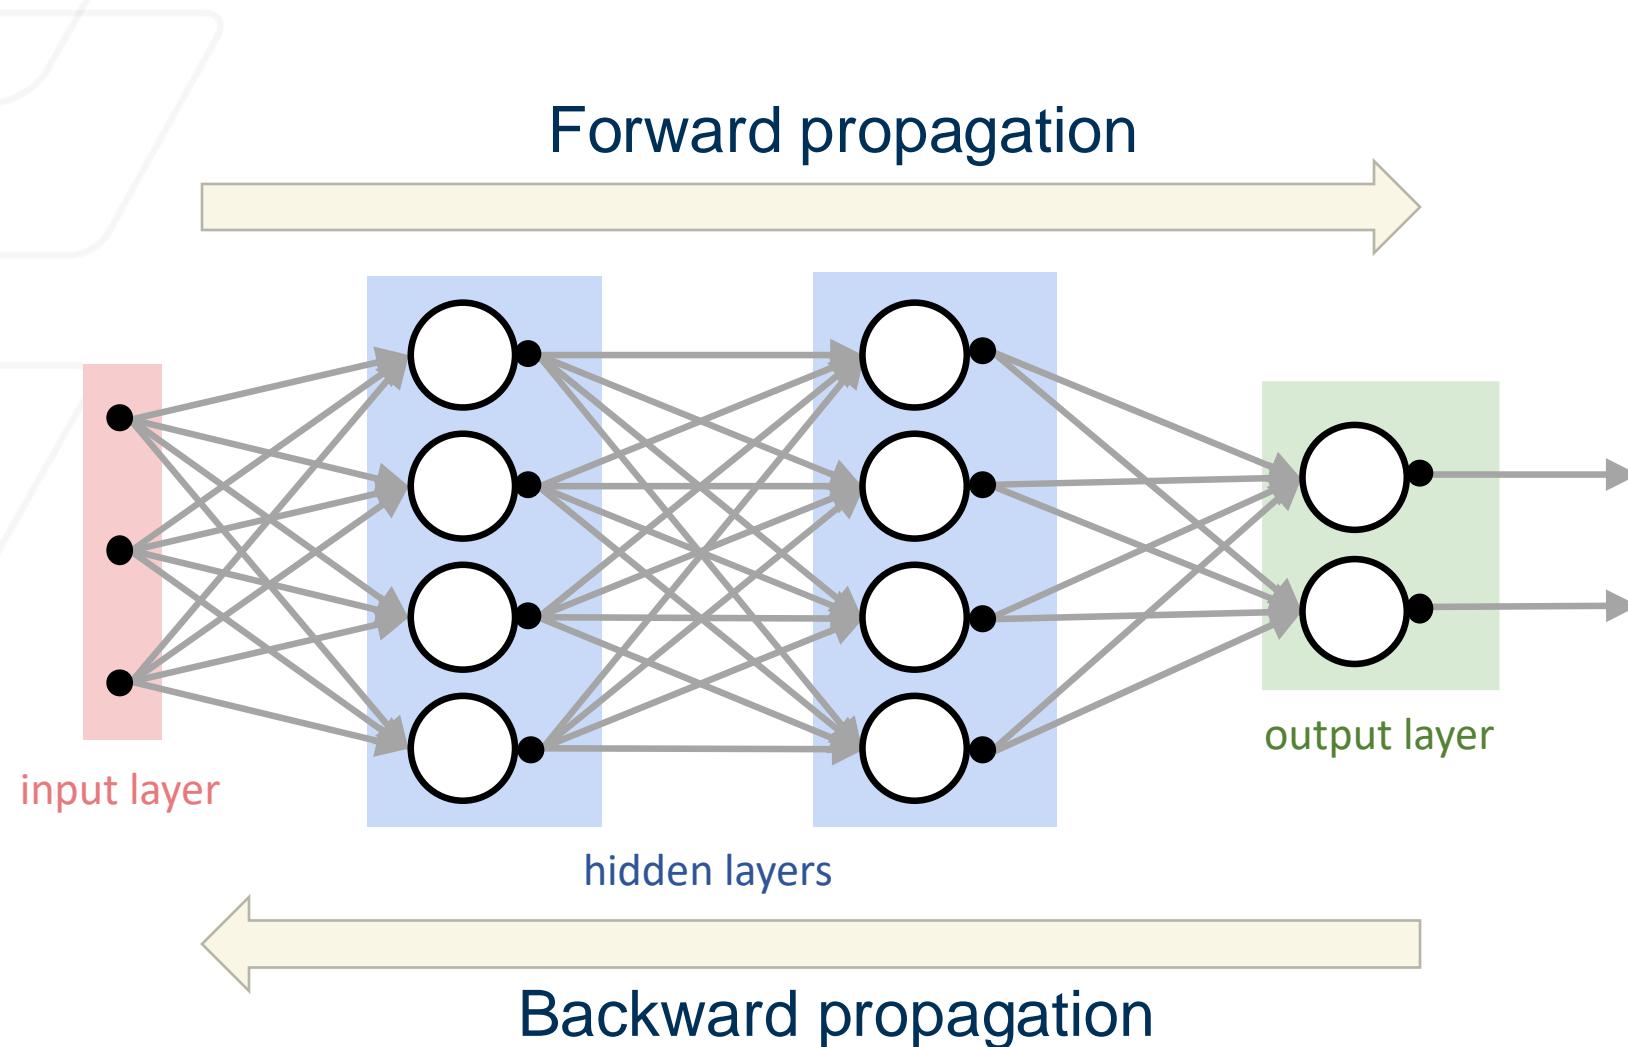
$$\sigma'(x) = \frac{\partial}{\partial x} (1 + e^{-x})^{-1} = -(1 + e^{-x})^{-2}(-e^{-x}) = \frac{1}{1 + e^{-x}} \frac{e^{-x}}{1 + e^{-x}} = \sigma(x)(1 - \sigma(x))$$

$$\frac{\partial}{\partial x} \sigma(f_x) = \sigma(f_x)(1 - \sigma(f_x)) \frac{\partial}{\partial x} f_x$$

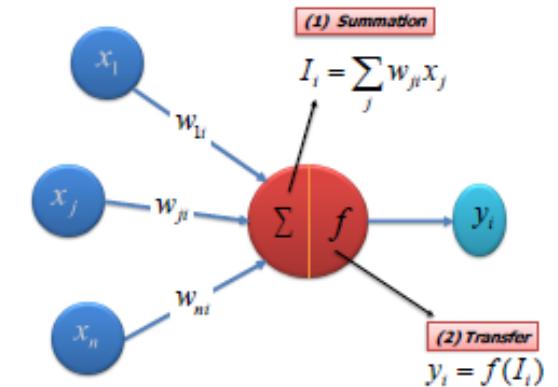


# Backpropagation

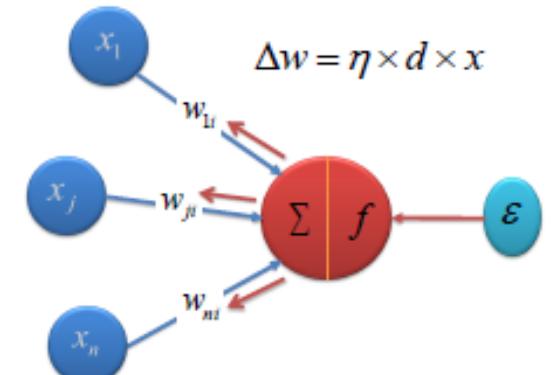
## Terminologies



Feedforward Input Data

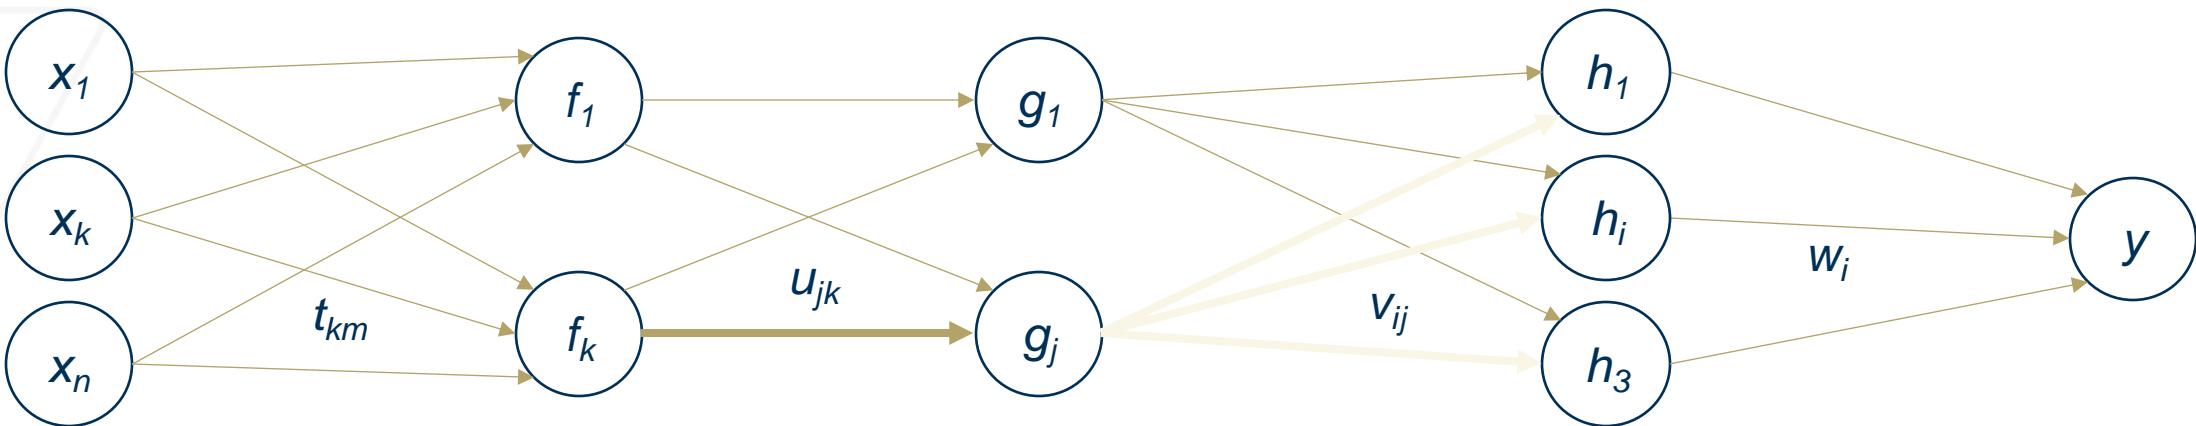


Backward Error Propagation



# Artificial Neural Networks

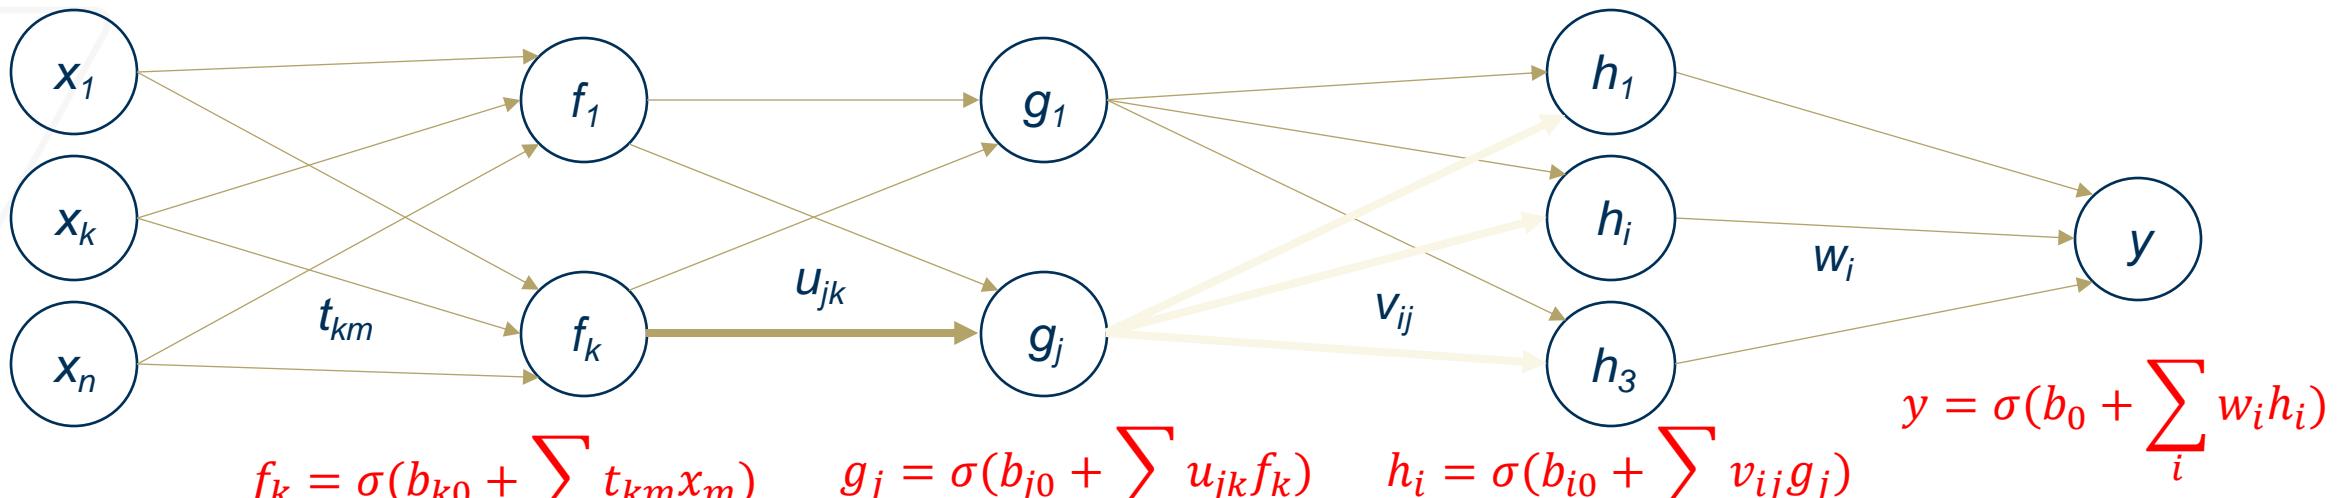
## Backpropagation



1. For a new sample  $x = [x_1, \dots, x_n]$
2. Feed forward: compute  $g_j$  based on units  $f_k$  from the previous layer using  $g_j = \sigma(b_{j0} + \sum_k u_{jk} f_k)$
3. Predict  $y$  as  $y^*$
4. Backpropagate error:  $e = y - y^*$ 
  - 1. Determine whether we need  $g_j$  to be lower or higher (assuming we already determined the better values for  $h$  from a previous step) by  $\frac{\partial e}{\partial g_j} = \sum_i \sigma'(h_i) v_{ij} \frac{\partial e}{\partial h_i}$ , or in other words by determining how  $h_i$  will change as  $g_j$  changes AND by inspecting if  $h_i$  was too high or too low
  - 2. Update weights: update the weight  $u_{jk}$  that feeds  $g_j$  by  $\frac{\partial e}{\partial u_{jk}} = \frac{\partial e}{\partial g_j} \sigma'(g_j) f_k$ , or in other words by determining if we want  $g_j$  to be higher or lower AND by determining how  $g_j$  will change if  $u_{jk}$  becomes higher or lower; now update the weight with a learning rate multiplier of  $\frac{\partial e}{\partial u_{jk}}$

# Artificial Neural Networks

## Backpropagation

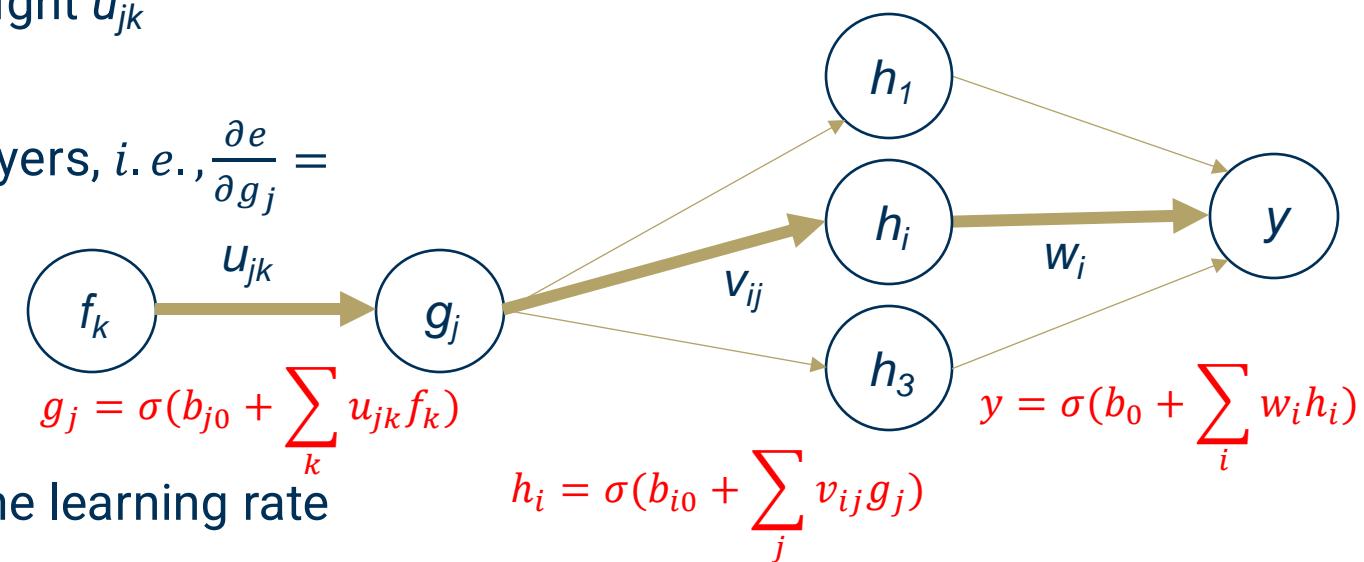


- $e = \frac{1}{2}(y - y^*)^2$   $f_k = \sigma(b_{k0} + \sum_m t_{km}x_m)$   $g_j = \sigma(b_{j0} + \sum_k u_{jk}f_k)$   $h_i = \sigma(b_{i0} + \sum_j v_{ij}g_j)$
- $\frac{\partial e}{\partial h_i} = (y - y^*) \frac{\partial y}{\partial h_i} = (y - y^*) y(1 - y)w_i$ ; the sum goes away because we are differentiating wrt to one of the elements
- $\frac{\partial e}{\partial g_j} = (y - y^*) \frac{\partial y}{\partial g_j} = (y - y^*) y(1 - y) \sum_i w_i \frac{\partial h_i}{\partial g_j} = (y - y^*) y(1 - y) \sum_i w_i h_i (1 - h_i) v_{ij}$
- By observing the nesting pattern, we can write  $\frac{\partial e}{\partial g_j} = \sum_i h_i (1 - h_i) v_{ij} \frac{\partial e}{\partial h_i}$  and so on
- $\frac{\partial e}{\partial u} = (y - y^*) \frac{\partial y}{\partial u} = (y - y^*) y(1 - y) \sum_i w_i h_i (1 - h_i) \sum_j \frac{\partial g_j}{\partial u} = (y - y^*) y(1 - y) \sum_i w_i h_i (1 - h_i) v_{ij} g_j (1 - g_j) f_k = g_j (1 - g_j) f_k \frac{\partial e}{\partial g_j}$

# Artificial Neural Networks

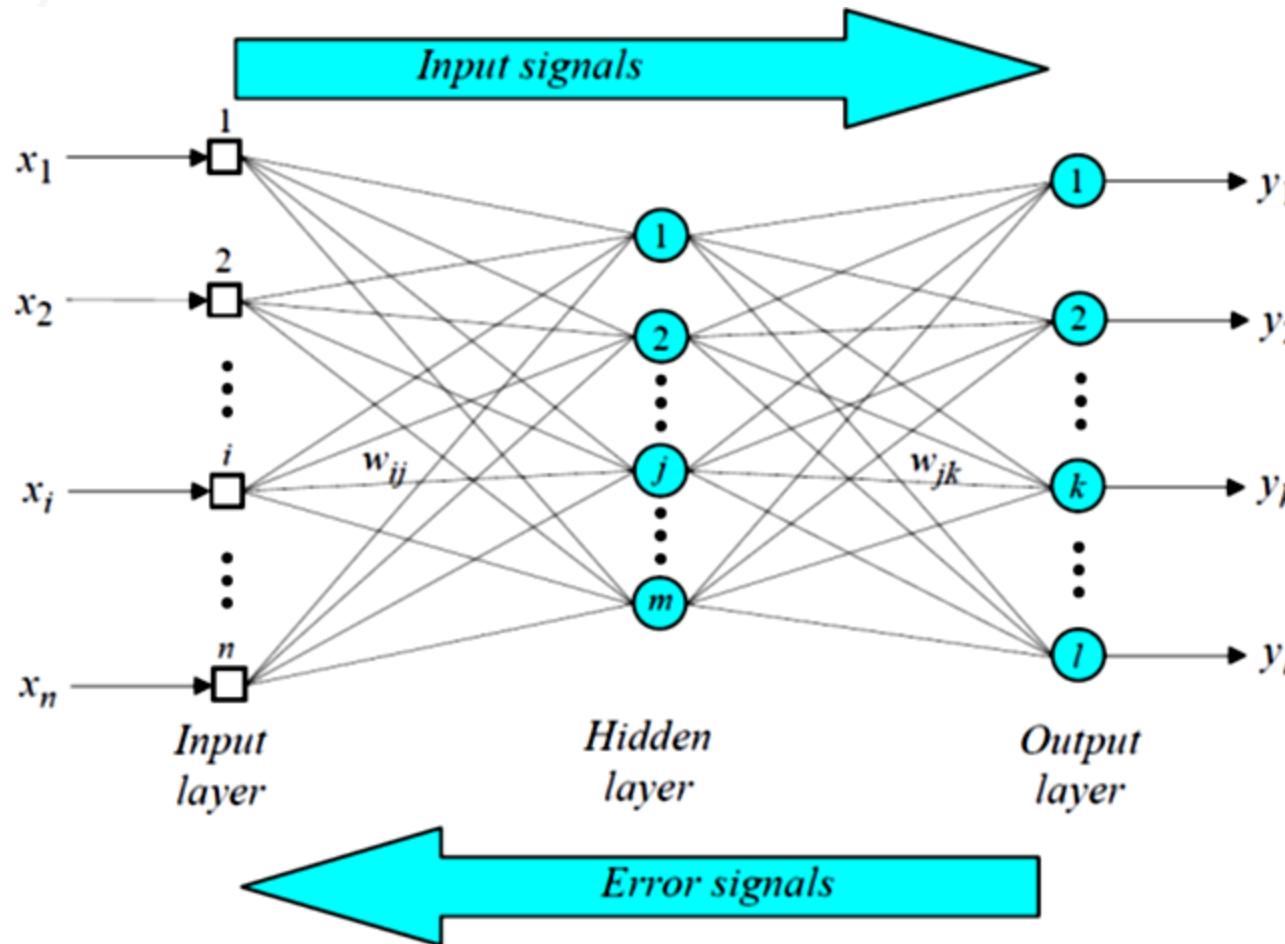
## Backpropagation Summary

- **Feed forward propagate** the sample  $x$ , compute activation outputs of neurons from previous layers and compute the predicted output  $y$
- Compute the **backpropagation error**:  $e = L(y, y^*)$ , where  $L$  is the loss function
- Assume using sigmoid activation  $\sigma$ , for a weight  $u_{jk}$  compute  $\frac{\partial e}{\partial u_{jk}}$  using **chain rule**:
  - Compute  $\frac{\partial e}{\partial g_j}$  by carrying  $\frac{\partial e}{\partial h_i}$  from later layers, i.e.,  $\frac{\partial e}{\partial g_j} = \sum_i h_i(1 - h_i)v_{ij}\frac{\partial e}{\partial h_i}$
  - Compute  $\frac{\partial e}{\partial u_{jk}} = g_j(1 - g_j)f_k \frac{\partial e}{\partial g_j}$
- Update weights using **gradient descent**
  - $u_{jk}(t + 1) = u_{jk}(t) - \alpha \frac{\partial e}{\partial u_{jk}}$ , where  $\alpha$  is the learning rate
- Iterate the above steps until the error  $e$  converges

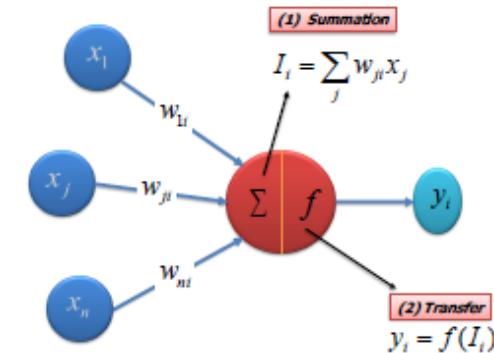


# Artificial Neural Networks

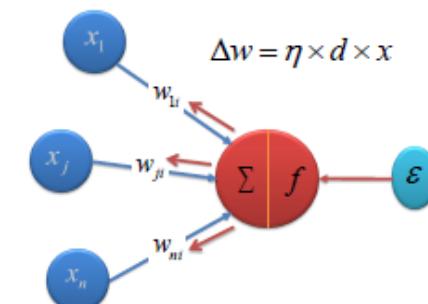
## Backpropagation Summary



Feedforward Input Data



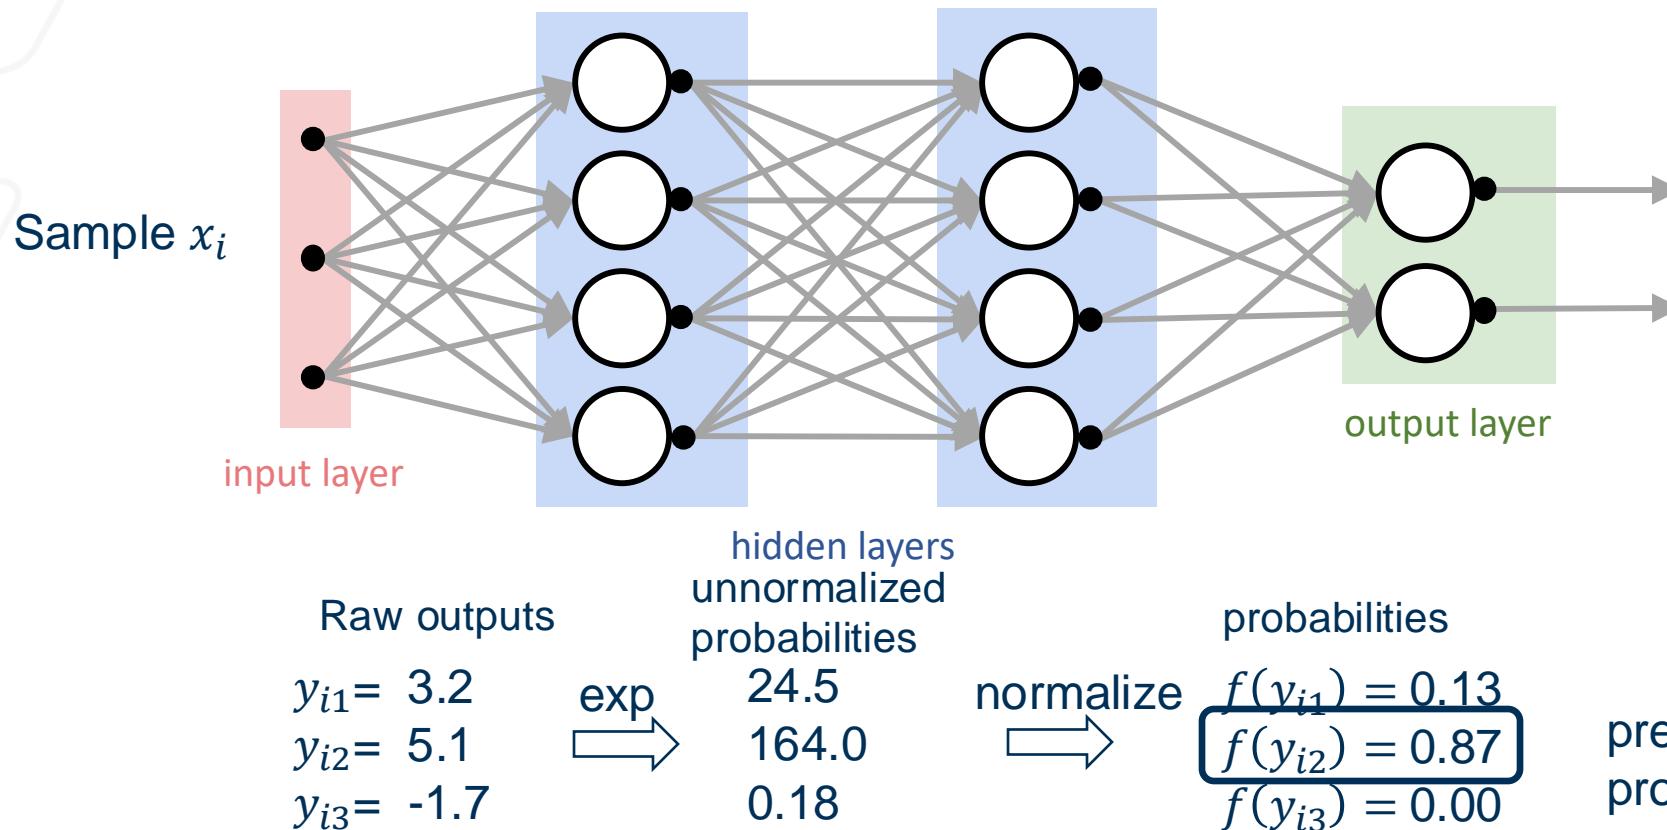
Backward Error Propagation



# Artificial Neural Networks

## Backpropagation Summary

During the inference, predict the class with the highest probability with SoftMax activation:  $\text{argmax}_j f(y_{ij})$



Assume the raw outputs  $y_i = [y_{i1}, y_{i2}, y_{i3}] = [3.2, 5.1, -1.7]$

**SoftMax activation:**

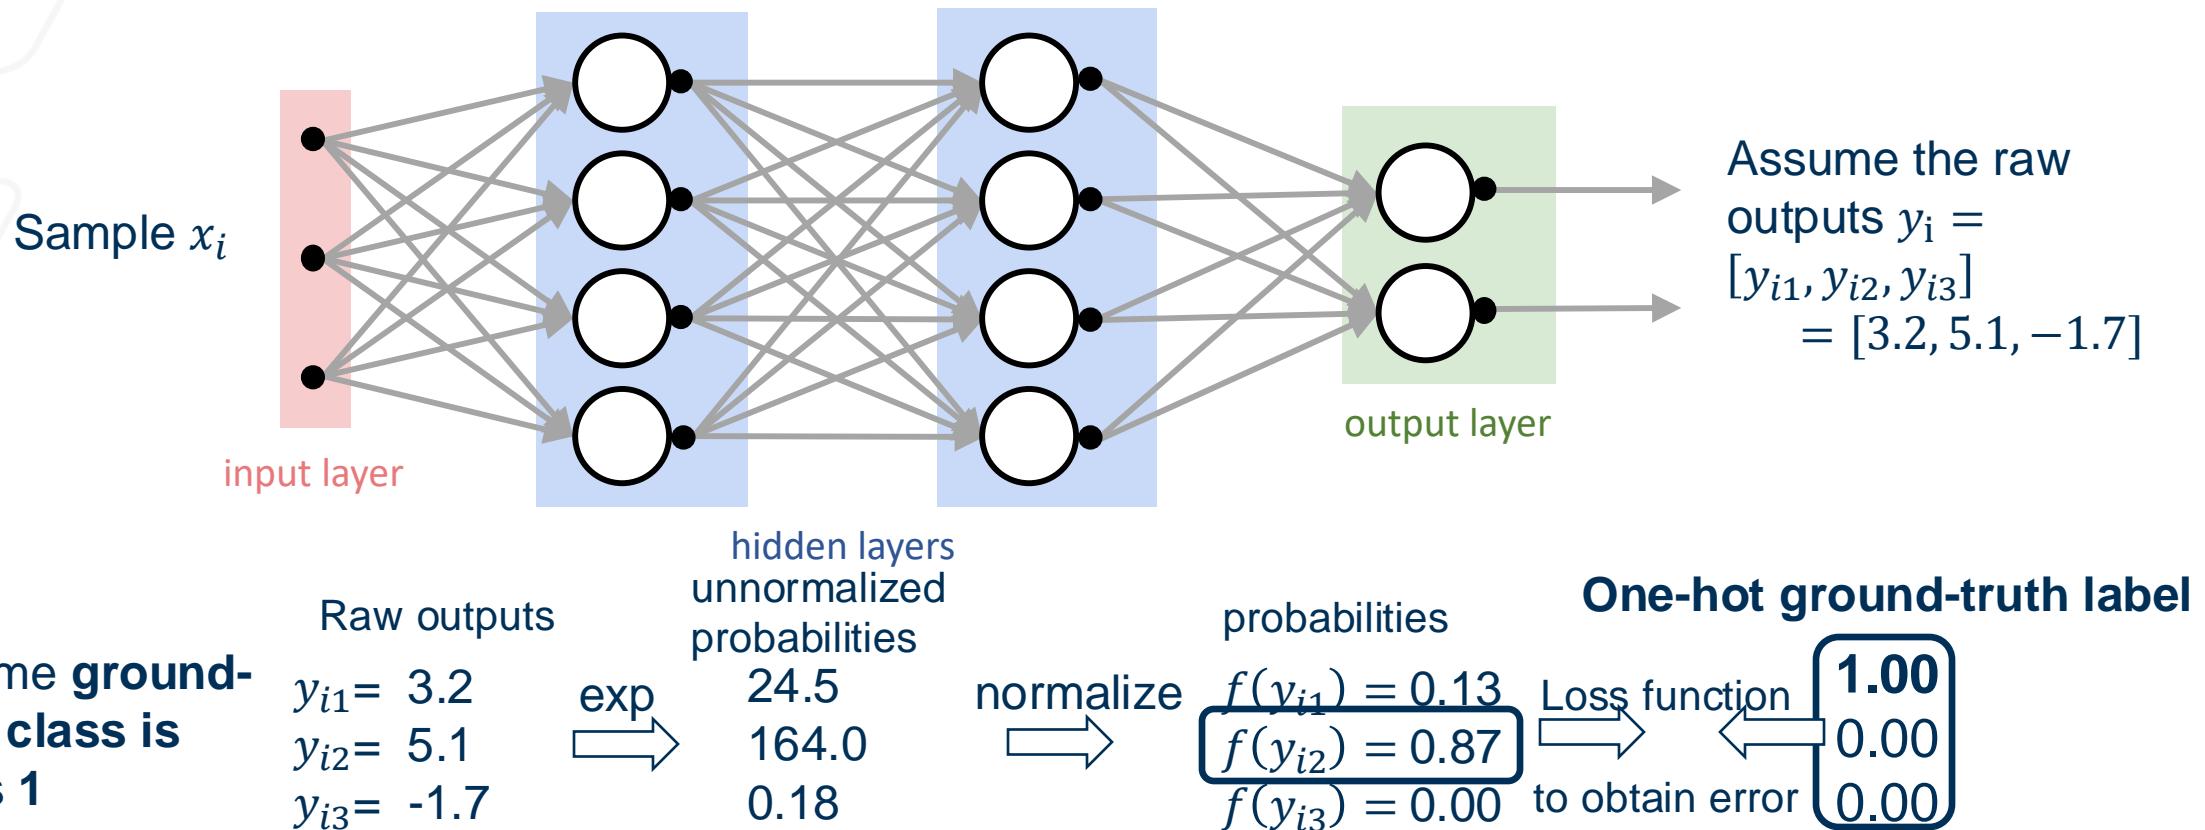
$$f(y_{ij}) = \frac{e^{y_{ij}}}{\sum_k e^{y_{ik}}}$$

predict class 2 with highest probability

# Artificial Neural Networks

# Backpropagation Summary

During training, for every sample, we set the ground-truth label as a one-hot vector  $[1, 0, \dots, 0]^T$  with 1 for the correct class and 0 for every other class. Backpropagate the error and repeat for every sample.



# Artificial Neural Networks

## Image Classification

**Image Classification:** mapping the image pixels to probabilities for each category

For simplicity, we take a linear function:

$$\hat{Y} = \phi(XW^T + b^T)$$

where

- $X \in \mathbb{R}^{N \times P}$ : dataset containing  $N$  vectorized images
- $P \in \mathbb{R}^{H \times W \times C}$ : the number of pixels (*features*) of each image
- $\hat{Y} \in \mathbb{R}^{N \times P^{(k)}}$ : the associated probabilities of each category  $1, 2, \dots, P^{(k)}$
- $\phi(X, W, b)$ : the activation function
- $W \in \mathbb{R}^{P \times P^{(k)}}$ : the weight matrix
- $b \in \mathbb{R}^{P^{(k)} \times 1}$ : the bias vector

# Artificial Neural Networks

## Image Classification

Single image binary classification (predicting dog/cat)



stretch pixels into single column vector

$$\phi \left( \begin{bmatrix} 0.2 & 2.1 \\ -0.5 & 0.0 \\ 0.1 & 0.25 \\ 2.0 & 0.2 \\ 1.5 & -0.3 \\ \vdots & \vdots \\ 1.3 & 1.2 \end{bmatrix}^T \begin{bmatrix} 56 \\ 231 \\ 24 \\ 188 \\ 75 \\ \vdots \\ 32 \end{bmatrix} + \begin{bmatrix} 0.2 \\ 2.4 \end{bmatrix} \right) = \begin{bmatrix} 0.8 \\ 0.2 \end{bmatrix}$$

Cat score  
Dog score

$w^T \in \mathbb{R}^{2 \times (32)(32)(3)}$   $x_i \in \mathbb{R}^{(32)(32)(3) \times 1}$

$b \in \mathbb{R}^{2 \times 1}$   $y_i \in \mathbb{R}^{2 \times 1}$

# Artificial Neural Networks

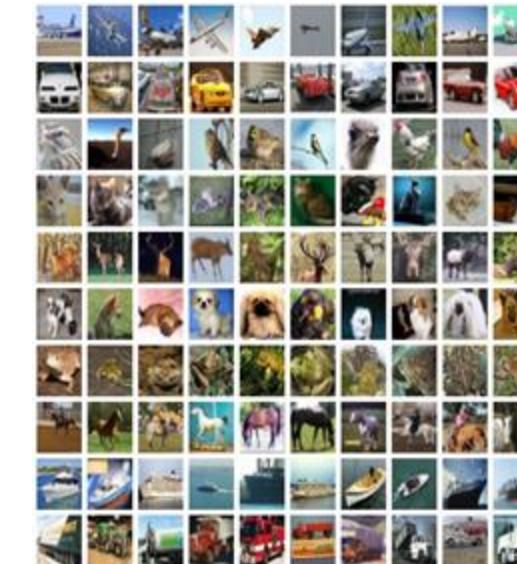
## Training MLP for Image Classification

### Example Datasets:

- MNIST ([hand-written digits](#))
  - # total: 70,000 grayscale images
  - # classes: 10
  - # size: 28x28
  - # training samples: 60,000 images
  - # test samples: 10,000 images
- CIFAR-10 (subsets of the [80 million tiny images](#))
  - # total: 60,000 color images
  - # classes: 10
  - # size: 32x32
  - # training samples: 50,000 images
  - # test samples: 10,000 images

0 0 0 0 0 0 0 0 0  
1 1 1 1 1 1 1 1 1  
2 2 2 2 2 2 2 2 2  
3 3 3 3 3 3 3 3 3  
4 4 4 4 4 4 4 4 4  
5 5 5 5 5 5 5 5 5  
6 6 6 6 6 6 6 6 6  
7 7 7 7 7 7 7 7 7  
8 8 8 8 8 8 8 8 8  
9 9 9 9 9 9 9 9 9

[https://www.researchgate.net/figure/Sample-images-of-MNIST-data\\_fig3\\_222834590](https://www.researchgate.net/figure/Sample-images-of-MNIST-data_fig3_222834590)



<https://www.kaggle.com/c/cifar-10>

## Appendix A: Notations

- $x_i$ : a single feature
- $\boldsymbol{x}_i$ : feature vector (data sample)
- $\boldsymbol{X}$ : matrix of feature vectors (dataset)
- $N$ : number of data samples
- $m$ : degree of polynomial
- $P$ : number of features in a feature vector
- $\theta_i$ : a single model coefficient (parameter)
- $\boldsymbol{\theta}$ : coefficient vector
- $\varepsilon$ : error margin
- $\alpha$ : learning rate
- $\gamma$ : bias factor
- Bold letter/symbol: vector
- Bold capital letters/symbol: matrix