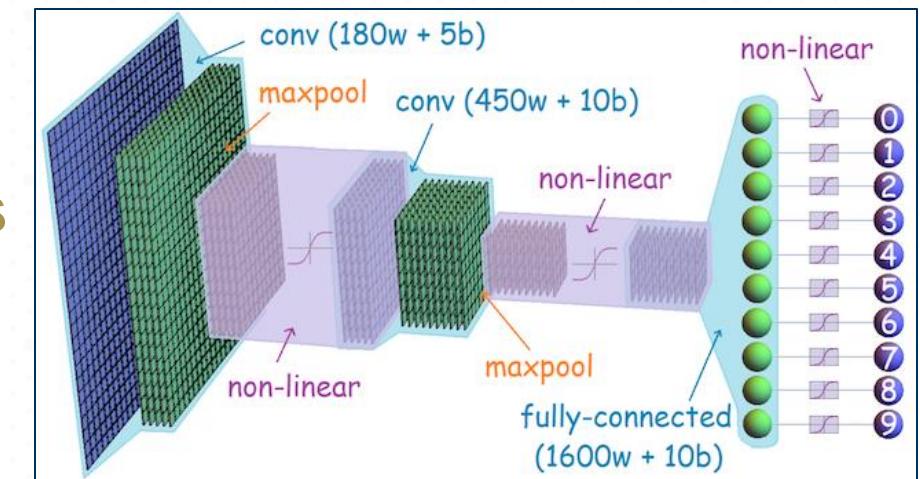


ECE 4803/8803: Fundamentals of Machine Learning (FunML)

Spring 2024

Lecture 13: Neural Networks



Logistics

- Reminders
 - HW 4, Deadline: Oct 11
- Midterms
 - Avg: 86
 - Std. Dev: 12
 - Maximum: 100 (4 students)
 - Your scores are posted.
 - You have 1 week (Oct 14) to ask for any regrades. Post a private question on Piazza (preferable) or email me (mohit.p@gatech.edu) with ECE 4252/8803 in subject line.

Logistics for 8803 Section

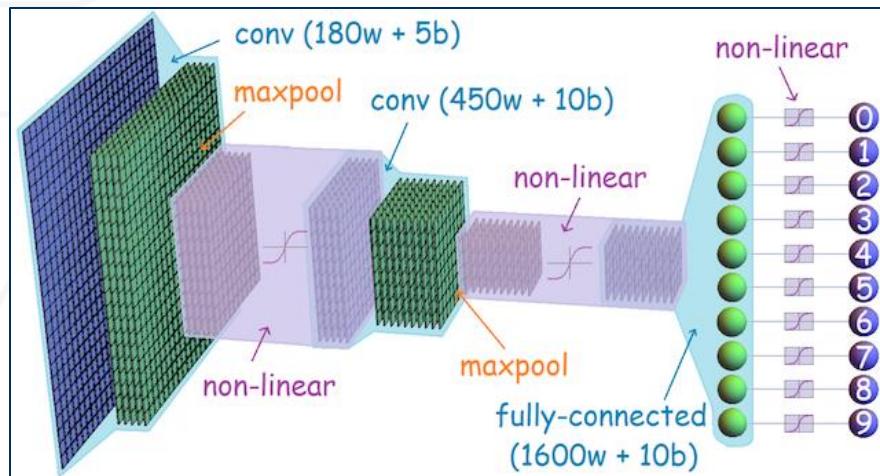
- Project: Biomarker Analysis on OLIVES dataset
- Logistics: Groups of 3
 - For a grad course, 20% project grade is roughly 40-50 hrs per person. Hence, project outputs must reflect 120-150 hrs of combined work
 - Grading is done on individual contributions
- Dataset download: <https://zenodo.org/records/7105232> (also available on AI makerspace)
- Goal: Detect Biomarkers given OCT images, and clinical labels
 - You are free to choose your train/test splits
 - You are free to choose your methods/evaluation
- Challenges you will be evaluated on
 - How will you handle multi-modal data (images, text etc.)?
 - What methods will you employ for evaluation?
 - **What analysis will you perform on results?**
 - How good are your results?

Logistics for 8803 Section

- Some resources:
 - Original paper: <https://arxiv.org/pdf/2209.11195>
 - 2023 SPS VIP competition: <https://alregib.ece.gatech.edu/competitions/2023-vip-cup/>
 - Competition structured dataset: <https://zenodo.org/records/8040573>
 - Example paper from the competition: <https://arxiv.org/pdf/2310.14005>
- Group members
 - Form groups here: <https://docs.google.com/spreadsheets/d/1LriuGQuiF4RDL8Db0lg-ZXMCrhokEjFPCT0AILJfOPs/edit?usp=sharing>
- Deliverables: (More instructions when we approach the deadline)
 - 25 Oct: “Progress” Report (PDF), needs to have the following
 - Team members (Representative team member uploads the PDF)
 - Problem description
 - Proposed approach (no results necessary)
 - No specific format
 - 29 Nov: Term Project Paper (PDF)
 - Max 4 pages of content + unlimited pages for references
 - IEEE conference style paper format
 - Code on Github/HuggingFace

Overview

In this Lecture..



Review: Artificial Neural Networks

- Activation Function
- Perceptron Network
- Multi-layer ANN

Backpropagation

- Feedforward and Backward Error Propagation
- Learning Algorithm
- Image Classification using ANNs
- Mini-batch Gradient descent
- Linear Classifier
- Logistic Regression

PyTorch

- Autograd
- Built-in Modules
- Examples

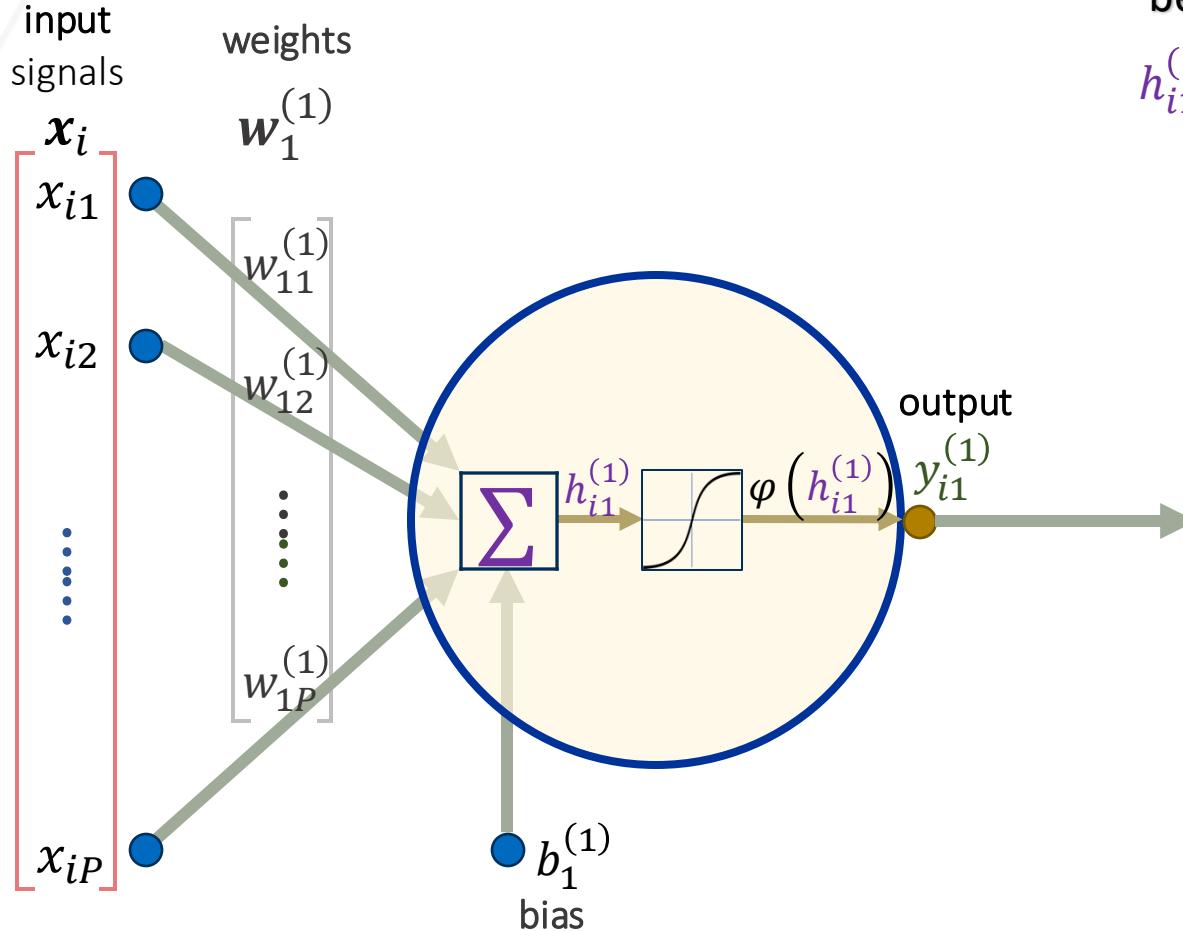
Classifier Comparison

Methods	Assumptions on Feat. Dist.	Feat. Normalization	Cost Function	Regularization	Linear Classifier	Prob. View of Prediction	Generative/Discriminative	Parametric/Non-parametric	Overfitting
Logistic Regression	No	Required	BCE (convex)	Additional term	Linear	Yes	Discriminative	Parametric	Not often
K Nearest Neighbors	No	Required	N/A	N/A	Non-linear	N/A	Discriminative	Non-parametric	when k is too small
Decision Trees	No	Not Required	N/A	N/A	Non-linear	N/A	Discriminative	Non-parametric	with large depth
Support Vector Machines	No	Required	Hinge (convex)	C (control robustness)	Linear/Non-linear(kernel)	N/A	Discriminative	Parametric	Not often
Naïve Bayes	Conditional independent	Not Required	N/A	N/A	Non-linear /Linear (Gaussian)	Yes	Generative	Parametric	Not often
Artificial Neural Networks	No	Required	Non-convex	Additional term	Non-linear	Yes	Discriminative/Generative	Parametric	with many layers

Artificial Neural Networks

The Perceptron

- Single-layer Perceptron (SLP)



weighted input to the neuron
before activation φ

$$h_{i1}^{(1)} = (\mathbf{w}_j^{(1)})^T \mathbf{x}_i + b_1^{(1)}$$

output of the neuron
after activation φ

$$\varphi(h_{i1}^{(1)})$$

The simplest form of the perceptron uses linear activation $\varphi(h_{i1}^{(1)}) = h_{i1}^{(1)}$, the outputs are binary (i.e. $y_{i1}^{(1)} \in \{1, -1\}$):

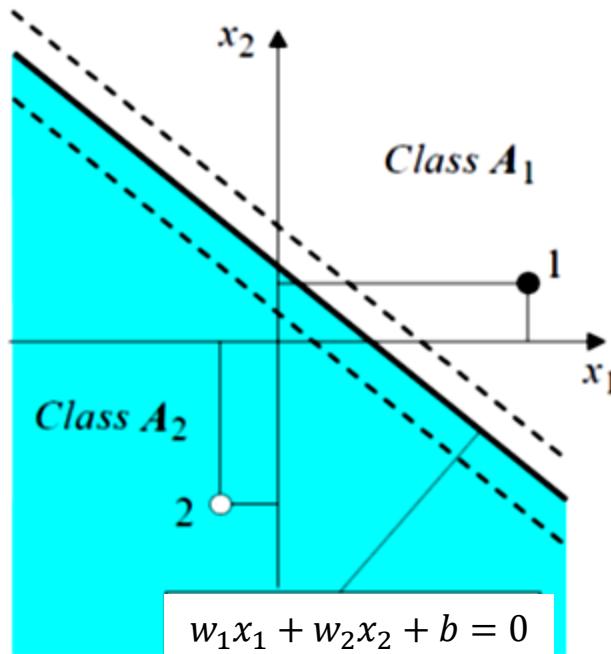
$$y_{i1}^{(1)} = \begin{cases} +1, & \text{if } (\mathbf{w}_j^{(1)})^T \mathbf{x}_i + b_1^{(1)} \geq 0 \\ -1, & \text{if } (\mathbf{w}_j^{(1)})^T \mathbf{x}_i + b_1^{(1)} < 0 \end{cases}$$

The above model applies to linearly separable data

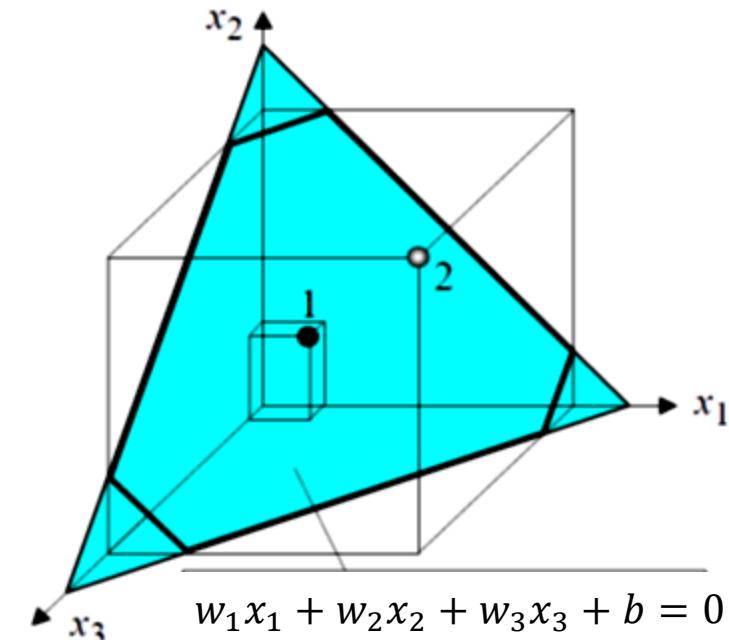
Artificial Neural Networks

The SLP Perceptron Model

- Linear separability in the perceptron



(a) Two-input perceptron.

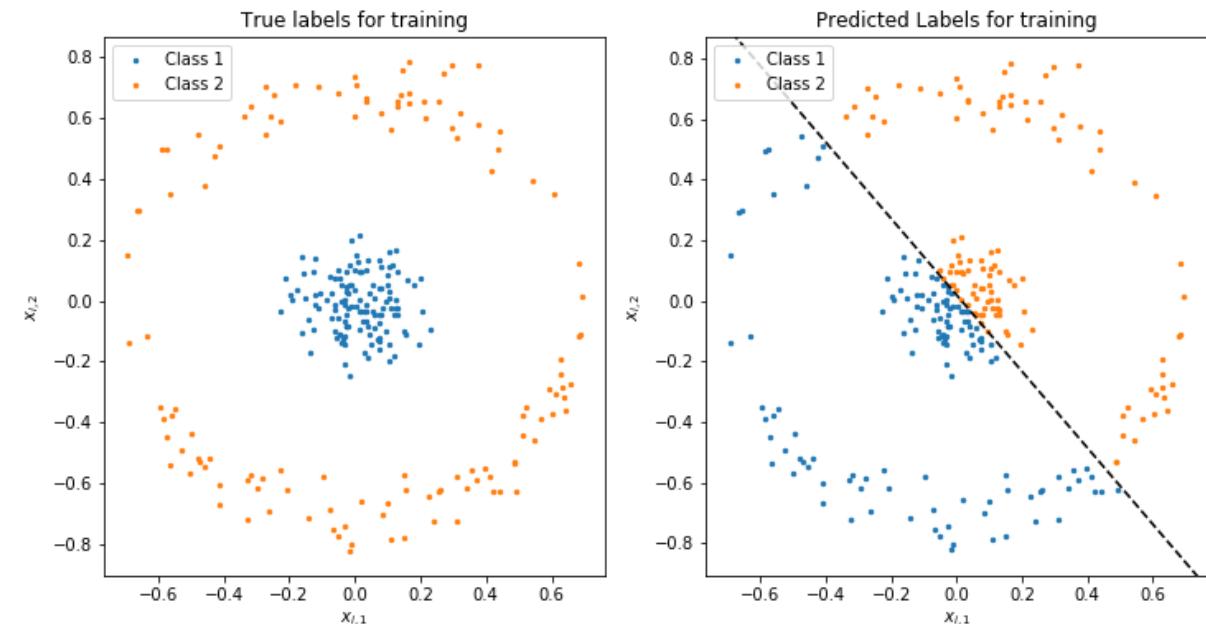


(b) Three-input perceptron.

Artificial Neural Networks

Hidden Layers

- Single neurons with **linear activation** fail to classify non-linearly separable dataset
- Non-linearly separable data requires
 - Non-linear Activation
 - Multi-layer networks of neurons (Multi-layer perceptron)

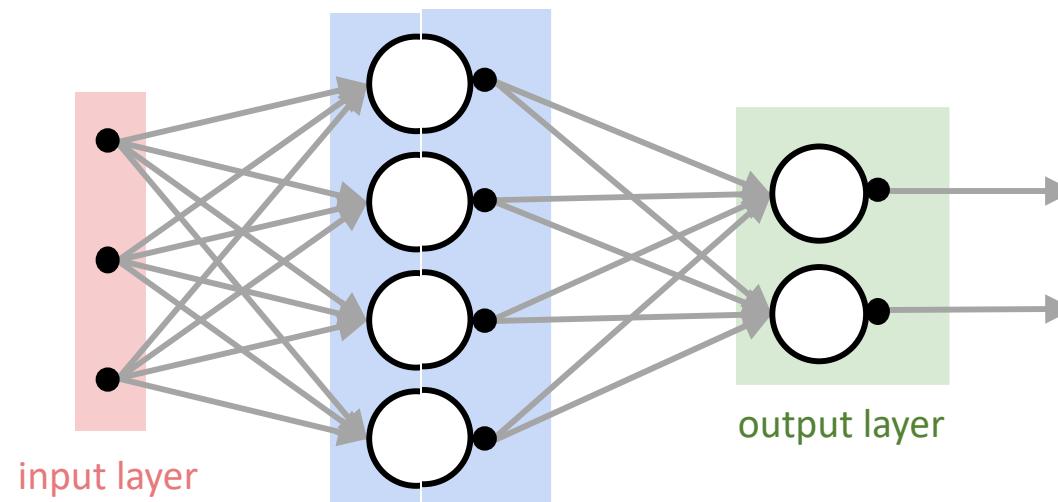


SLP with linear activation classifying non-linearly separable data

Artificial Neural Networks

Hidden Layers

- A NN with one hidden layer can represent:
 - any bounded continuous function (to some arbitrary ϵ) [Universal Approximation Theorem, Cybenko 1989]
 - any Boolean Function, but it requires 2^k hidden units for $1K$ inputs



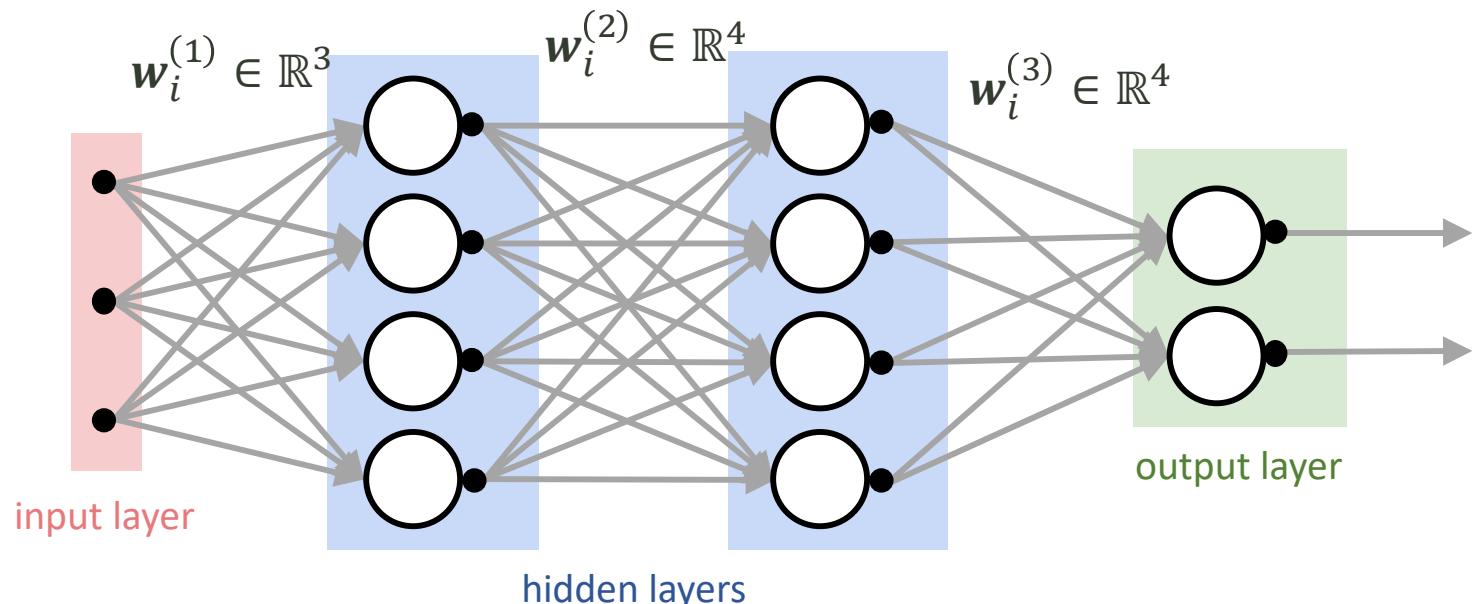
Artificial Neural Networks

Multiple Hidden Layers

Layer-wise organization: MLP consists of *fully-connected* layers in which neurons between two adjacent layers are fully pairwise connected, while neurons within a single layer share no connections.

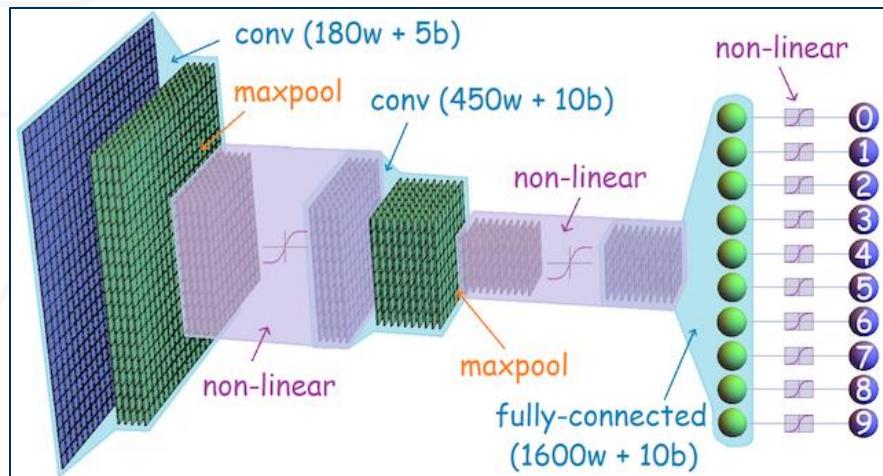
Sizing MLP: the number of neurons or the number of parameters (more commonly)

- The network has $4 + 4 + 2 = 10$ neurons, $[3 \times 4] + [4 \times 4] + [4 \times 2] = 12 + 16 + 8 = 36$ weights and $4 + 4 + 2 = 10$ biases, for a total of 46 learnable parameters.



Overview

In this Lecture..



Review: Artificial Neural Networks

- Activation Function
- Perceptron Network
- Multi-layer ANN

Backpropagation

- Feedforward and Backward Error Propagation
- Learning Algorithm
- Image Classification using ANNs
- Mini-batch Gradient descent
- Linear Classifier
- Logistic Regression

PyTorch

- Autograd
- Built-in Modules
- Examples

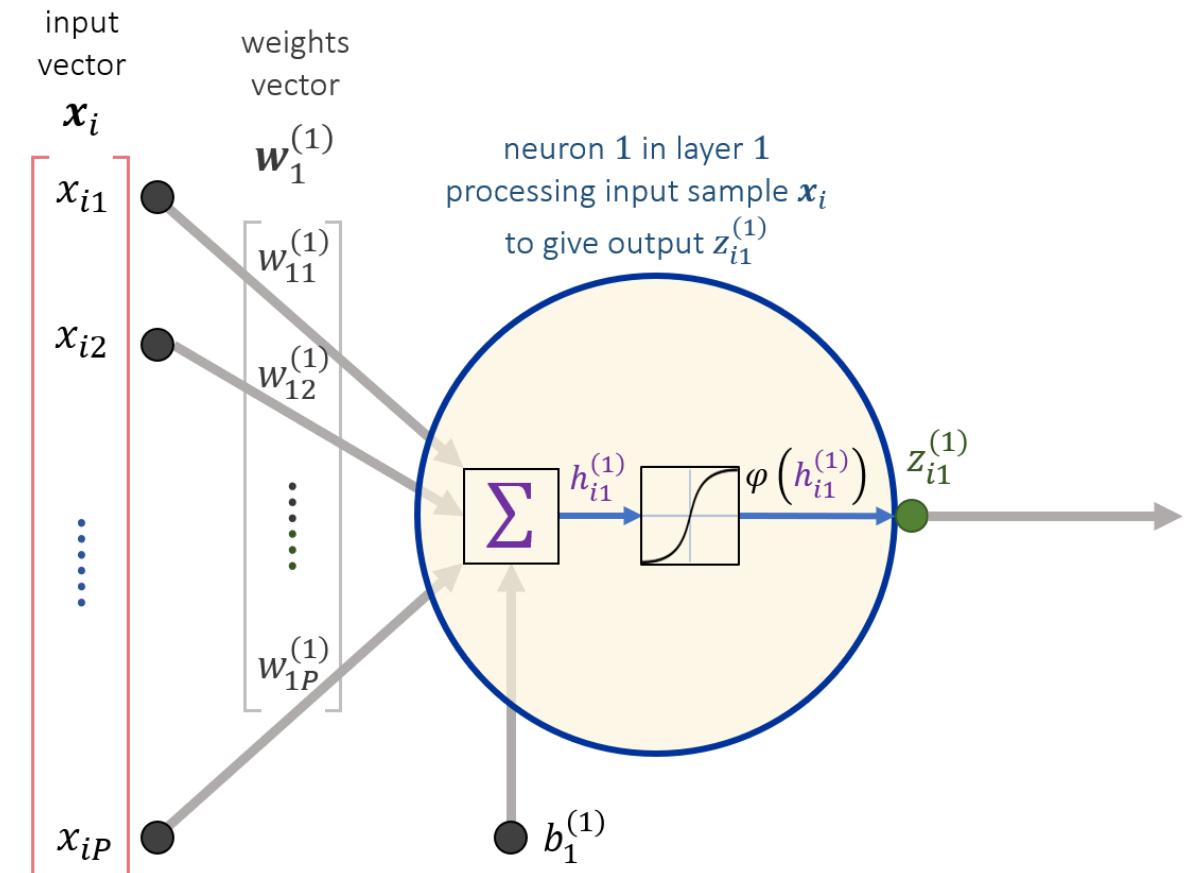
Backpropagation

Notations

Neuron 1 in layer 1:

- Let $\mathbf{w}_1^{(1)} = \begin{bmatrix} w_{11}^{(1)} \\ w_{12}^{(1)} \\ \vdots \\ w_{1P}^{(1)} \end{bmatrix}$ be the neuron's weights vector, where P is the number of features in the input layer
- Let $b_1^{(1)}$ be the neuron's bias
- Let $h_{i1}^{(1)} = (\mathbf{w}_j^{(1)})^T \mathbf{x}_i + b_1^{(1)}$ be the weighted input to the neuron before activation
- Let $z_{i1}^{(1)} = \varphi(h_{i1}^{(1)})$ be the activation associated with the neuron using any activation function φ such as the sigmoid function, σ :

$$z_{i1}^{(1)} = \varphi(h_{i1}^{(1)}) = \sigma(h_{i1}^{(1)}) = \frac{1}{1 + e^{-h_{i1}^{(1)}}}$$

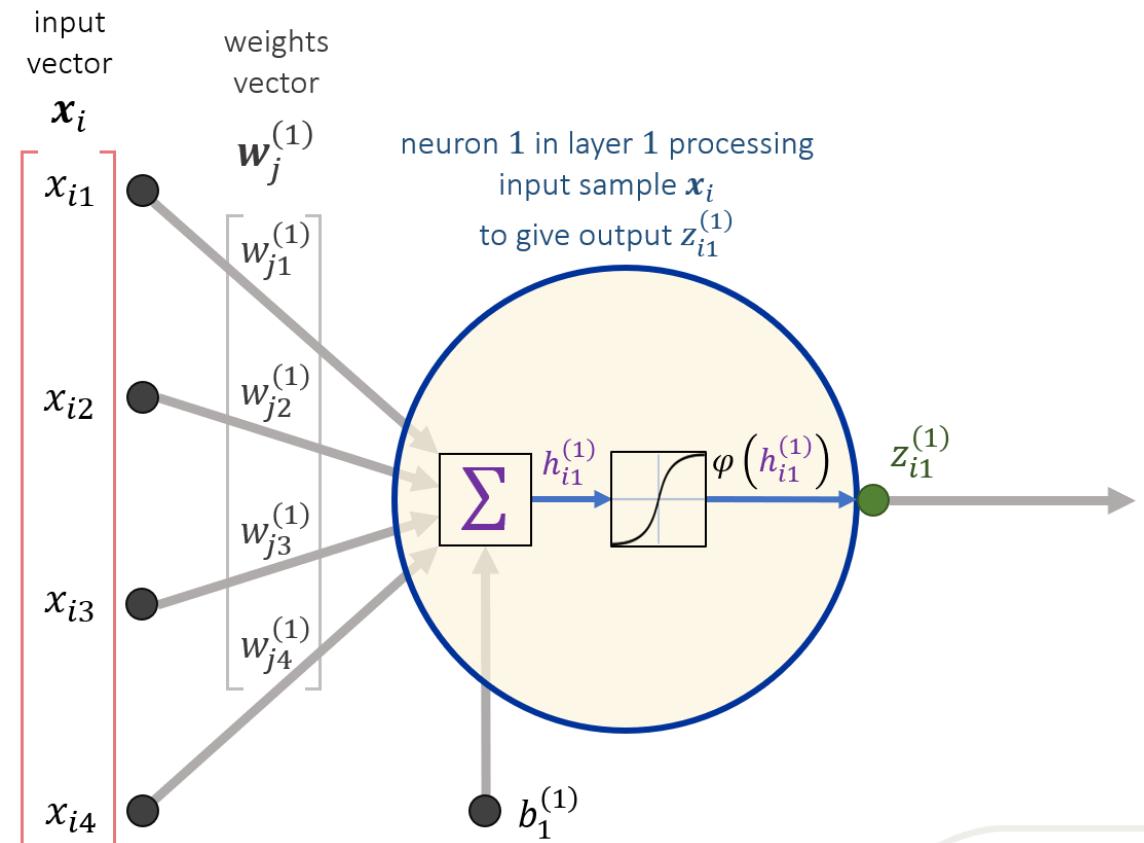


Backpropagation

Notations

Example: The 1st neuron in layer 1, processing a 4-dimensional input

- Input: $x_i = \begin{bmatrix} x_{i1} \\ x_{i2} \\ x_{i3} \\ x_{i4} \end{bmatrix}$, Weights: $w_1^{(1)} = \begin{bmatrix} w_{11}^{(1)} \\ w_{12}^{(1)} \\ w_{13}^{(1)} \\ w_{14}^{(1)} \end{bmatrix}$
- Bias: $b_1^{(1)}$
- Weighted input: $h_{i1}^{(1)} = (w_1^{(1)})^T x_i + b_1^{(1)}$
- Output: $z_{i1}^{(1)} = \varphi(h_{i1}^{(1)}) = \sigma(h_{i1}^{(1)}) = \frac{1}{1+e^{-h_{i1}^{(1)}}}$



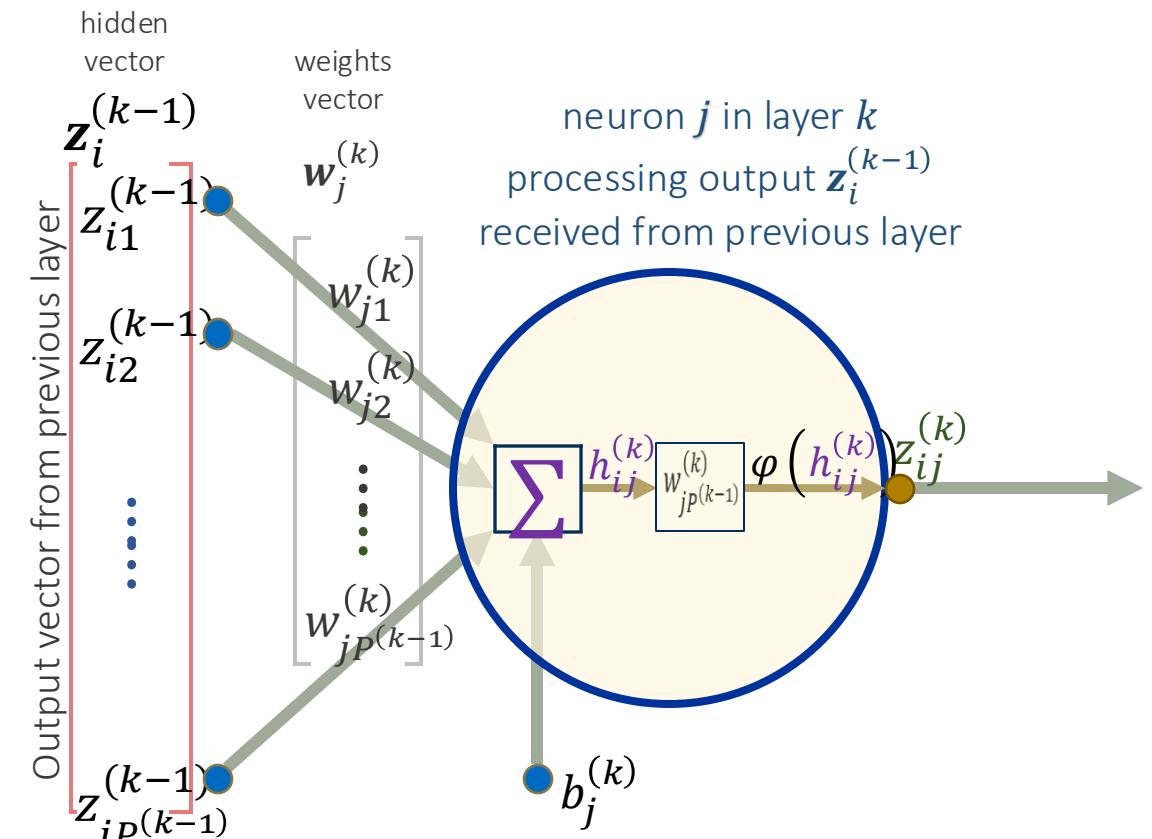
Backpropagation

Notations

General notation, neuron j in layer k :

- Let $\mathbf{w}_j^{(k)} = \begin{bmatrix} w_{j1}^{(k)} \\ w_{j2}^{(k)} \\ \vdots \\ w_{jP^{(k-1)}}^{(k)} \end{bmatrix}$ be the neuron's weights vector, where $P^{(k-1)}$ is the number of neurons in the previous layer $k - 1$
($P^{(0)} = P$, the number of features in the input vector)
- Let $b_j^{(k)}$ be the neuron's bias
- Let $h_{ij}^{(k)} = (\mathbf{w}_j^{(k)})^T \mathbf{z}_i^{(k-1)} + b_j^{(k)}$ be the weighted input to the neuron before activation
- Let $z_{ij}^{(k)} = \varphi(h_{ij}^{(k)})$ be the activation associated with the neuron using any activation function φ such as the sigmoid function, σ :

$$z_{ij}^{(k)} = \varphi(h_{ij}^{(k)}) = \sigma(h_{ij}^{(k)}) = \frac{1}{1 + e^{-h_{ij}^{(k)}}}$$

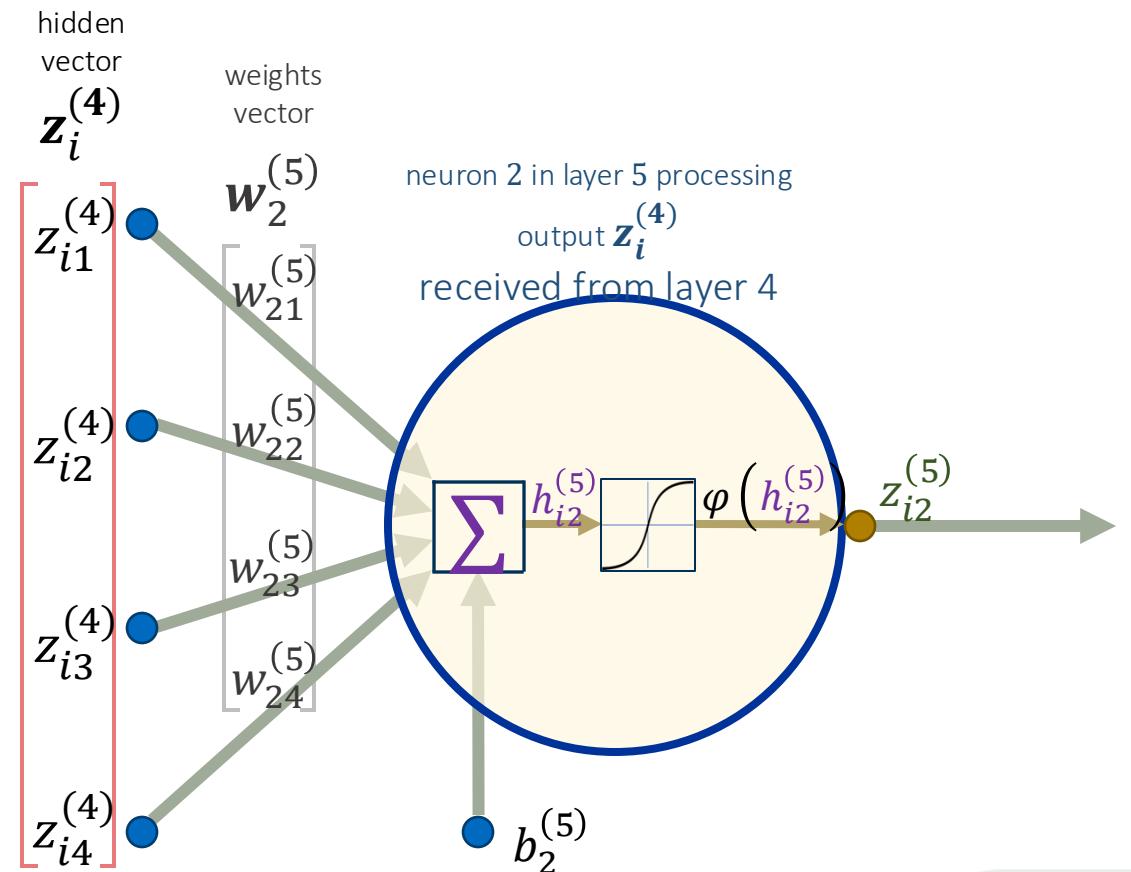


Backpropagation

Notations

Example: The 2nd neuron in layer 5, processing a 4-dimensional input

- Input: $\mathbf{z}_i^{(4)} = \begin{bmatrix} z_{i1} \\ z_{i2} \\ z_{i3} \\ z_{i4} \end{bmatrix}$, Weights: $\mathbf{w}_2^{(5)} = \begin{bmatrix} w_{21}^{(5)} \\ w_{22}^{(5)} \\ w_{23}^{(5)} \\ w_{24}^{(5)} \end{bmatrix}$
- Bias: $b_2^{(5)}$
- Weighted input: $h_{i2}^{(5)} = (\mathbf{w}_2^{(5)})^T \mathbf{z}_i^{(4)} + b_2^{(5)}$
- Output: $z_{i2}^{(5)} = \varphi(h_{i2}^{(5)}) = \sigma(h_{i2}^{(5)}) = \frac{1}{1+e^{-h_{i2}^{(5)}}}$



Backpropagation

Optimum Weights via Gradient Descent

- For a simple **linear** perceptron model, we can try to find the optimum w and b by minimizing a **Least Squares** cost function
- Let $L(\theta)$ be the **MSE loss** function defined over the entire dataset such that:

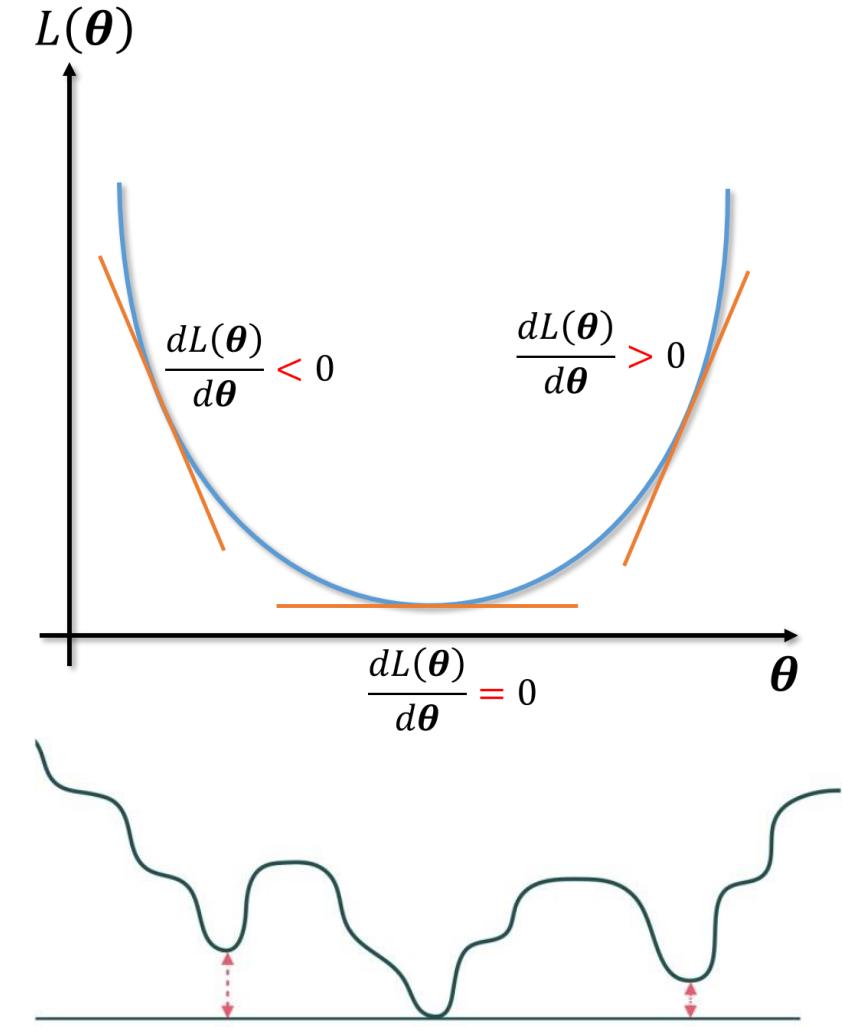
$$L(\theta) = \text{MSE}(X, \theta) = \frac{1}{N} \|\hat{Y} - Y\|_F^2$$

where Y is a matrix of all target outputs y and \hat{Y} is a matrix of all predicted outputs \hat{y} , and $\|X\|_F^2 = \text{Tr}(XX^T)$ is the Frobenius norm of a matrix which is the square root of the sum of the absolute squares of its elements.

- Let θ^* be the set of optimum weights and biases such that:

$$\theta^* = \underset{\theta}{\operatorname{argmin}} L(\theta)$$

- Considering the change in the loss function with respect to θ , the loss function is minimum when $\frac{dL(\theta)}{d\theta} = 0$.
- Finding the θ^* where the loss is minimum can be achieved either:
 - Solving the **Normal Equation**
 - Training by **Gradient Descent**



In general, cost functions of NNs are non-convex with flat areas, and many saddle points.

Backpropagation

Optimum Weights for a Single Neuron

- Recall that $X \in \mathcal{R}^{N \times P}$, $Y \in \mathcal{R}^{N \times K}$, $W \in \mathcal{R}^{P \times K}$, where we have N samples, P features, and K output classes per sample
- To simplify the math, let the bias $b_1^{(1)} = 0$, thus $\theta = W$ and $\hat{Y} = XW$
- Assume that we use MSE as the loss function

$$L(\theta) = MSE(X, \theta) = \frac{1}{N} \|\hat{Y} - Y\|_F^2$$

where $\|X\|_F^2 = Tr(XX^T)$ is the Frobenius norm of a matrix which is the square root of the sum of the absolute squares of its elements.

- Its derivative with respect to the weights θ is:

$$\frac{\partial L(\theta)}{\partial \theta} = \frac{\partial L(\theta)}{\partial W} = \frac{1}{N} \frac{\partial}{\partial W} (Tr \{ (\hat{Y} - Y)^T (\hat{Y} - Y) \})$$

- The loss is minimum when $\frac{\partial L(\theta)}{\partial \theta} = 0$. That is:

$$W = (X^T X)^{-1} X^T Y$$

- The above is the **Normal Form** which gives a direct solution for the weight vector θ
- However, this method can be inefficient
- **Gradient Descent** is an iterative method that can find optimal θ more efficiently.

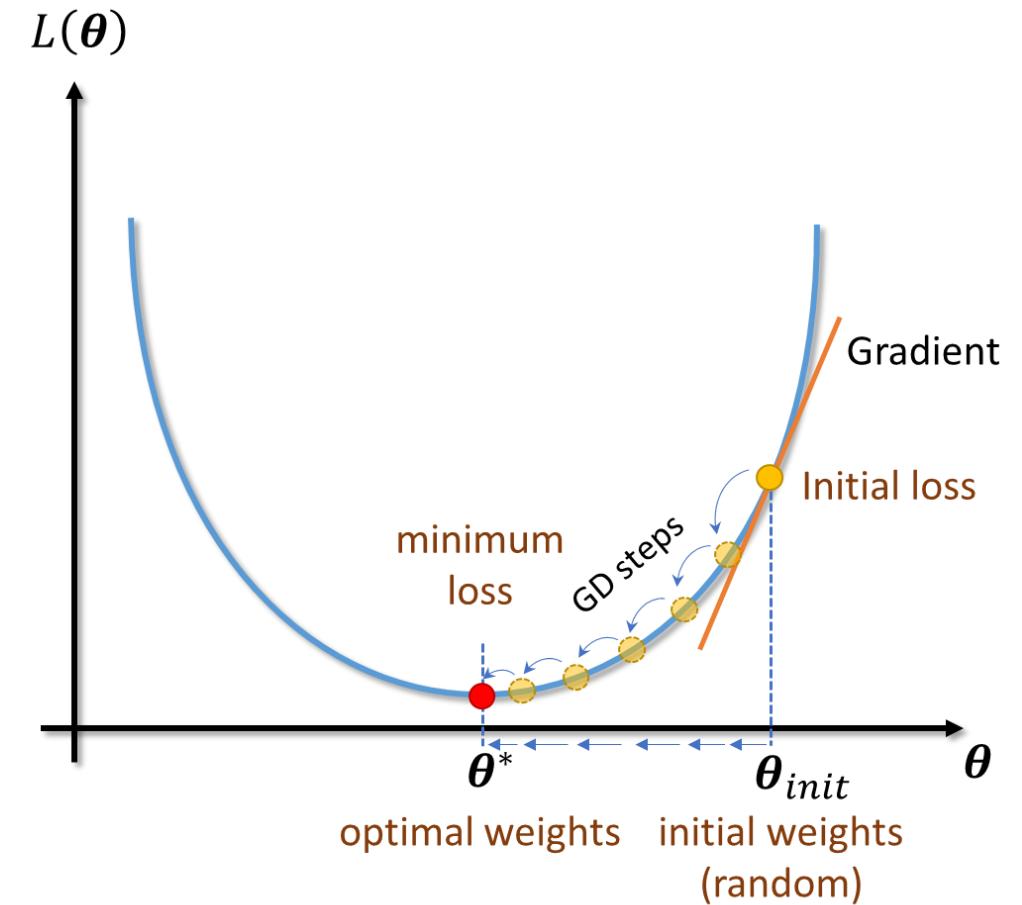
Backpropagation

Optimum Weights for a Single Neuron

- An iterative process where the neuron learns iteratively from training samples to find the optimum weights at which the loss function reaches its minimum.
- The basic idea is to compute the gradient (derivative) of the loss function in a sequence of epochs (intake of the dataset) and update the weights **in proportion to the gradient** according to the update rule:

$$\begin{aligned} \mathbf{W}(t+1) &= \mathbf{W}(t) - \alpha \frac{\partial L(\theta)}{\partial \mathbf{W}} \\ &= \mathbf{W}(t) - \alpha \frac{2}{N} \mathbf{X}^T (\hat{\mathbf{Y}} - \mathbf{Y}) \end{aligned}$$

where $t + 1$ and t indicate the new and current epochs, respectively, and α is a predetermined learning rate usually smaller than 0.5



Backpropagation

Optimum Weights for a Single Neuron

- Accounting for non-zero bias,

$$\widehat{\mathbf{Y}} = \mathbf{X}\mathbf{W} + \mathbf{o}\mathbf{b}^T$$

where $\mathbf{o} \in \mathbb{R}^{N \times 1} = \mathbf{1}_N$ is a vector of ones

- Again, using MSE as the loss function

$$L(\boldsymbol{\theta}) = MSE(\mathbf{X}, \boldsymbol{\theta}) = \frac{1}{N} \left\| \widehat{\mathbf{Y}} - \mathbf{Y} \right\|_F^2$$

- Its derivative with respect to the bias \mathbf{b} is:

$$\frac{\partial L(\boldsymbol{\theta})}{\partial \mathbf{b}} = \frac{\partial L(\boldsymbol{\theta})}{\partial \mathbf{b}} = \frac{1}{N} \frac{\partial}{\partial \mathbf{b}} (Tr \left\{ (\widehat{\mathbf{Y}} - \mathbf{Y})(\widehat{\mathbf{Y}} - \mathbf{Y})^T \right\}) = \frac{2}{N} \mathbf{X}^T (\widehat{\mathbf{Y}} - \mathbf{Y}) \mathbf{o}$$

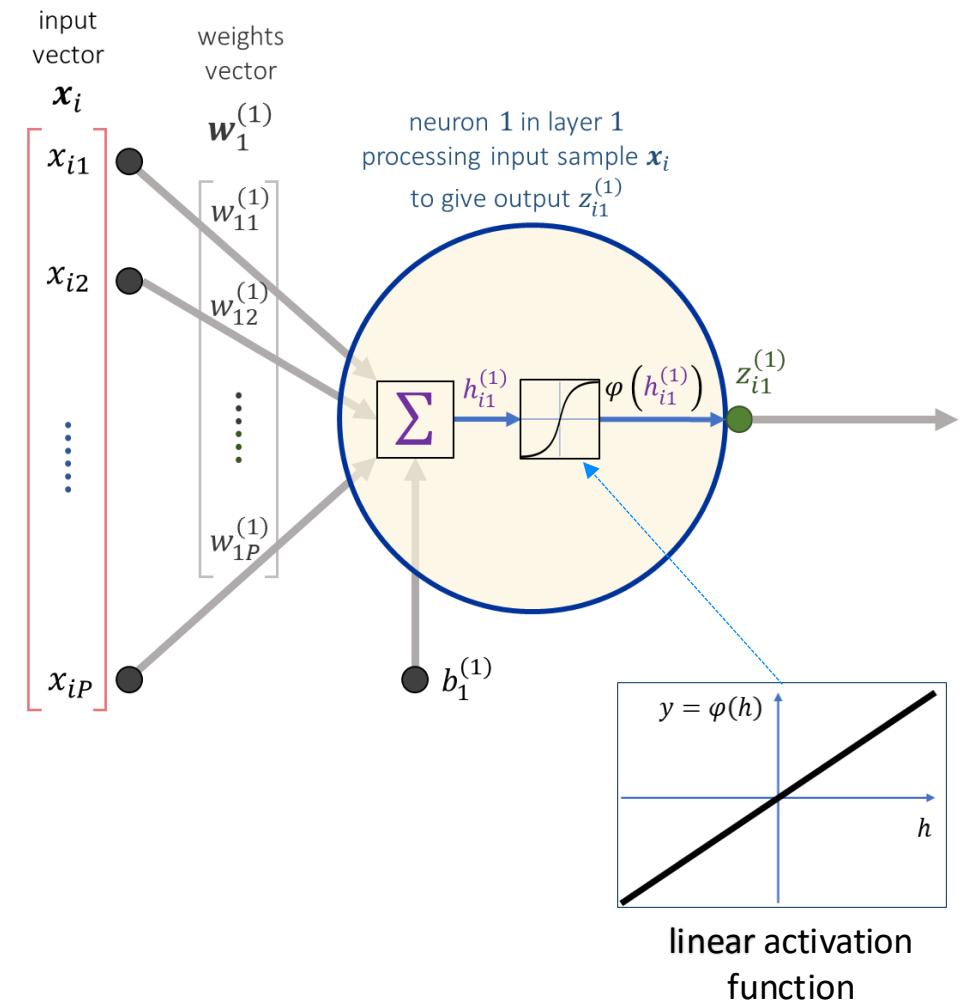
- Thus, the update rule for the Bias is:

$$\mathbf{b}(t+1) = \mathbf{b}(t) - \alpha \frac{\partial L(\boldsymbol{\theta})}{\partial \mathbf{b}} = \boxed{\mathbf{b}(t) - \alpha \frac{2}{N} (\widehat{\mathbf{Y}}^T - \mathbf{Y}^T) \mathbf{o}}$$

Backpropagation

Optimum Weights for a Single Neuron (GD)

- Steps:
 1. Initialize the weights and the bias
 2. While no convergence condition is met
 - i. Calculate the predicted output matrix: $\hat{Y} = \mathbf{XW} + \mathbf{ob}^T$
 - ii. Update weights: $\mathbf{W}(t + 1) = \mathbf{W}(t) - \alpha \frac{2}{N} (\hat{Y}^T - \mathbf{Y}^T) \mathbf{X}$
 - iii. Update bias: $\mathbf{b}(t + 1) = \mathbf{b}(t) - \alpha \frac{2}{N} (\hat{Y}^T - \mathbf{Y}^T) \mathbf{o}$
 - Convergence:
 - The epoch loss $\frac{1}{N} \|\hat{Y} - \mathbf{Y}\|_F^2$ is less than a threshold γ , or
 - A predetermined number of epochs have been completed.
 - Notes:
 - bias \mathbf{b} is updated similarly to the weights \mathbf{W}
 - This is *Batch GD*; two more variants are also used.



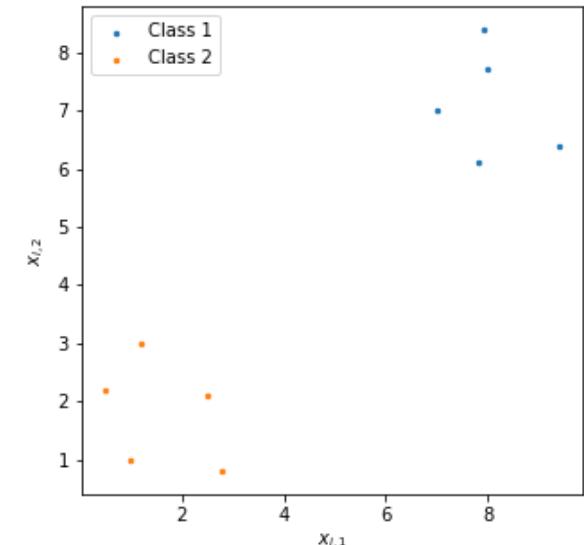
Backpropagation

Optimum Weights for a Single Neuron (GD)

Setup:

- Input: $X = \begin{bmatrix} x_{1,1} & x_{1,2} \\ x_{2,1} & x_{2,2} \\ \vdots & \vdots \\ x_{10,1} & x_{10,2} \end{bmatrix}$, $Y = \begin{bmatrix} y_{1,1} \\ y_{2,1} \\ \vdots \\ y_{10,1} \end{bmatrix}$
- Weights: $\mathbf{W}^{(1)} = [w_{11}^{(1)} \quad w_{12}^{(1)}]$
- Bias: $\mathbf{b}_1^{(1)} = [b_1^{(k)}]$
- Output: $\hat{Y} = X(\mathbf{W}^{(1)}) + o(\mathbf{b}_1^{(1)})^T = \begin{bmatrix} \hat{y}_{1,1} \\ \hat{y}_{2,1} \\ \vdots \\ \hat{y}_{10,1} \end{bmatrix}$
- MSE loss: $\frac{1}{N} \|\hat{Y} - Y\|_F^2$

x_{i1}	x_{i2}	y_i
1.0	1.0	1
9.4	6.4	0
2.5	2.1	1
8.0	7.7	0
0.5	2.2	1
7.9	8.4	0
7.0	7.0	0
2.8	0.8	1
1.2	3.0	1
7.8	6.1	0



A dataset for perceptron classification

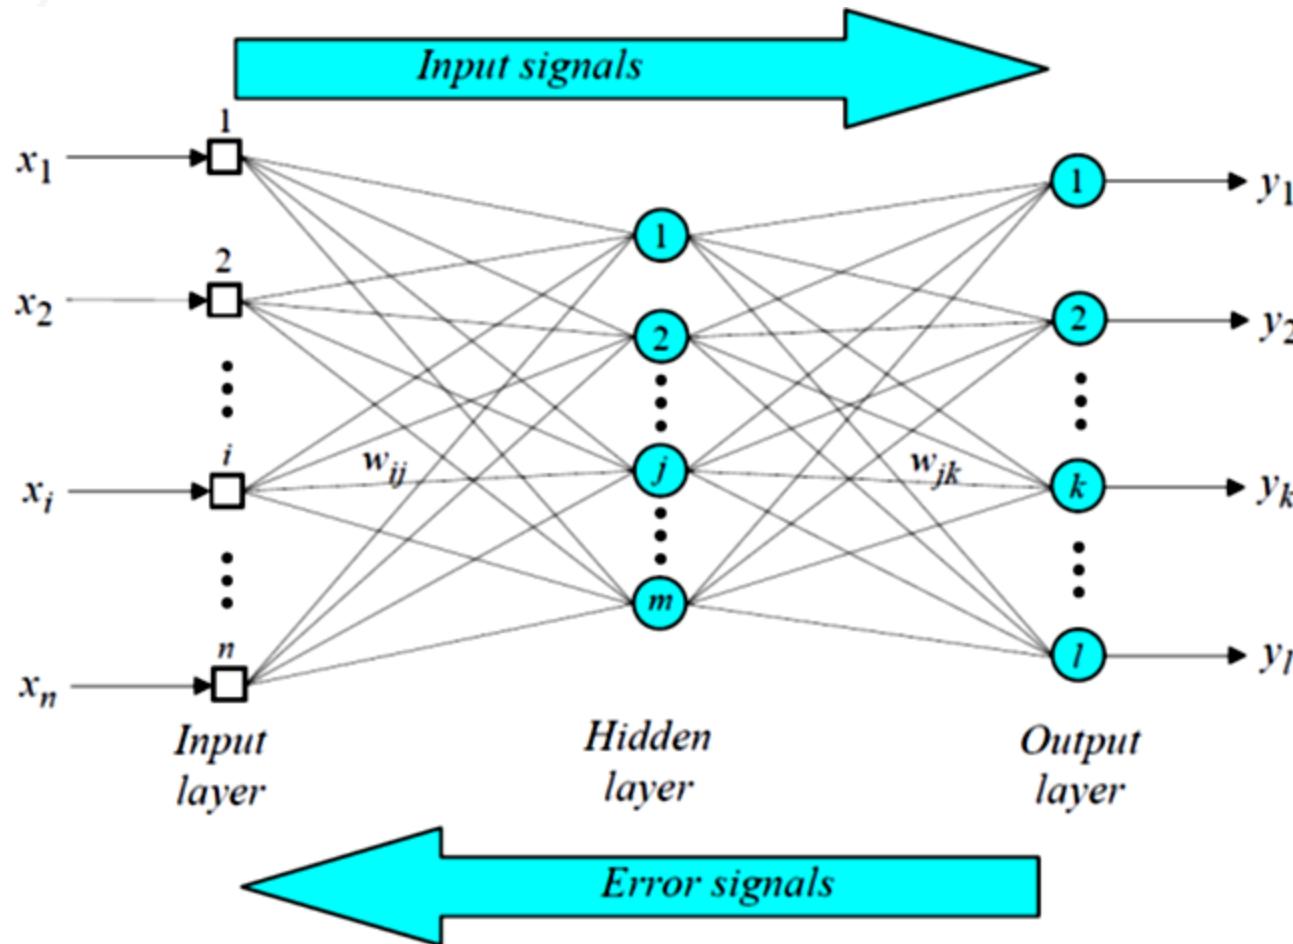
Gradient Descent:

Updating Weights: $\frac{\partial L(\theta)}{\partial \mathbf{W}^{(1)}} = \frac{2}{N} (\hat{Y}^T - Y^T) X$

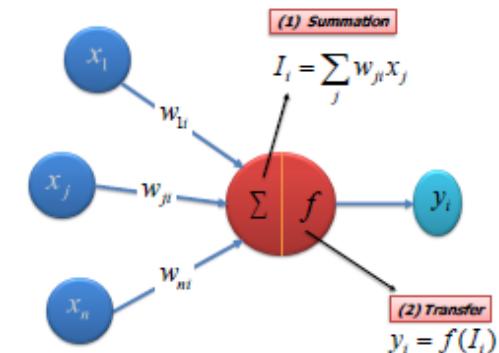
Updating Bias: $\frac{\partial L(\theta)}{\partial \mathbf{b}_1^{(1)}} = \frac{2}{N} (\hat{Y}^T - Y^T) o$

Backpropagation

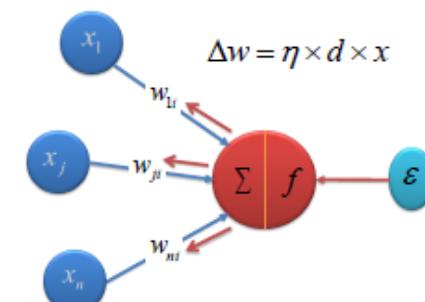
Terminologies



Feedforward Input Data

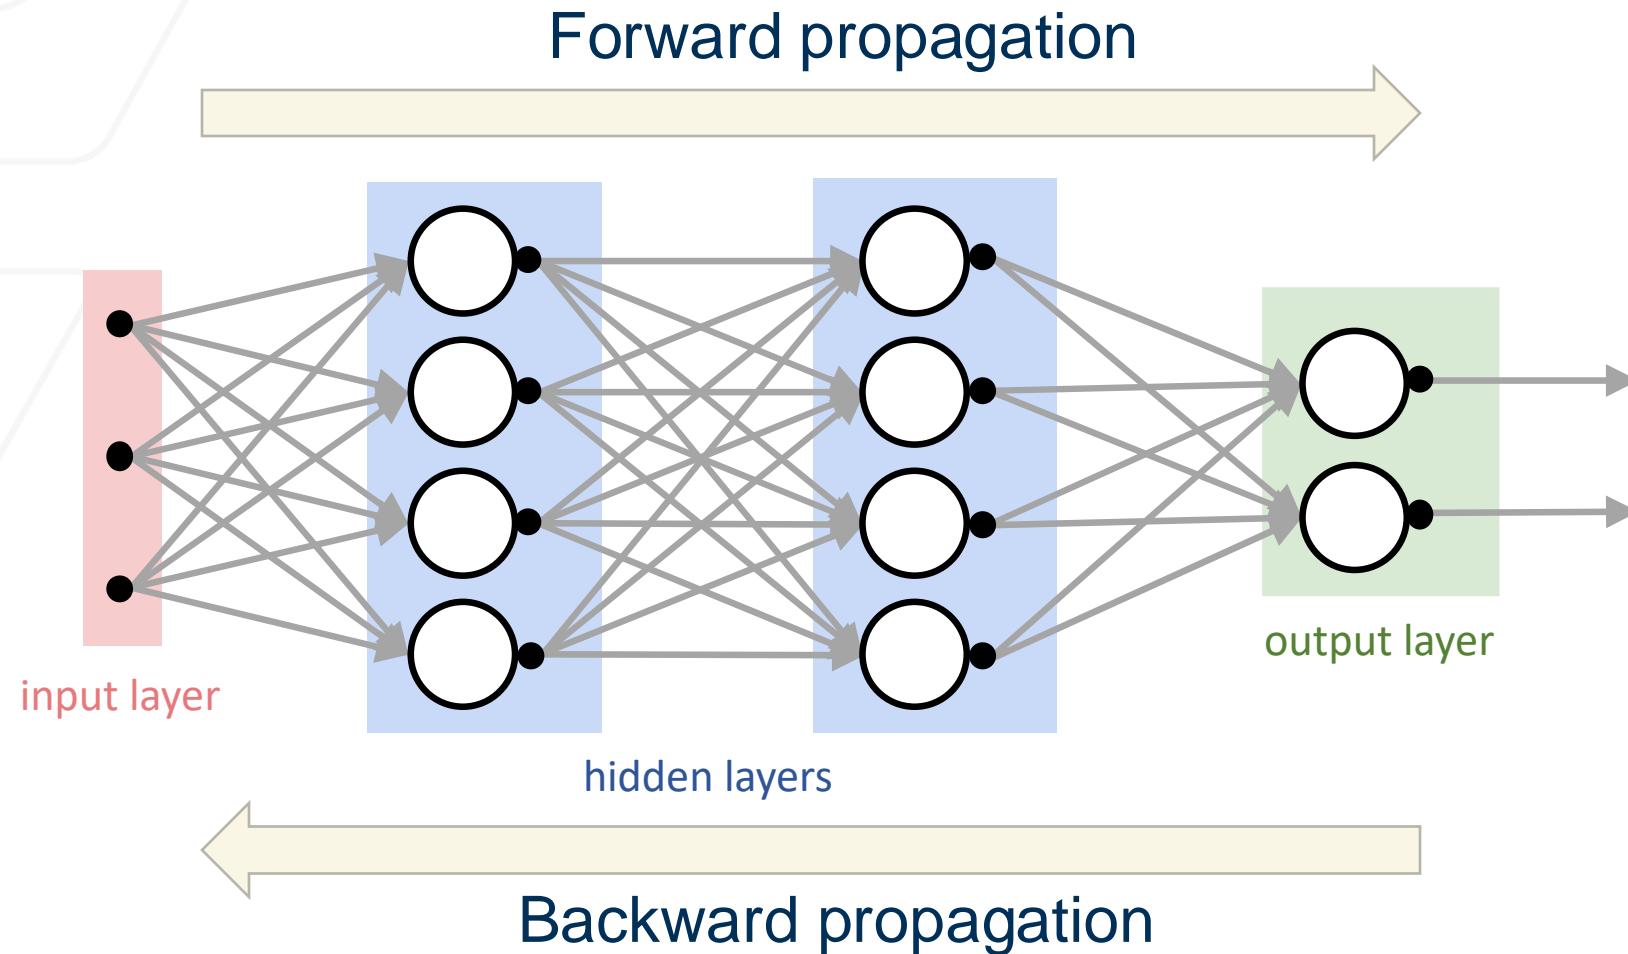


Backward Error Propagation

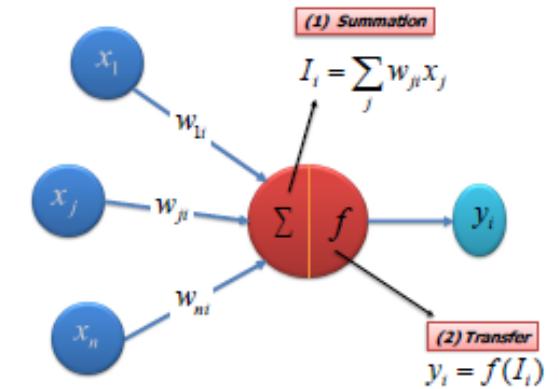


Backpropagation

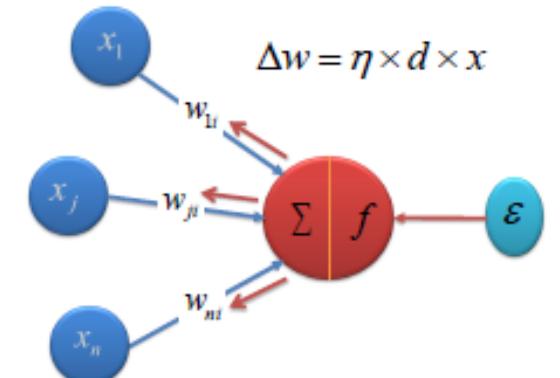
Terminologies



Feedforward Input Data



Backward Error Propagation



Neural Networks

Quick Review of Derivatives

Derivative of a function, $f(\theta)$:

- how fast f changes around θ
- will f increase or decrease if we increase θ
- is θ higher or lower than it should be (θ^*); shall we make it bigger or smaller to be where it should be

$$\frac{\partial}{\partial \theta} f(g(\theta)) = \frac{\partial}{\partial g} f(g(\theta)) \cdot \frac{\partial}{\partial \theta} g(\theta) \text{ chain rule}$$

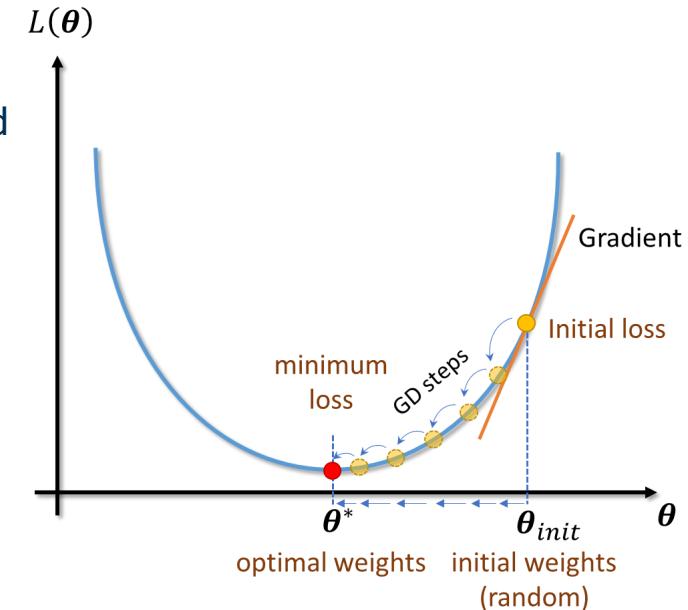
$$\frac{\partial}{\partial \theta} \sum_i f_i(\theta) = \sum_i \frac{\partial}{\partial \theta} f_i(\theta) \text{ derivative of the sum is the sum of the derivatives}$$

$$\frac{\partial}{\partial \theta_k} \sum_i a_i f(\theta_i) = a_k \frac{\partial}{\partial \theta_k} f(\theta_k) \text{ derivative wrt one element of the sum collapses the sum}$$

The sigmoid has a special property:

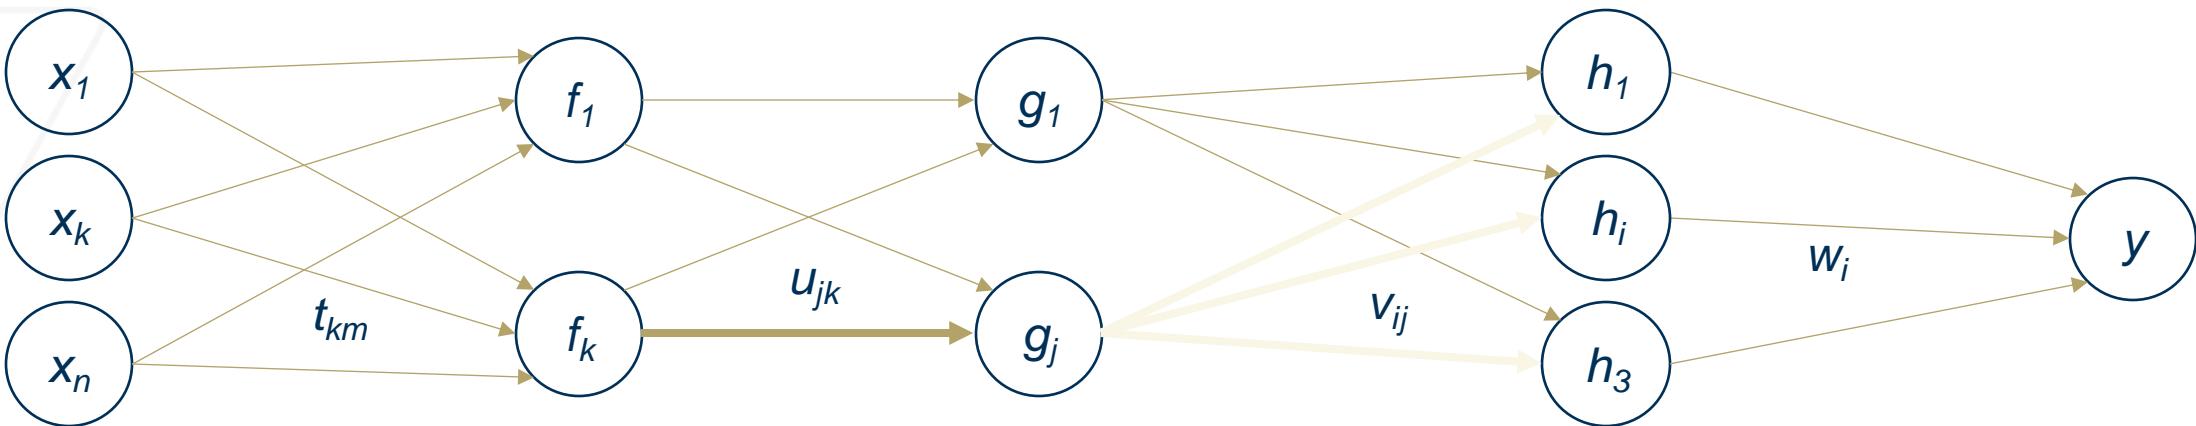
$$\sigma'(x) = \frac{\partial}{\partial x} (1 + e^{-x})^{-1} = -(1 + e^{-x})^{-2}(-e^{-x}) = \frac{1}{1 + e^{-x}} \frac{e^{-x}}{1 + e^{-x}} = \sigma(x)(1 - \sigma(x))$$

$$\frac{\partial}{\partial x} \sigma(f_x) = \sigma(f_x)(1 - \sigma(f_x)) \frac{\partial}{\partial x} f_x$$



Neural Networks

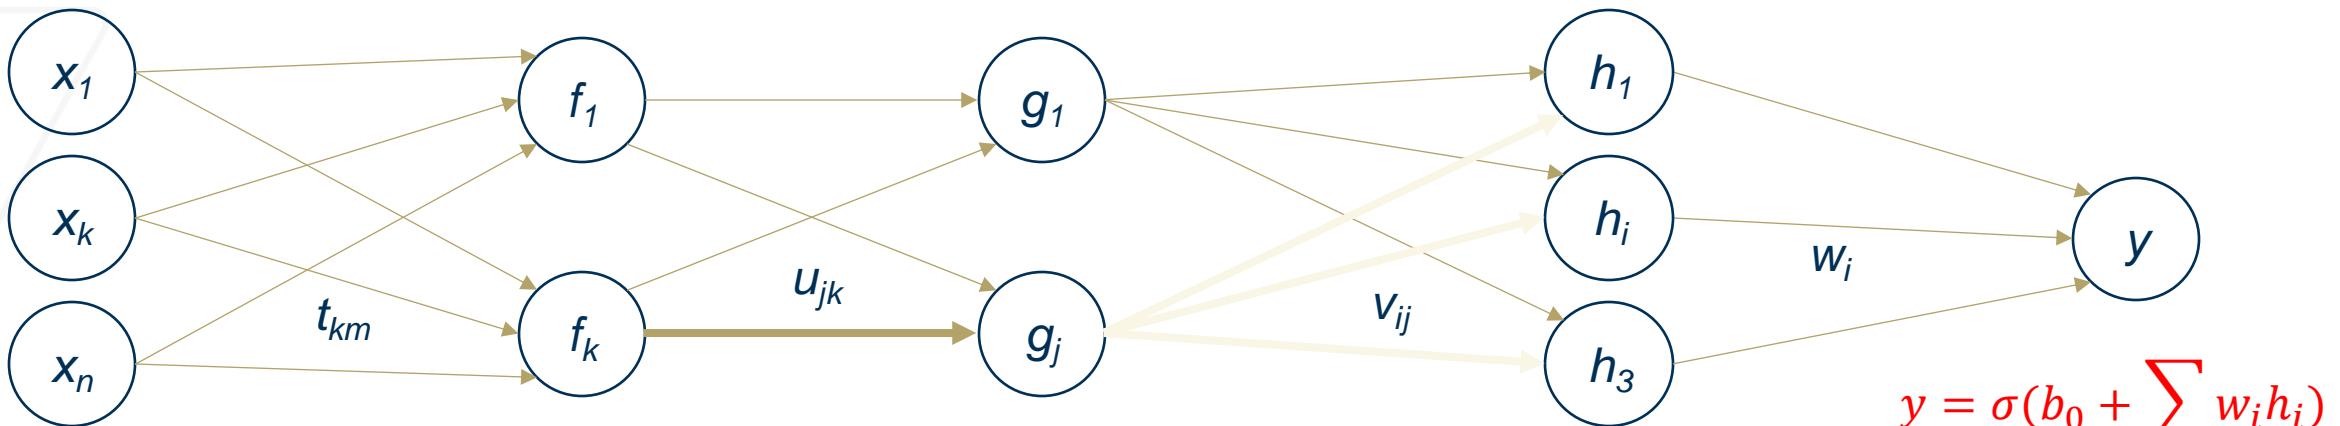
Backpropagation



1. For a new sample $x = [x_1, \dots x_n]$
2. Feed forward: compute g_j based on units f_k from the previous layer using $g_j = \sigma(b_{j0} + \sum_k u_{jk} f_k)$
3. Predict y as y^*
4. Backpropagate error: $e = y - y^*$
 1. Determine whether we need g_j to be lower or higher (assuming we already determined the better values for h from a previous step) by $\frac{\partial e}{\partial g_j} = \sum_i \sigma'(h_i) v_{ij} \frac{\partial e}{\partial h_i}$, or in other words by determining how h_i will change as g_j changes AND by inspecting if h_i was too high or too low
 2. Update weights: update the weight u_{jk} that feeds g_j by $\frac{\partial e}{\partial u_{jk}} = \frac{\partial e}{\partial g_j} \sigma'(g_j) f_k$, or in other words by determining if we want g_j to be higher or lower AND by determining how g_j will change if u_{jk} becomes higher or lower; now update the weight with a learning rate multiplier of $\frac{\partial e}{\partial u_{jk}}$

Neural Networks

Backpropagation

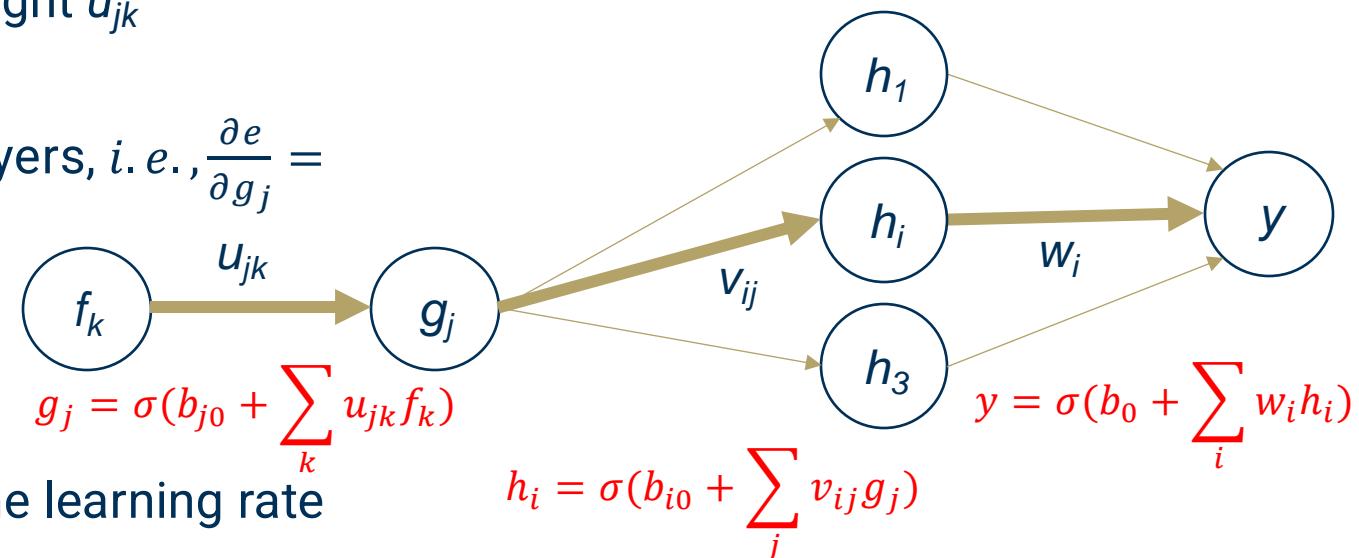


- $e = \frac{1}{2}(y - y^*)^2$ $f_k = \sigma(b_{k0} + \sum_m t_{km}x_m)$ $g_j = \sigma(b_{j0} + \sum_k u_{jk}f_k)$ $h_i = \sigma(b_{i0} + \sum_j v_{ij}g_j)$
- $\frac{\partial e}{\partial h_i} = (y - y^*) \frac{\partial y}{\partial h_i} = (y - y^*) y(1 - y) w_i$; the sum goes away because we are differentiating wrt to one of the elements
- $\frac{\partial e}{\partial g_j} = (y - y^*) \frac{\partial y}{\partial g_j} = (y - y^*) y(1 - y) \sum_i w_i \frac{\partial h_i}{\partial g_j} = (y - y^*) y(1 - y) \sum_i w_i h_i (1 - h_i) v_{ij}$
- By observing the nesting pattern, we can write $\frac{\partial e}{\partial g_j} = \sum_i h_i (1 - h_i) v_{ij} \frac{\partial e}{\partial h_i}$ and so on ; compare this equation with the one in 4.1 on the previous slide
- $\frac{\partial e}{\partial u} = (y - y^*) \frac{\partial y}{\partial u} = (y - y^*) y(1 - y) \sum_i w_i h_i (1 - h_i) \sum_j \frac{\partial g_j}{\partial u} = (y - y^*) y(1 - y) \sum_i w_i h_i (1 - h_i) v_{ij} g_j (1 - g_j) f_k = g_j (1 - g_j) f_k \frac{\partial e}{\partial g_j}$

Neural Networks

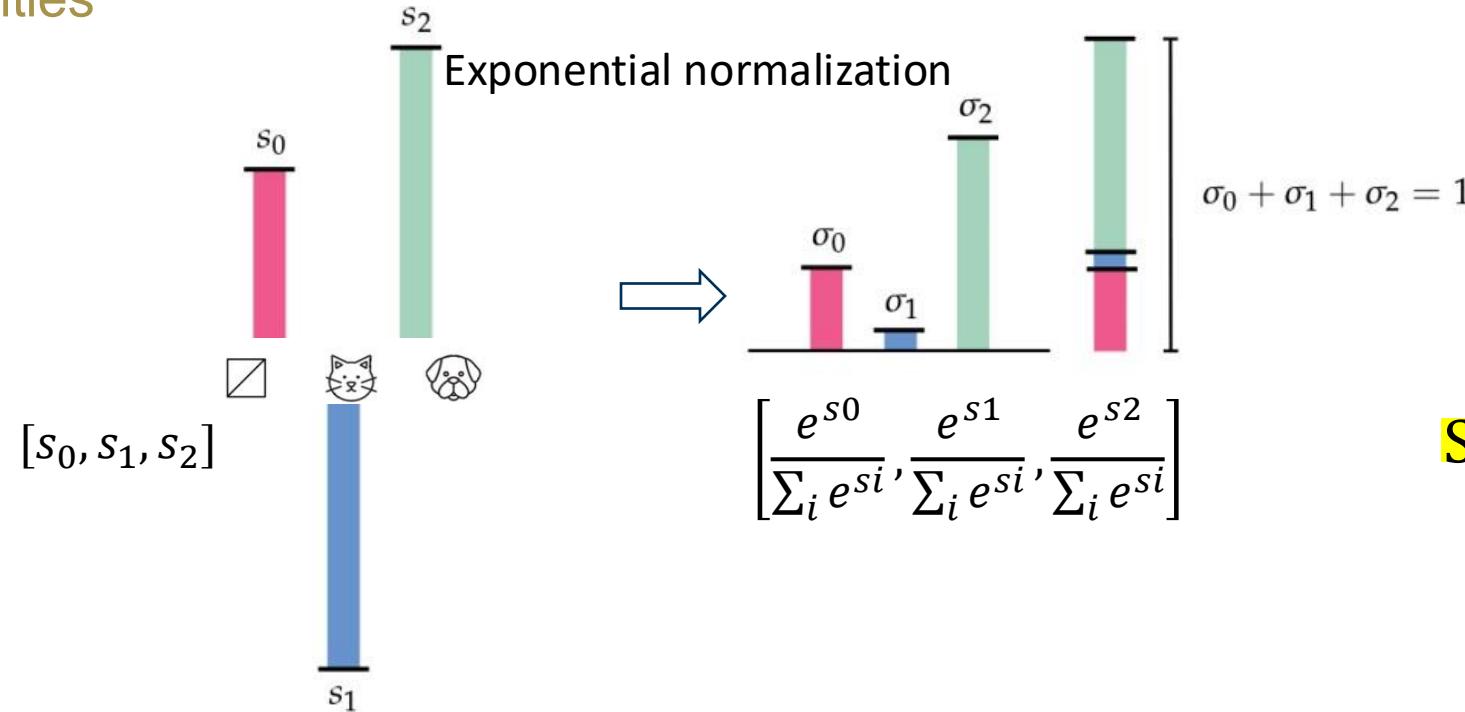
Backpropagation Summary

- **Feed forward propagate** the sample x , compute activation outputs of neurons from previous layers and compute the predicted output y
- Compute the **backpropagation error**: $e = L(y, y^*)$, where L is the loss function
- Assume using sigmoid activation σ , for a weight u_{jk} compute $\frac{\partial e}{\partial u_{jk}}$ using **chain rule**:
 - Compute $\frac{\partial e}{\partial g_j}$ by carrying $\frac{\partial e}{\partial h_i}$ from later layers, i.e., $\frac{\partial e}{\partial g_j} = \sum_i h_i(1 - h_i)v_{ij}\frac{\partial e}{\partial h_i}$
 - Compute $\frac{\partial e}{\partial u_{jk}} = g_j(1 - g_j)f_k \frac{\partial e}{\partial g_j}$
- Update weights using **gradient descent**
 - $u_{jk}(t + 1) = u_{jk}(t) - \alpha \frac{\partial e}{\partial u_{jk}}$, where α is the learning rate
- Iterate the above steps until the error e converges



Neural Networks

Softmax Probabilities



- Suppose the raw outputs of a 3-class ANN is $s = [s_0, s_1, s_2]$, where $s_1 < 0$. The raw outputs can not be interpreted as probability distribution.
- Exponentiating each element in s gives nonnegative values $[e^{s_0}, e^{s_1}, e^{s_2}]$, maintaining the ordering of each element. If we normalize the values by sum, we obtain the **discrete probability distributions**

$$\left[\frac{e^{s_0}}{\sum_i e^{s_i}}, \frac{e^{s_1}}{\sum_i e^{s_i}}, \frac{e^{s_2}}{\sum_i e^{s_i}} \right]$$

of the model outputs. All entries now sum to 1

- This exponential normalization is referred to as SoftMax activation

Neural Networks

Categorical Labels

- Suppose that we have numerical label values $y \in \{0, 1, \dots, C - 1\}$. The One-hot encoded vectors for y takes the form:

$$y_p = 0 \leftarrow \mathbf{y}_p = [1, 0, \dots, 0, 0]$$

$$y_p = 1 \leftarrow \mathbf{y}_p = [0, 1, \dots, 0, 0]$$

⋮

$$y_p = C - 1 \leftarrow \mathbf{y}_p = [0, 0, \dots, 0, 1].$$

- Each one-hot encoded vector is a *categorical label* with length C and contains all zeros except a 1 in the index equal to the true class (from 0 to $C-1$)

Neural Networks

Cross-entropy with Softmax

- Suppose for sample x_i with class p , a multi-class classifier predicts **discrete probability distribution** $f(s) = [f(s_0), \dots, f(s_{c-1})]$ via exponential normalization.
- The desired prediction will be the one-hot vector $y_i = [0, \dots 1, \dots 0]$
- Recall the cost function we used for Logistic Regression:
 - $-\log(\sigma(x))$, if $y = 1$
 - $-\log(1 - \sigma(x))$, if $y = 0$
- Here the log cost can be written as $L = -y_i^T \log(f(s))$

p-th entry

becomes binary cross entropy
when $y_i \in \{[0,1], [1,0]\}$

Note that y_i is a one-hot vector, where $y_{ip} = 1$ and all other entries 0. The form simplifies:

$$L = -\log(f(s_p))$$

- Plug in $f(s_j) = \frac{e^{s_j}}{\sum_j e^{s_j}}$, the form is precisely $L = -\log\left(\frac{e^{s_j}}{\sum_j e^{s_j}}\right)$

We then form a cost function by averaging the entire N data samples:

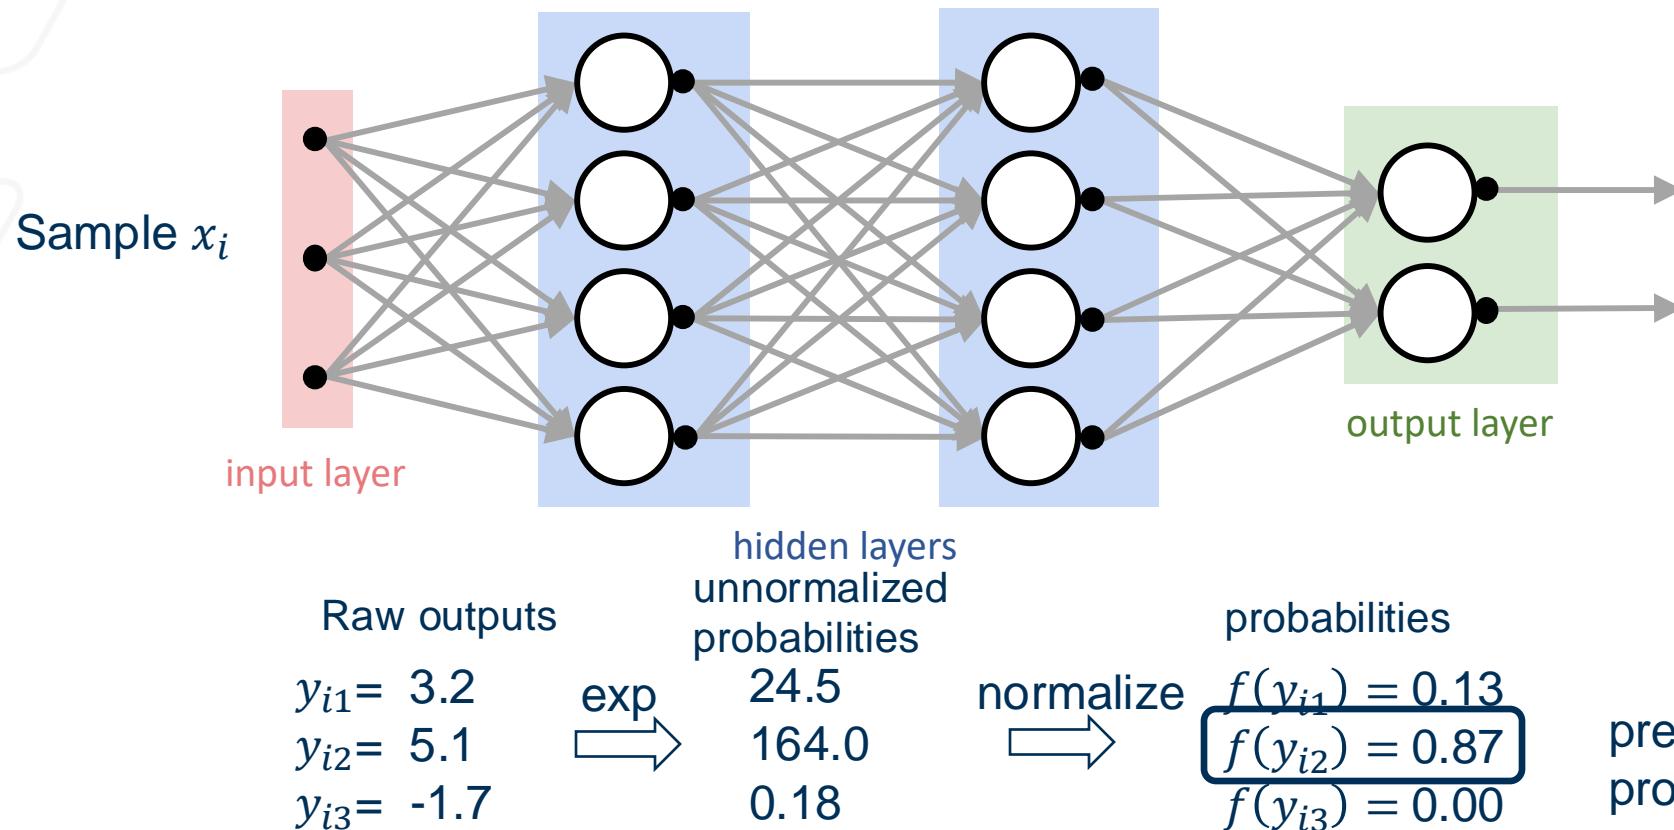
$$L = -\frac{1}{N} \sum_i \log\left(\frac{e^{s_{ip}}}{\sum_j e^{s_{ij}}}\right)$$

Cross-entropy with SoftMax

Review: Artificial Neural Networks

Backpropagation Summary

During the inference, predict the class with the highest probability with SoftMax activation: $\text{argmax}_j f(y_{ij})$



Assume the raw outputs $y_i = [y_{i1}, y_{i2}, y_{i3}] = [3.2, 5.1, -1.7]$

SoftMax activation:

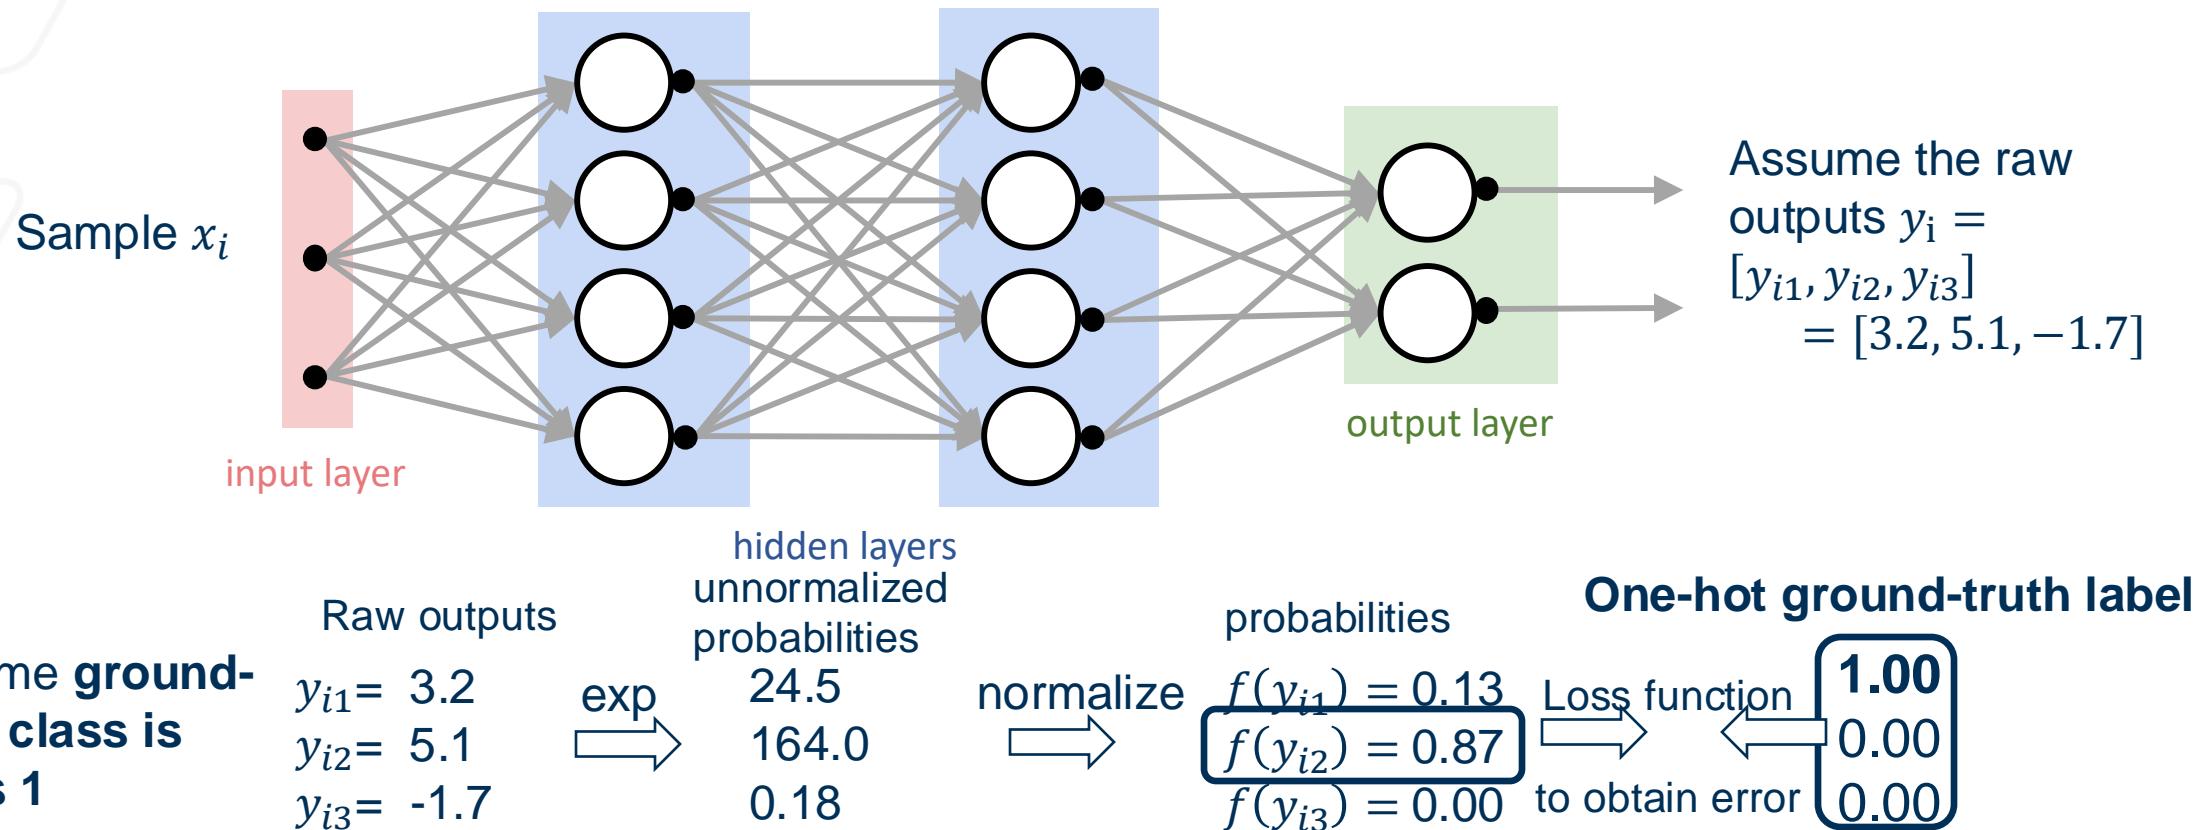
$$f(y_{ij}) = \frac{e^{y_{ij}}}{\sum_k e^{y_{ik}}}$$

predict class 2 with highest probability

Review: Artificial Neural Networks

Backpropagation Summary

During training, for every sample, we set the ground-truth label as a one-hot vector $[1, 0, \dots, 0]^T$ with 1 for the correct class and 0 for every other class. Backpropagate the error and repeat for every sample.



Assume ground-truth class is class 1

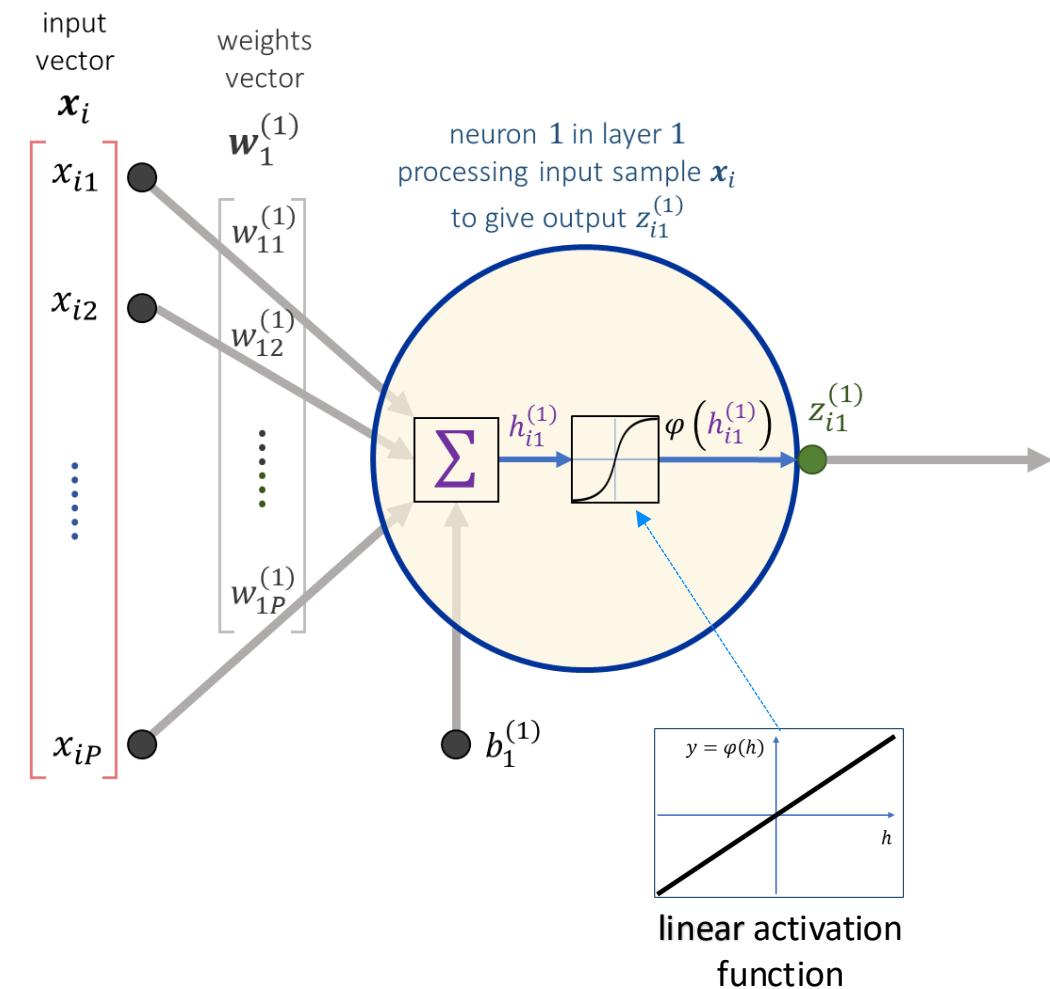
Backpropagation

Linear Classifier

- Given the simple case of a NN having:
 - a single layer with
 - a single neuron with
 - a linear activation function.
- The weight matrix is the **single-column matrix** $\mathbf{W} \in \mathbb{R}^{P \times 1}$
- $\mathbf{W}^{(1)} = [\mathbf{w}_1^{(1)}] = \mathbf{w}_1^{(1)}$
- The bias $[\mathbf{b}_1^{(1)}] = b_1^{(1)}$
- The output $\mathbf{z}^{(1)} = [(\mathbf{z}_1^{(1)})^T] = z_1^{(1)} = \varphi(h_{i1}^{(1)}) = h_{i1}^{(1)}$
- Given an **input matrix** $\mathbf{X} \in \mathbb{R}^{N \times P}$ as defined earlier, the neuron's output matrix $\hat{\mathbf{Y}} \in \mathbb{R}^{N \times 1}$ is:

$$\hat{\mathbf{Y}} = \varphi(\mathbf{X}(\mathbf{W}^{(1)})^T + \mathbf{o}(\mathbf{b}^{(1)})^T)$$

where $\mathbf{o} \in \mathbb{R}^{N \times 1}$ $\mathbf{X}(\mathbf{W}^{(1)})^T + \mathbf{o}(\mathbf{b}^{(1)})^T$ is a vector of ones $\mathbf{w}_1^{(1)} + \mathbf{b}_1^{(1)}$



Backpropagation

Single Layer MLPs

1. System setup:

$$\hat{Y} = \sigma\left(X(\mathbf{W}^{(1)})^T + \mathbf{o}(\mathbf{b}^{(1)})^T\right)$$

For a single sample $x_i \in \mathbb{R}^{p \times 1}$, the output of the ANN is given by:

$$\begin{aligned}\hat{y}_i &= \sigma(h_i) = \frac{1}{1 + e^{-h_i}} \\ h_i &= \mathbf{w}^T \mathbf{x}_i + b = \mathbf{w} \mathbf{x}_i + b\end{aligned}$$

where, $h_i, \hat{y}_i \in \mathbb{R}$, $\mathbf{w}^{(1)} \in \mathbb{R}^{1 \times p}$ is the weights vector, and $b^{(1)} \in \mathbb{R}$ is the bias term

2. Derivation of the loss

Recall binary cross-entropy loss:

$$L(y_i, \hat{y}_i) = -\frac{1}{N} \sum_{i=1}^N y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)$$

Backpropagation

Single Layer MLPs

2.1. derivative of the loss w.r.t. \mathbf{w} :

$$\begin{aligned}\frac{\partial L}{\partial \mathbf{w}} &= -\frac{1}{N} \sum_{i=1}^N y_i \left(\frac{\partial}{\partial \mathbf{w}} \log \hat{y}_i \right) + (1 - y_i) \left(\frac{\partial}{\partial \mathbf{w}} \log(1 - \hat{y}_i) \right) \\ &= -\frac{1}{N} \sum_{i=1}^N y_i \left(\frac{1}{\hat{y}_i} \frac{\partial \hat{y}_i}{\partial \mathbf{w}} \right) + (1 - y_i) \left(\frac{-1}{1 - \hat{y}_i} \frac{\partial \hat{y}_i}{\partial \mathbf{w}} \right)\end{aligned}$$

using $\hat{y}_i = \sigma(h_i)$ and $\frac{\partial \sigma(h)}{\partial h} = \sigma(h)[1 - \sigma(h)]$

$$\begin{aligned}\frac{\partial L}{\partial \mathbf{w}} &= -\frac{1}{N} \sum_{i=1}^N y_i \left(\frac{1}{\hat{y}_i} \hat{y}_i [1 - \hat{y}_i] \frac{\partial h_i}{\partial \mathbf{w}} \right) + (1 - y_i) \left(\frac{-1}{1 - \hat{y}_i} \hat{y}_i [1 - \hat{y}_i] \frac{\partial h_i}{\partial \mathbf{w}} \right) \\ &= -\frac{1}{N} \sum_{i=1}^N y_i \left([1 - \hat{y}_i] \frac{\partial}{\partial \mathbf{w}} (\mathbf{w}_i + b) \right) + (1 - y_i) \left(-\hat{y}_i \frac{\partial}{\partial \mathbf{w}} (\mathbf{w}_i + b) \right)\end{aligned}$$

Therefore $\frac{\partial L}{\partial \mathbf{w}} = -\frac{1}{N} \sum_{i=1}^N y_i ([1 - \hat{y}_i] \mathbf{x}_i^T) + (1 - y_i)(-\hat{y}_i \mathbf{x}_i^T)$

Backpropagation

Single Layer MLPs

2.1. derivative of the loss w.r.t. \mathbf{w} :

$$\begin{aligned}\frac{\partial L}{\partial \mathbf{w}} &= -\frac{1}{N} \sum_{i=1}^N y_i ([1 - \hat{y}_i] \mathbf{x}_i^T) + (1 - y_i)(-\hat{y}_i \mathbf{x}_i^T) \\ &= -\frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i) \mathbf{x}_i^T \\ &= -\frac{1}{N} (\mathbf{Y}^T - \hat{\mathbf{Y}}^T) \mathbf{X} = \frac{1}{N} (\hat{\mathbf{Y}}^T - \mathbf{Y}^T) \mathbf{X}\end{aligned}$$

2.2 derive of the loss w.r.t. b :

Derivation of b is similar to \mathbf{w} , the final derivative:

$$\begin{aligned}\frac{\partial L}{\partial b} &= -\frac{1}{N} \sum_{i=1}^N y_i ([1 - \hat{y}_i]) + (1 - y_i)(-\hat{y}_i) \\ &= -\frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i) \\ &= -\frac{1}{N} (\mathbf{Y}^T - \hat{\mathbf{Y}}^T) \mathbf{o} = \frac{1}{N} (\hat{\mathbf{Y}}^T - \mathbf{Y}^T) \mathbf{o}\end{aligned}$$

Image Classification: mapping the image pixels to probabilities for each category

For simplicity, we take a linear function:

$$\hat{Y} = \phi(XW^T + b^T)$$

where

- $X \in \mathbb{R}^{N \times P}$: dataset containing N vectorized images
- $P \in \mathbb{R}^{H \times W \times C}$: the number of pixels (*features*) of each image
- $\hat{Y} \in \mathbb{R}^{N \times P^{(k)}}$: the associated probabilities of each category $1, 2, \dots, P^{(k)}$
- $\phi(X, W, b)$: the activation function
- $W \in \mathbb{R}^{P \times P^{(k)}}$: the weight matrix
- $b \in \mathbb{R}^{P^{(k)} \times 1}$: the bias vector

$P^{(k)} = C$ (Number of classes) for Classification

Single image binary classification (predicting dog/cat)



stretch pixels into single column vector

$$\phi \left(\begin{bmatrix} 0.2 & 2.1 \\ -0.5 & 0.0 \\ 0.1 & 0.25 \\ 2.0 & 0.2 \\ 1.5 & -0.3 \\ \vdots & \vdots \\ 1.3 & 1.2 \end{bmatrix}^T \begin{bmatrix} 56 \\ 231 \\ 24 \\ 188 \\ 75 \\ \vdots \\ 32 \end{bmatrix} + \begin{bmatrix} 0.2 \\ 2.4 \end{bmatrix} \right) = \begin{bmatrix} 0.8 \\ 0.2 \end{bmatrix}$$

Cat score
Dog score

$w \in \mathbb{R}^{2 \times (32)(32)(3)}$ $x_i \in \mathbb{R}^{(32)(32)(3) \times 1}$

$b \in \mathbb{R}^{2 \times 1}$ $y_i \in \mathbb{R}^{2 \times 1}$

Neural Networks

Training MLP for Image Classification

Example Datasets:

- MNIST ([hand-written digits](#))
 - # total: 70,000 grayscale images
 - # classes: 10
 - # size: 28x28
 - # training samples: 60,000 images
 - # test samples: 10,000 images
- CIFAR-10 (subsets of the [80 million tiny images](#))
 - # total: 60,000 color images
 - # classes: 10
 - # size: 32x32
 - # training samples: 50,000 images
 - # test samples: 10,000 images

0 0 0 0 0 0 0 0 0
1 1 1 1 1 1 1 1 1
2 2 2 2 2 2 2 2 2
3 3 3 3 3 3 3 3 3
4 4 4 4 4 4 4 4 4
5 5 5 5 5 5 5 5 5
6 6 6 6 6 6 6 6 6
7 7 7 7 7 7 7 7 7
8 8 8 8 8 8 8 8 8
9 9 9 9 9 9 9 9 9

https://www.researchgate.net/figure/Sample-images-of-MNIST-data_fig3_222834590

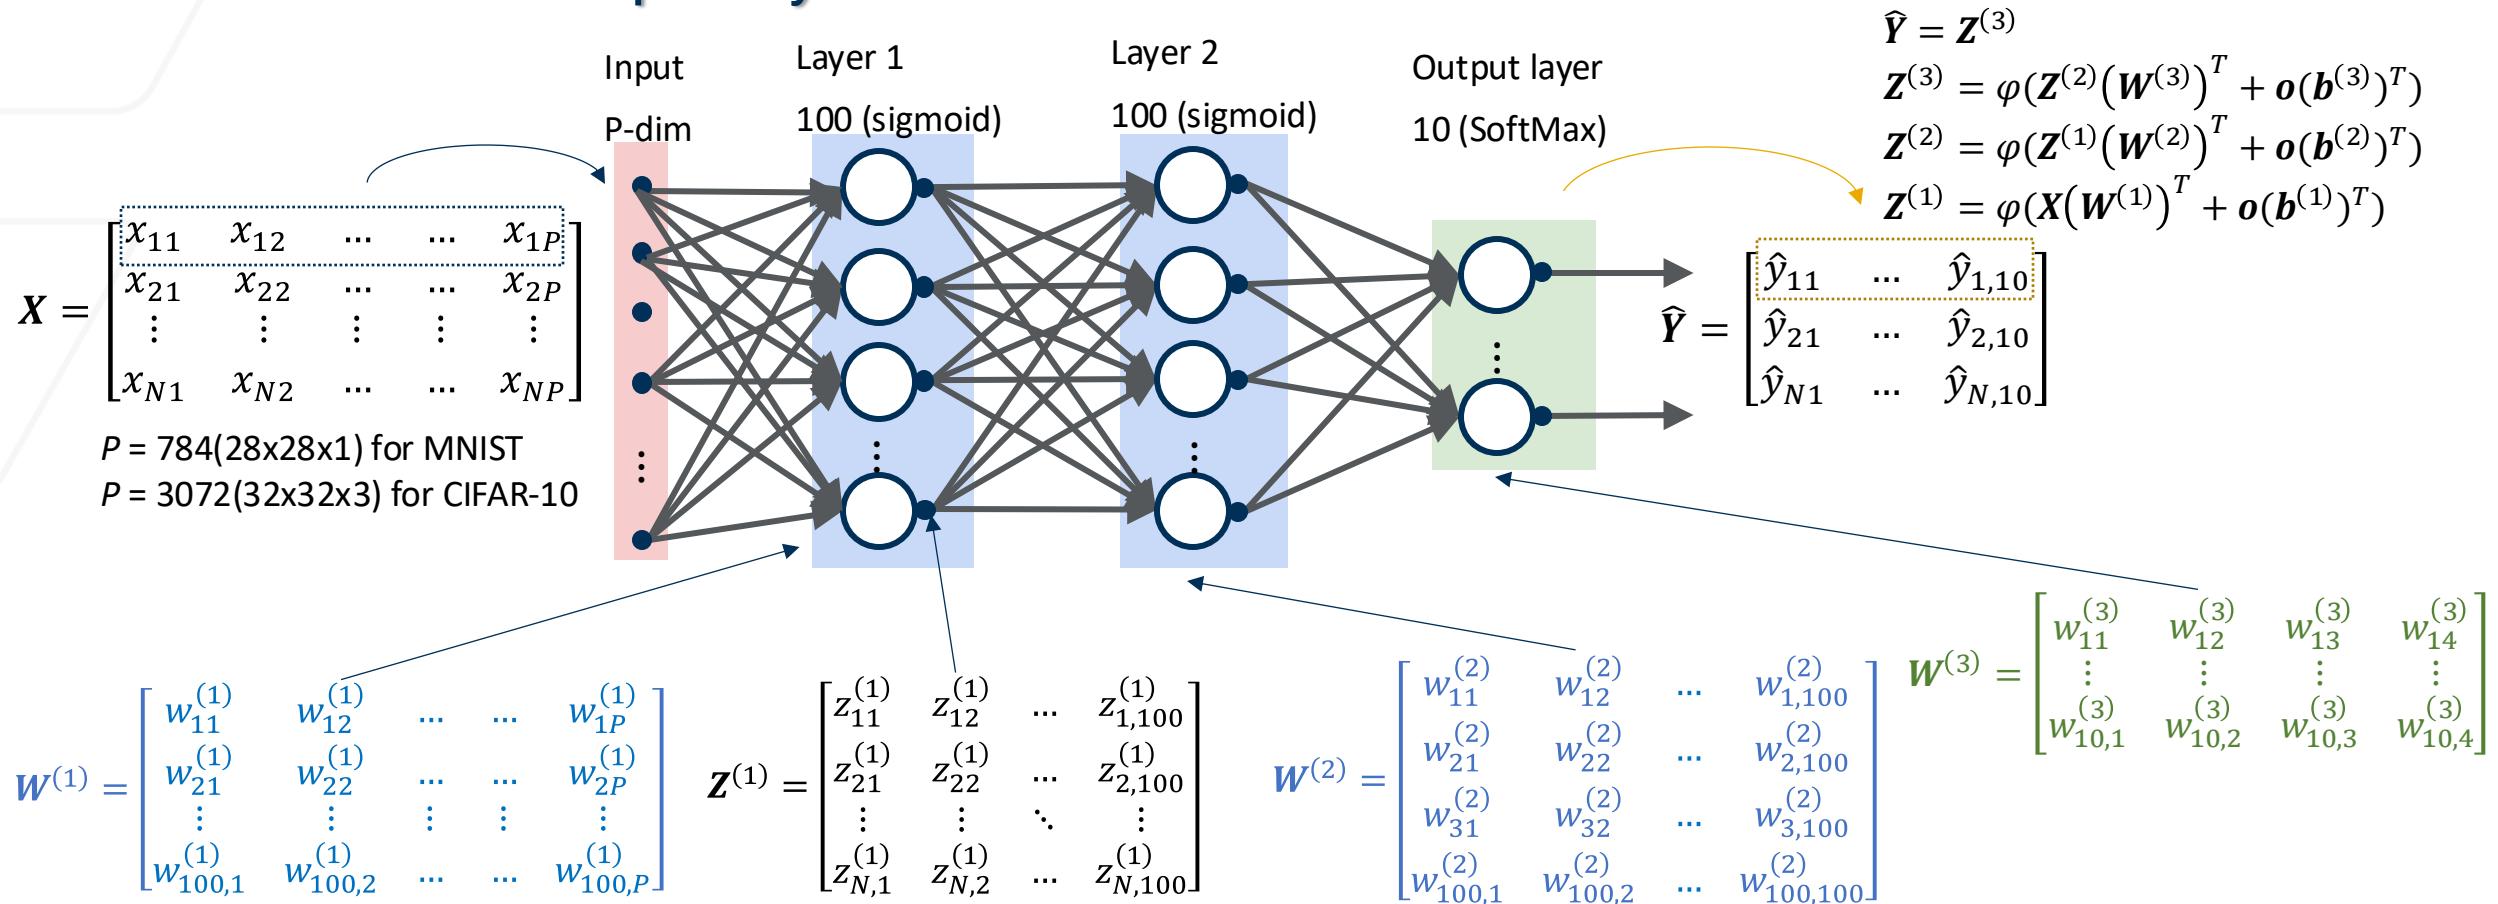


<https://www.kaggle.com/c/cifar-10>

Neural Networks

Training MLP for Image Classification

- P -dimensional input vectors, 2 hidden layers with 100 neurons in each, and 10 neurons in output layer.



Neural Networks

Training MLP for Image Classification

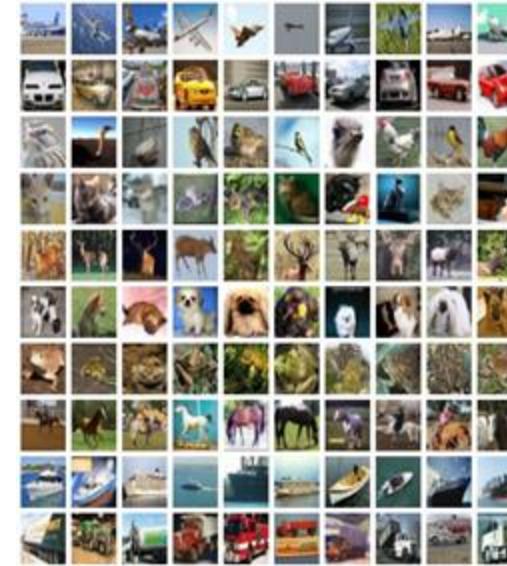


A 4x4 grid of handwritten digits from the MNIST dataset. The digits are black on a white background and are arranged in a 4x4 pattern.

0 0 0 0 0 0 0 0
1 1 1 1 1 1 1 1
2 2 2 2 2 2 2 2
3 3 3 3 3 3 3 3
4 4 4 4 4 4 4 4
5 5 5 5 5 5 5 5
6 6 6 6 6 6 6 6
7 7 7 7 7 7 7 7
8 8 8 8 8 8 8 8
9 9 9 9 9 9 9 9

https://www.researchgate.net/figure/Sample-images-of-MNIST-data_fig3_222834590

MNIST acc: ~72%



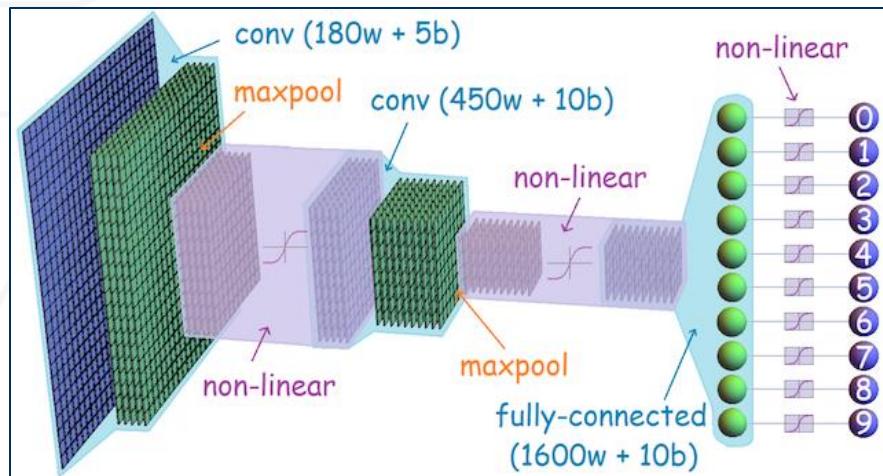
<https://www.kaggle.com/c/cifar-10>

CIFAR-10 acc:~53 %

- MLP tends to recognize better on the images in which objects are well-centered, with simple shape/texture information.
- need more powerful models to address complex image datasets.

Overview

In this Lecture..



Review: Artificial Neural Networks

- Activation Function
- Perceptron Network
- Multi-layer ANN

Backpropagation

- Feedforward and Backward Error Propagation
- Learning Algorithm
- Image Classification using ANNs
- Mini-batch Gradient descent
- Linear Classifier
- Logistic Regression

PyTorch

- Autograd
- Built-in Modules
- Examples

- What is PyTorch?
 - A Python-based scientific computing package for executing dynamic computational graphs over Tensor objects that behave similarly as Numpy arrays.
- Why PyTorch?
 - It comes with a powerful *automatic differentiation* engine that removes the need for manual back-propagation

PyTorch

AutoGrad: Automatic Differentiation

- Given a vector valued function $y = f(x)$, $x \in \mathbb{R}^n$, $y \in \mathbb{R}^m$, the gradient of y w.r.t. x is a Jacobian matrix:

$$J = \begin{bmatrix} \frac{\partial y_1}{\partial x_1} & \dots & \frac{\partial y_1}{\partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial y_m}{\partial x_1} & \dots & \frac{\partial y_m}{\partial x_n} \end{bmatrix}$$

- Generally, `torch.autograd` is an engine for computing *vector-Jacobian* product. That is, given any vector $v = \begin{bmatrix} v_1 \\ v_2 \\ \vdots \\ v_m \end{bmatrix}$, compute the product $v^T J$. If v happens to be the gradient of a scalar function $l =$

$g(y)$, that is, $v = \begin{bmatrix} \frac{\partial l}{\partial y_1} \\ \frac{\partial l}{\partial y_2} \\ \vdots \\ \frac{\partial l}{\partial m} \end{bmatrix}$, then by the chain rule, the *vector-Jacobian* product would be the gradient of l w.r.t. x .

PyTorch

AutoGrad: Toy Example

Given $y = \|x_1 - x_2\|_2^2$

Compare $2(x_1 - x_2)$ and $\frac{\partial y}{\partial x_1}$ using PyTorch *autograd*

```
import torch
x1 = torch.randn(2, 2).requires_grad_(True)
x2 = torch.randn(2, 2).requires_grad_(True)

y = torch.sum((x1-x2)**2)
y.backward()

print(2*(x1-x2))
print(x1.grad)
# ===== output =====
# tensor([[ 0.1411,  3.9669], [-1.2993,  2.1217]])
# tensor([[ 0.1411,  3.9669], [-1.2993,  2.1217]] ,
grad_fn=<MulBackward0>)
```

manually computed
gradient $2(x_1 - x_2)$

$\frac{\partial y}{\partial x_1}$ using PyTorch autograd

PyTorch

Numpy vs PyTorch

```
for i in range(max_iter):  
  
    Y_hat = X@W.T + o@b.T  
  
    # MSE loss  
  
    loss = np.sum((Y_hat-Y)**2)/num_samples  
    #FrobeniusNorm  
  
    W_gradients = (2/num_samples)*(Y_hat.T - Y.T)@X  
    b_gradients = (2/num_samples)*(Y_hat.T - Y.T)@o  
  
    # follow the negative gradients  
  
    W = W - alpha*W_gradients # update the weight  
matrix  
  
    b = b - alpha*b_gradients # update the biases  
matrix  
  
    Y_hat = X@W.T + o@b.T # predicted  
"probabilities"  
  
#W gradients: [ 0.06089246 -0.01254302]  
#b gradients: -0.34077861476403615  
  
Iteration:120 | Loss = 0.1104 | Train Acc = 0.80
```

Given example, we can compare the gradients computed manually with NumPy and PyTorch autograd (see highlighted lines)

```
for i in range(max_iter):  
  
    y_hat = X_tensor@W_tensor.t() + o_tensor@b_tensor.t()  
  
    # MSE loss  
  
    loss = F.mse_loss(y_hat, y_tensor)  
  
    loss.backward()  
  
    W_gradient = W_tensor.grad  
    b_gradient= b_tensor.grad  
  
    # updating the weights  
  
    with torch.no_grad():  
  
        W_tensor.data = W_tensor.data - alpha*W_gradient  
        b_tensor.data = b_tensor.data - alpha*b_gradient  
  
    y_hat = X_tensor@W_tensor.t() + o_tensor@b_tensor.t()  
    # predicted "probabilities"  
  
#W gradients: [ 0.06089246 -0.01254302]  
#b gradients: -0.34077861476403615  
  
Iteration:120 | Loss = 0.1104 | Train Acc = 0.80
```

PyTorch

Numpy vs PyTorch: Sigmoid and BCE Loss

```
for i in range(max_iter):  
  
    H = X@W.T + o@b.T  
  
    Y_hat = 1. / (1. + np.exp(-H))  
  
    # BCE loss  
  
    loss = -1. * np.sum(Y * np.log(Y_hat) + (1-Y) *  
np.log(1-Y_hat)) / num_samples  
  
    W_gradients = (1./num_samples)*(Y_hat.T - Y.T)@X  
    b_gradients = (1./num_samples)*(Y_hat.T - Y.T)@o  
  
    # follow the negative gradients  
  
    W = W - alpha*W_gradients # update the weight  
matrix  
  
    b = b - alpha*b_gradients # update the biases  
matrix  
  
#W gradients: [0.07499208 -0.05085141]  
#b gradients: -0.16729619094308953  
  
Iteration:120 | Loss = 0.3512 | Train Acc = 0.90
```

Given example, we can compare the gradients computed manually with NumPy and PyTorch autograd (see highlighted lines)

```
for i in range(max_iter):  
  
    h_hat = X_tensor@W_tensor.t() + o_tensor@b_tensor.t()  
  
    y_hat = torch.sigmoid(h_hat)  
  
    # BCE loss  
  
    loss = F.binary_cross_entropy(y_hat, y_tensor)  
  
    loss.backward()  
  
    W_gradient = W_tensor.grad  
    b_gradient = b_tensor.grad  
  
    # updating the weights  
  
    with torch.no_grad():  
  
        W_tensor.data = W_tensor.data - alpha*W_gradient  
        b_tensor.data = b_tensor.data - alpha*b_gradient  
  
#W gradients: [0.07499208 -0.05085141]  
#b gradients: -0.16729619095070608  
  
Iteration:120 | Loss = 0.3512 | Train Acc = 0.90
```

Basic Modules in PyTorch:

nn.Module API: defines arbitrary network architectures, while automatically tracking every learnable parameters.

torch.optim package: implements all the common optimizers, such as SGD, RMSProp, Adagrad, and Adam, etc.

PyTorch

Basic Modules

To use the Module API,
follow the steps below:

1. Subclass nn.Module

2. In the constructor
`__init__()` define all the
layers and components
you need as class
attributes

3. In the forward() method,
define the connectivity of
your network

```
# Give your network class an intuitive name like
`SNC(Sigle Neuron Classifier)`.

class SNC(nn.Module):
    def __init__(self):
        super(SNC, self).__init__()
        # define a single neuron
        self.single_neuron= nn.Linear(in_features=2,
out_features=1, bias=True) # the layer will learn an
additive bias
        self.non_linearity = nn.Sigmoid()

    def forward(self, x):
        x = self.single_neuron(x)
        x = self.non_linearity(x)
        return x
```

PyTorch example – single neuron classifier

PyTorch

Implementing Single Neuron Classifier

```
# Single Neuron Nonlinear Classifier

class SNC(nn.Module):

    def __init__(self): # define the layer(s) and network components
        super(SNC, self).__init__()

        self.single_neuron= nn.Linear(in_features=num_features,
out_features=1, bias=True) # the layer will learn an additive bias.

        self.non_linearity = nn.Sigmoid()

    def forward(self, x): # here we define the forward pass for network

        x = self.single_neuron(x)

        x = self.non_linearity(x)

        return x

network = SNC()

optimizer = optim.SGD(network.parameters(), lr=2) #Stochastic gradient
descend

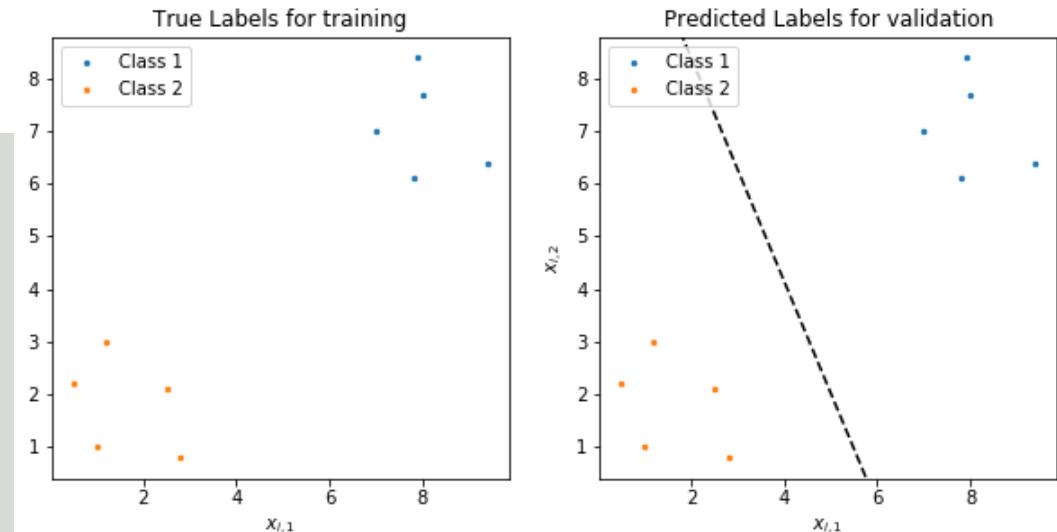
criterion = nn.BCELoss() # using BCE loss for training

for i in range(max_iter):

    network.zero_grad() # clearing all gradients in the network
    y_hat_tensor = network(X_tensor) #forward pass
    loss = criterion(y_hat_tensor, y_tensor)

    loss.backward()
    optimizer.step()
```

example



:

Iteration 20#, training acc = 1.00

Loss: $L(\theta)= 0.0321$

W gradients:

$$\frac{\partial L(\theta)}{\partial W^{(1)}} = [0.0083, 0.0355]^T$$

b gradients:

$$\frac{\partial L(\theta)}{\partial b^{(1)}} = -0.0133$$



PyTorch

Implementing MLP

```
# Multi-Layer non-linear Classifier

class MLP(nn.Module):

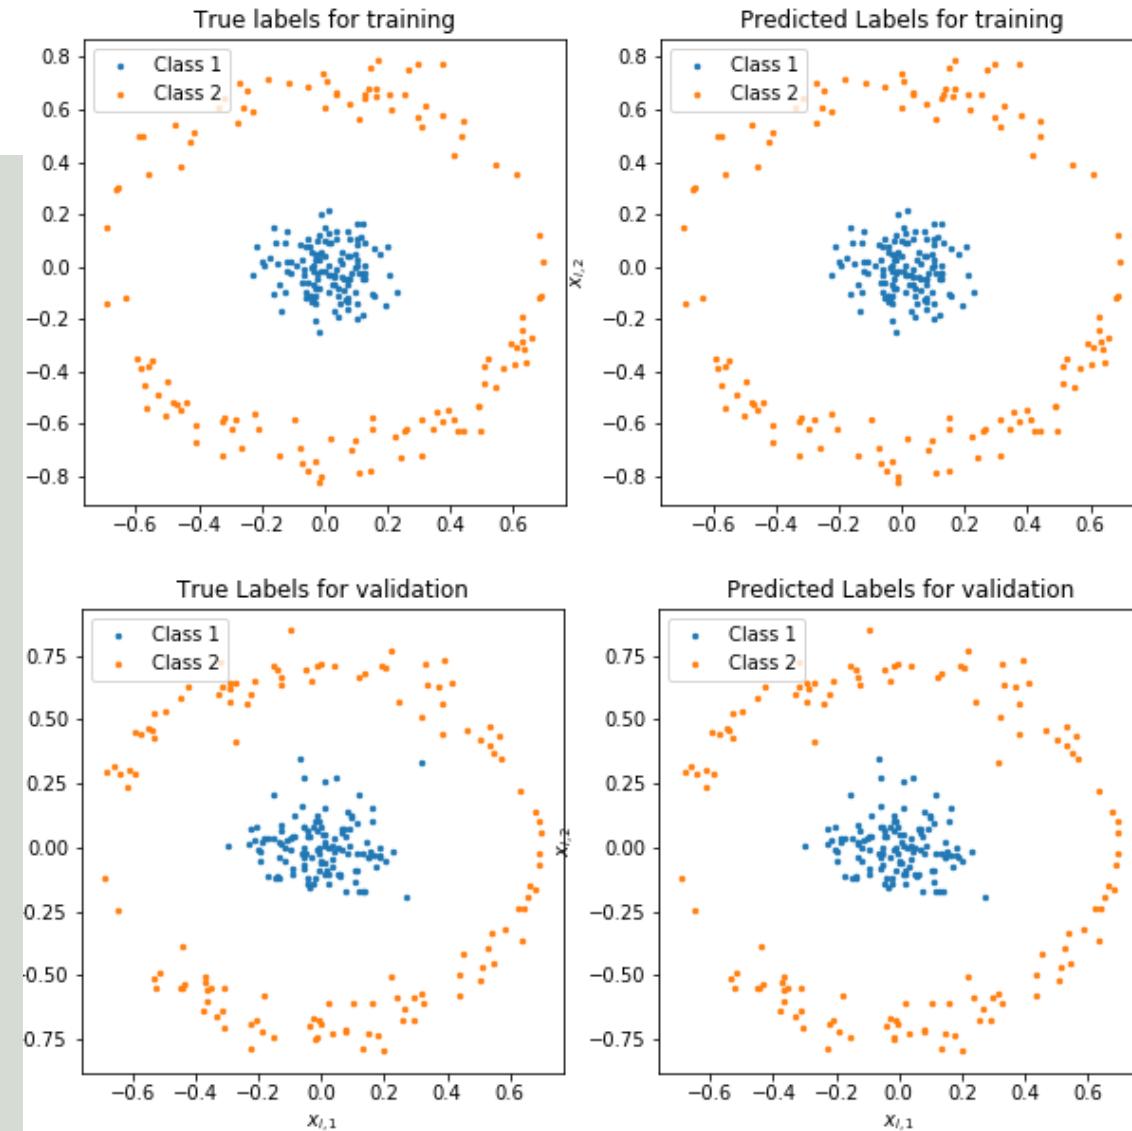
    def __init__(self, num_classes): # define the layer(s) and network components
        super(MLP, self).__init__()
        self.layer1= nn.Linear(in_features=num_features, out_features=4, bias=True)
        self.out= nn.Linear(in_features=4, out_features=1, bias=True)
        self.non_linearity = nn.Sigmoid()

    def forward(self, x): # here we define the forward pass for network
        y = self.layer1(x)
        y = self.non_linearity(y)
        y = self.out(y)
        y = self.non_linearity(y)
        return y

network = MLP(num_classes=1)

optimizer = optim.SGD(network.parameters(), lr=1.2) criterion = nn.BCELoss()

for i in range(max_iter):
    network.zero_grad() # clearing all gradients in the network
    y_train_hat = network(X_train_tensor) #forward pass
    loss = criterion(y_train_hat, y_train_tensor)
    loss.backward()
    optimizer.step()
```



Appendix A: Notations

- x_i : a single feature
- \boldsymbol{x}_i : feature vector (a data sample)
- $\boldsymbol{x}_{:,i}$: feature vector of all data samples
- \boldsymbol{X} : matrix of feature vectors (dataset)
- N : number of data samples
- P : number of features in a feature vector
- $k^{(i)}$: number of neurons in the i^{th} layer
- α : learning rate
- Bold letter/symbol: vector
- Bold capital letters/symbol: matrix