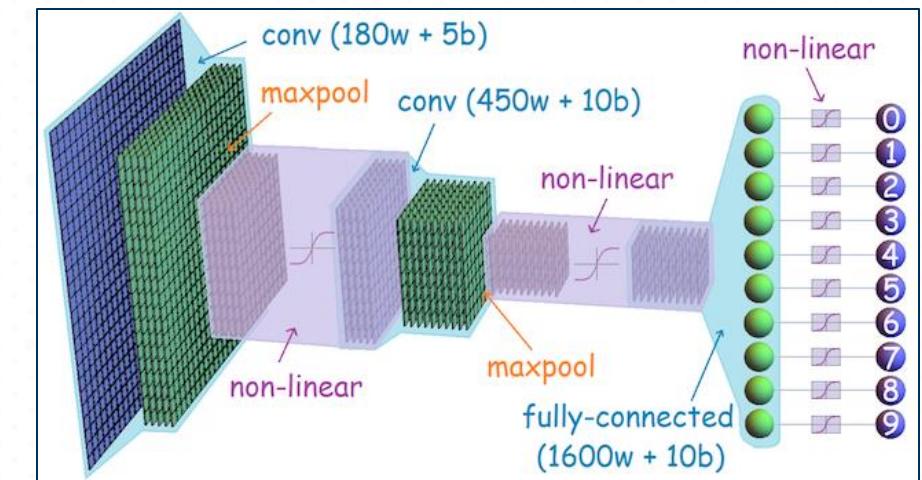


ECE 4252/8803: Fundamentals of Machine Learning (FunML)

Fall 2024

Lecture 14: Convolutional Neural Networks



Overview

In this Lecture..

Locally Connected Layer

- Limitations of ANNNs
- Convolutional Layers

Introduction to Convolutional Neural Networks

Layers used to build Convolutional Neural Networks

- Convolutional layer
- Pooling Layer
- Fully-connected layers

Examples of Deep CNNs Architectures

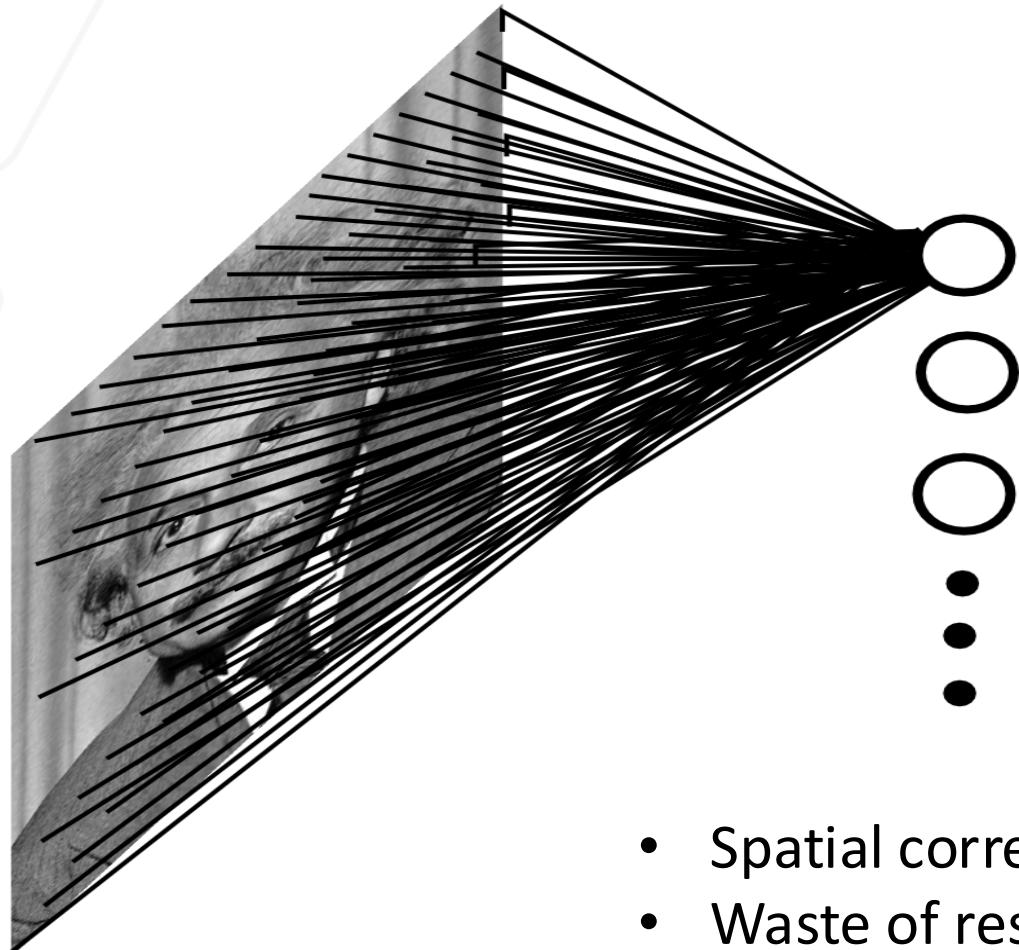
Training CNNs

- Optimization

Visualization of Convolutional Neural Networks

MLPs/ANNs

Limitations of Multi Layer Perceptrons



Example of a single layer Neural Network:

Input: 200x200 pixels image
(vectorized image $x \in \mathbb{R}^{4 \times 10^4}$)

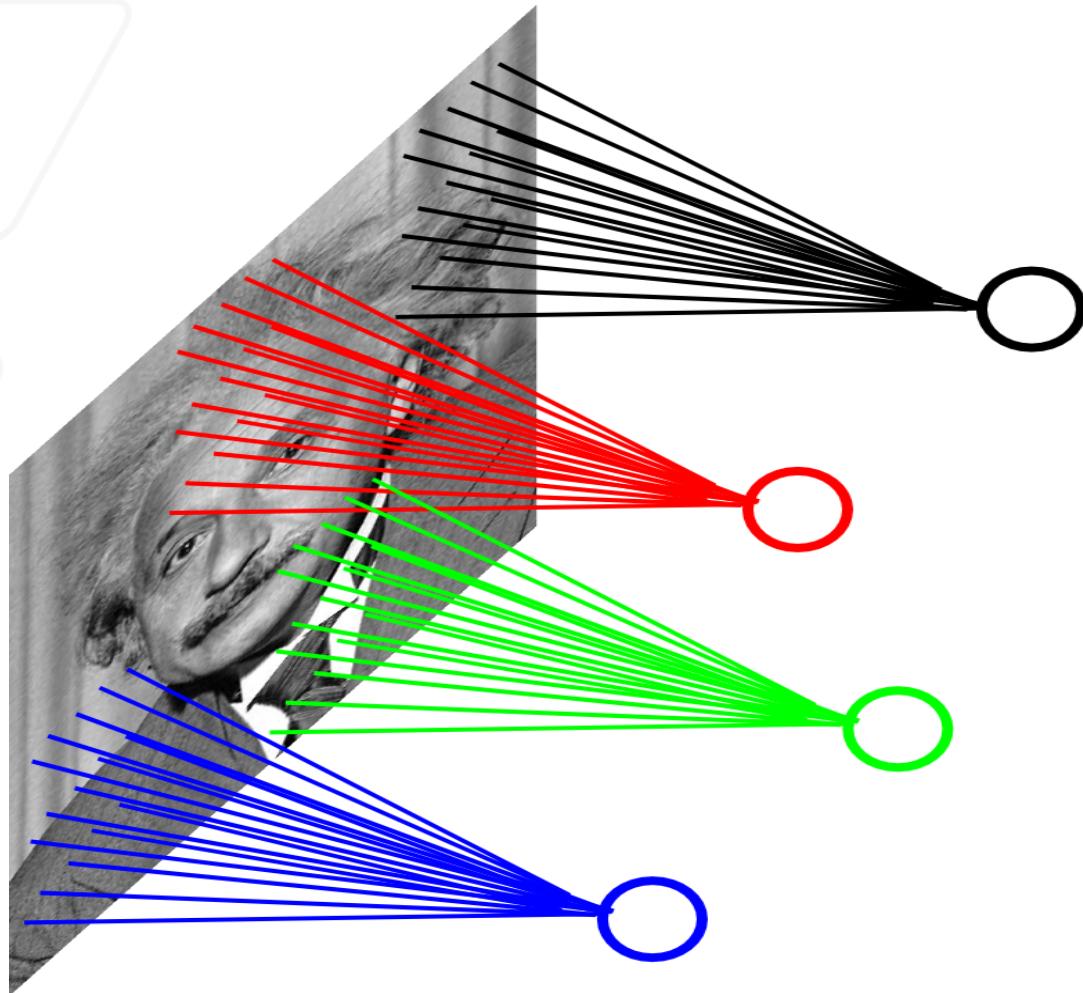
For a single neuron, $w_i \in \mathbb{R}^{4 \times 10^4}$ (4×10^4 weights)

For 100 neurons: $4 \times 10^4 \times 100$
~4M parameters!!!

- Spatial correlation is local
- Waste of resources while not enough training samples...

Locally Connected Neural networks

Local Connectivity



Example of a locally connected layer:

Input: 200x200 pixels image (vectorized image $x \in \mathbb{R}^{4 \times 10^4}$)

For a single neuron (filter), $w_i \in \mathbb{R}^{3 \times 3}$ (**9 shared weights**)

For 100 neurons: 9 x100
900 parameters

Locally Connected Neural networks

Local Connectivity

This is good when:

- The inputs are (roughly) registered (e.g., face recognition)

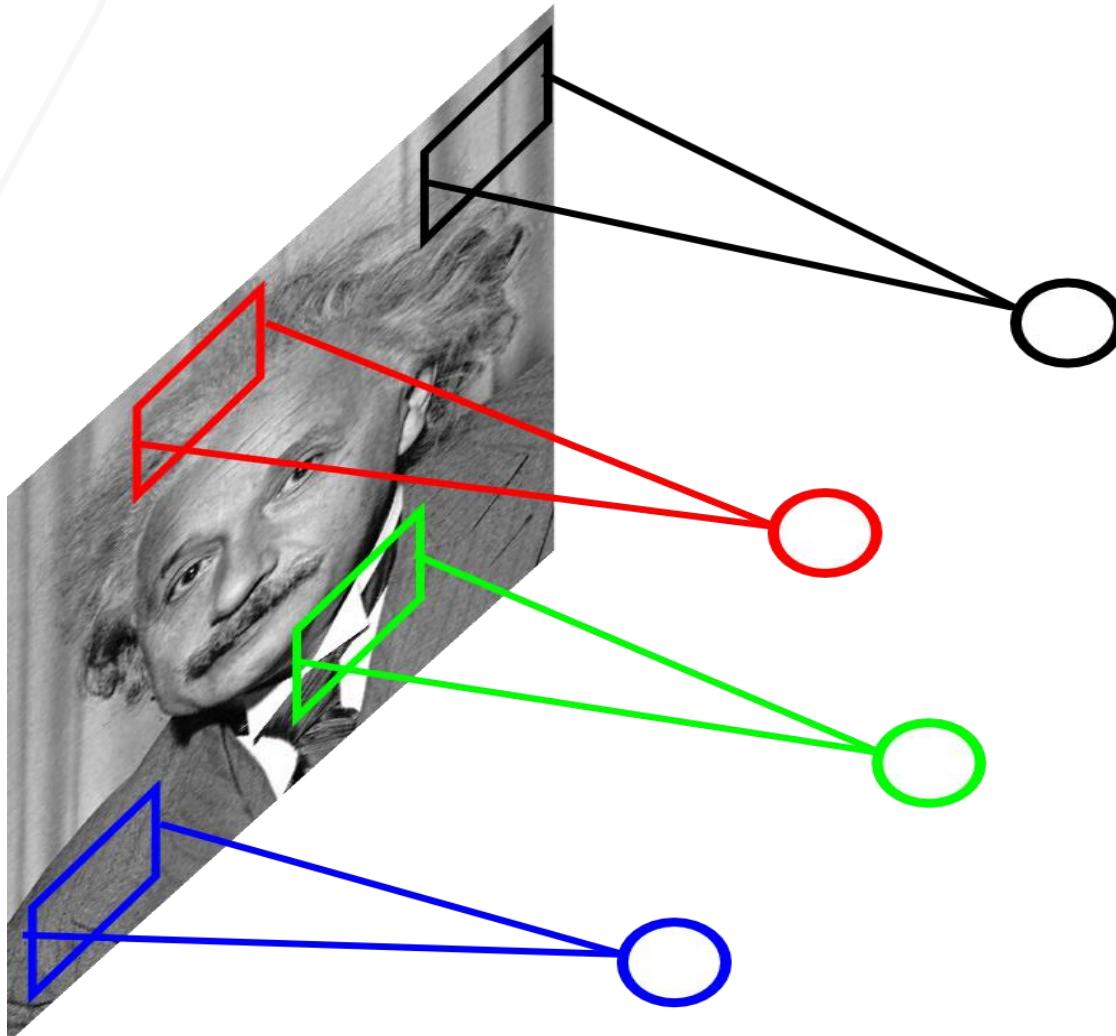


- The object can be anywhere



Locally Connected Neural networks

Local Connectivity

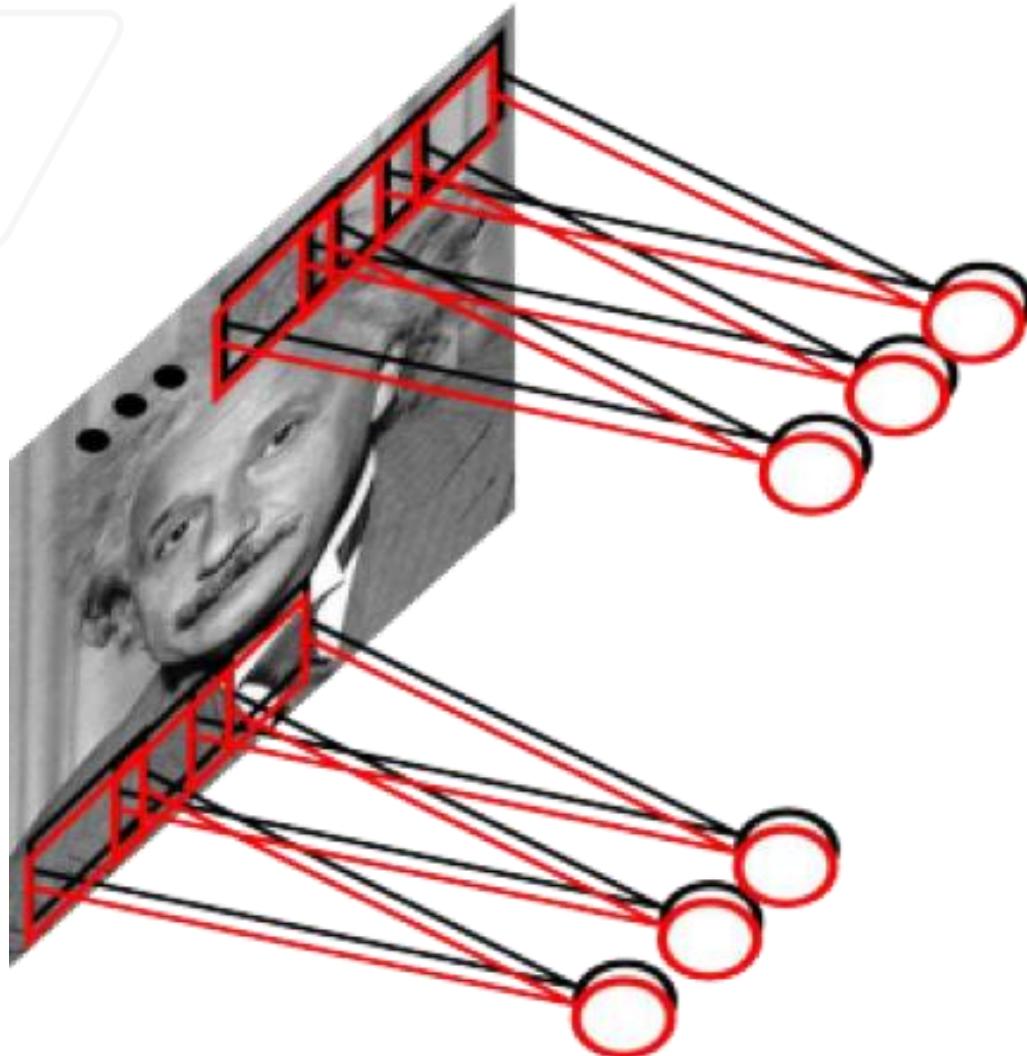


STATIONARITY?

Statistics are similar at different locations

Locally Connected Neural networks

Parameter Sharing



- Connect each neuron to an input patch
- Share the same parameters across different locations (assuming input is stationary):

Convolutions with learned kernels

Locally Connected Neural networks

ANNs vs CNNs

Artificial Neural Networks:

- number of parameters scales quadratically with the size of the input
- does not leverage stationarity

Convolutional layers:

- connect each neuron to a patch of the input
- share the weight across spatial locations

Convolution

1D and 2D Convolutions

- One-dimensional:

$$y(t) = \int_{-\infty}^{+\infty} x(t-a)w(a)da$$

- Two-dimensional:

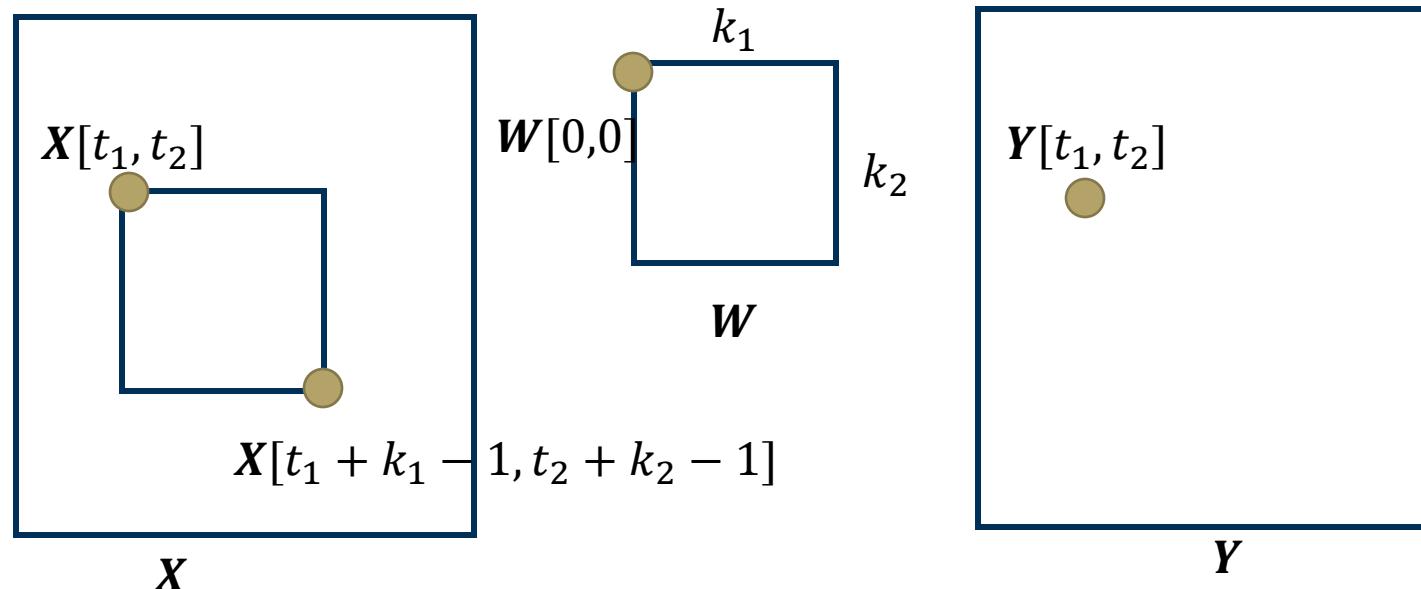
$$y(t_1, t_2) = \int_{-\infty}^{+\infty} \int_{-\infty}^{+\infty} x(t_1 - a, t_2 - b)w(a, b) da db$$

Convolution

2D Discrete Convolution

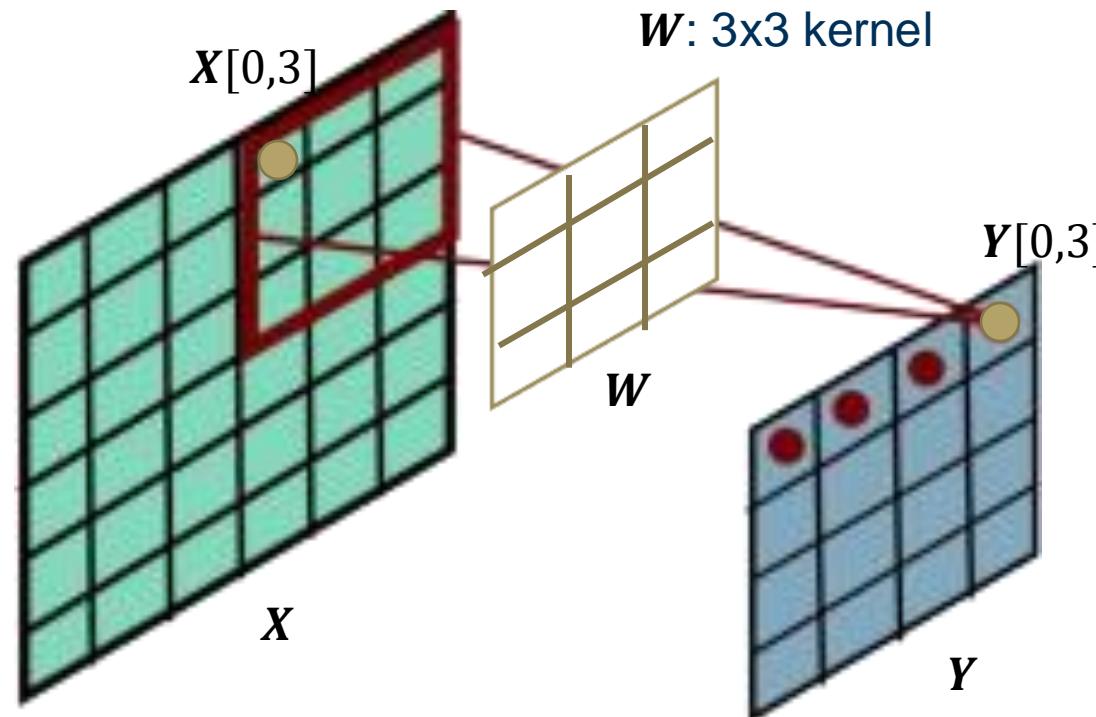
- Two-dimensional (discrete):

$$Y[t_1, t_2] = \sum_{a=0}^{k_1-1} \sum_{b=0}^{k_2-1} X[t_1 + a, t_2 + b] W[a, b]$$



Convolutional Layer

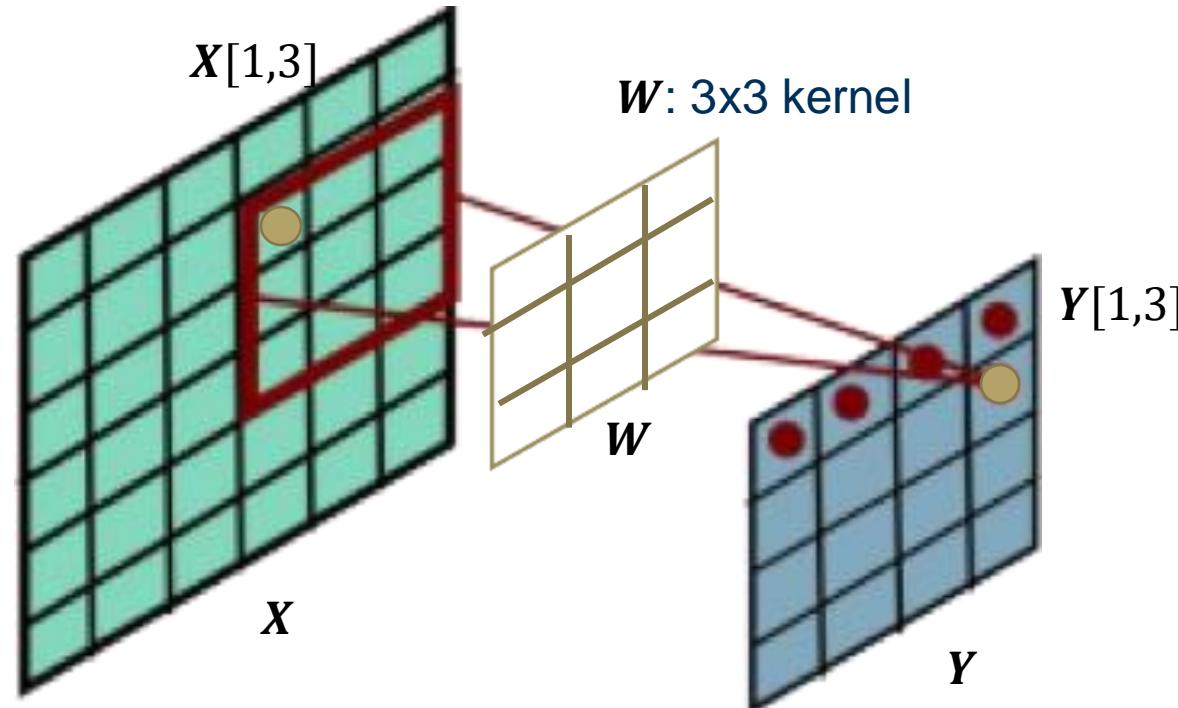
2D Discrete Convolution



$$Y[0,3] = \sigma \left(\sum_{a=0}^{3-1} \sum_{b=0}^{3-1} X[0 + a, 3 + b] W[a, b] \right) \quad (\sigma : \text{non-linear activation})$$

Convolutional Layer

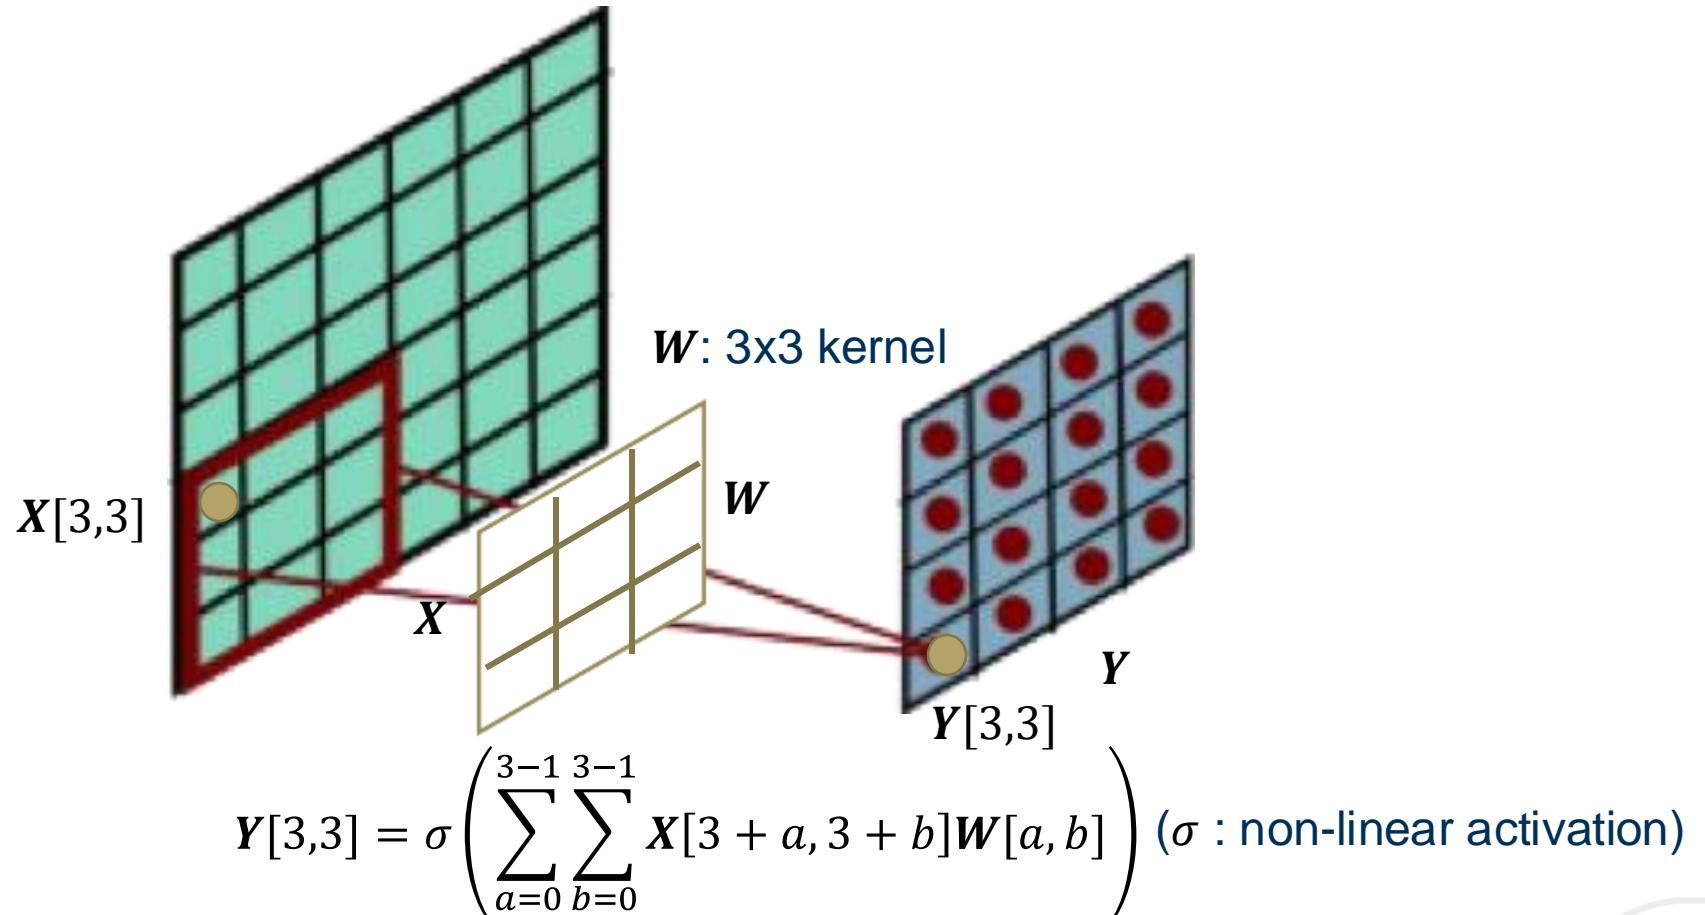
2D Discrete Convolution



$$Y[1,3] = \sigma \left(\sum_{a=0}^{3-1} \sum_{b=0}^{3-1} X[1+a, 3+b] W[a, b] \right) \quad (\sigma : \text{non-linear activation})$$

Convolutional Layer

2D Discrete Convolution



Convolutional Layer

Kernel Intuition

A 3x3 kernel slides over the input features, performing an element-wise multiplication with the current patch of the input, summing up the results into a single output at every location it slides over.

3 ₀	3 ₁	2 ₂	1	0
0 ₂	0 ₂	1 ₀	3	1
3 ₀	1 ₁	2 ₂	2	3
2	0	0	2	2
2	0	0	0	1

input features (5x5 matrix)

12.0	12.0	17.0
10.0	17.0	19.0
9.0	6.0	14.0

output features (3x3 matrix)

Kernel:
$$\begin{bmatrix} 0 & 1 & 2 \\ 2 & 2 & 0 \\ 0 & 1 & 2 \end{bmatrix}$$

Convolutional Layer

Kernel

The Convolutional layer is the core building block of a Convolutional Network that does most of the computational heavy lifting.

The properties of convolutional layer:

- Local Connectivity
- Parameter sharing
- Spatial arrangement

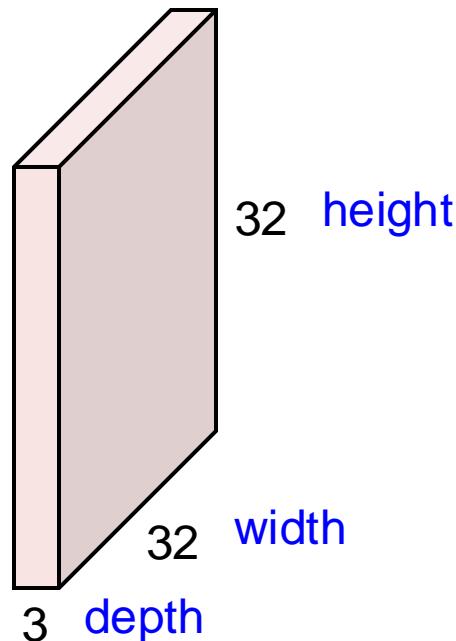
In this course: Kernels = Filters

Convolutional Layer

Spatial Arrangement

(32x32)x3 image → preserve spatial structure

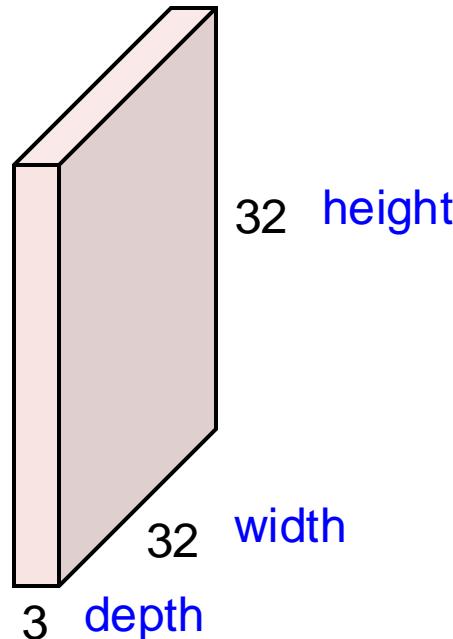
Multi-dimensional input



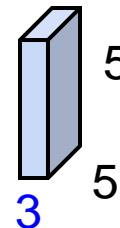
Convolutional Layer

Spatial Arrangement

(32x32)x3 image



(5x5)x3 filters*



Filters always extend the full depth of the input volume

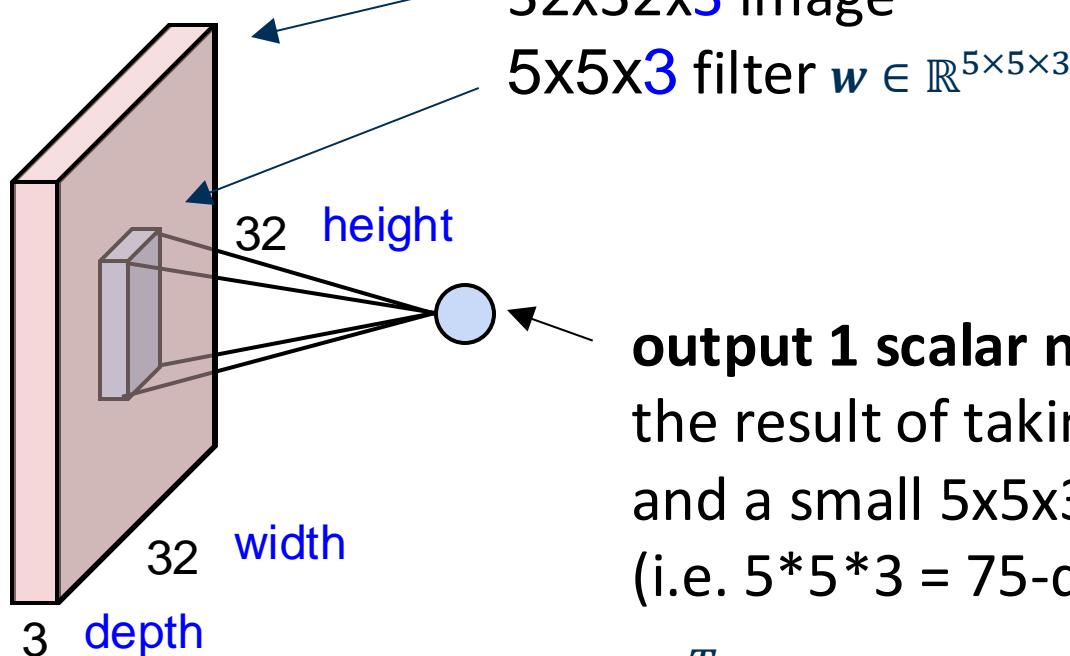
Convolve the filter with the image i.e. “slide over the image spatially, computing dot products”

* A collection of 3 unique 2-dimensional filters, of which dimension is 5x5

Convolutional Layer

Spatial Arrangement

consider the first filter



output 1 scalar number:

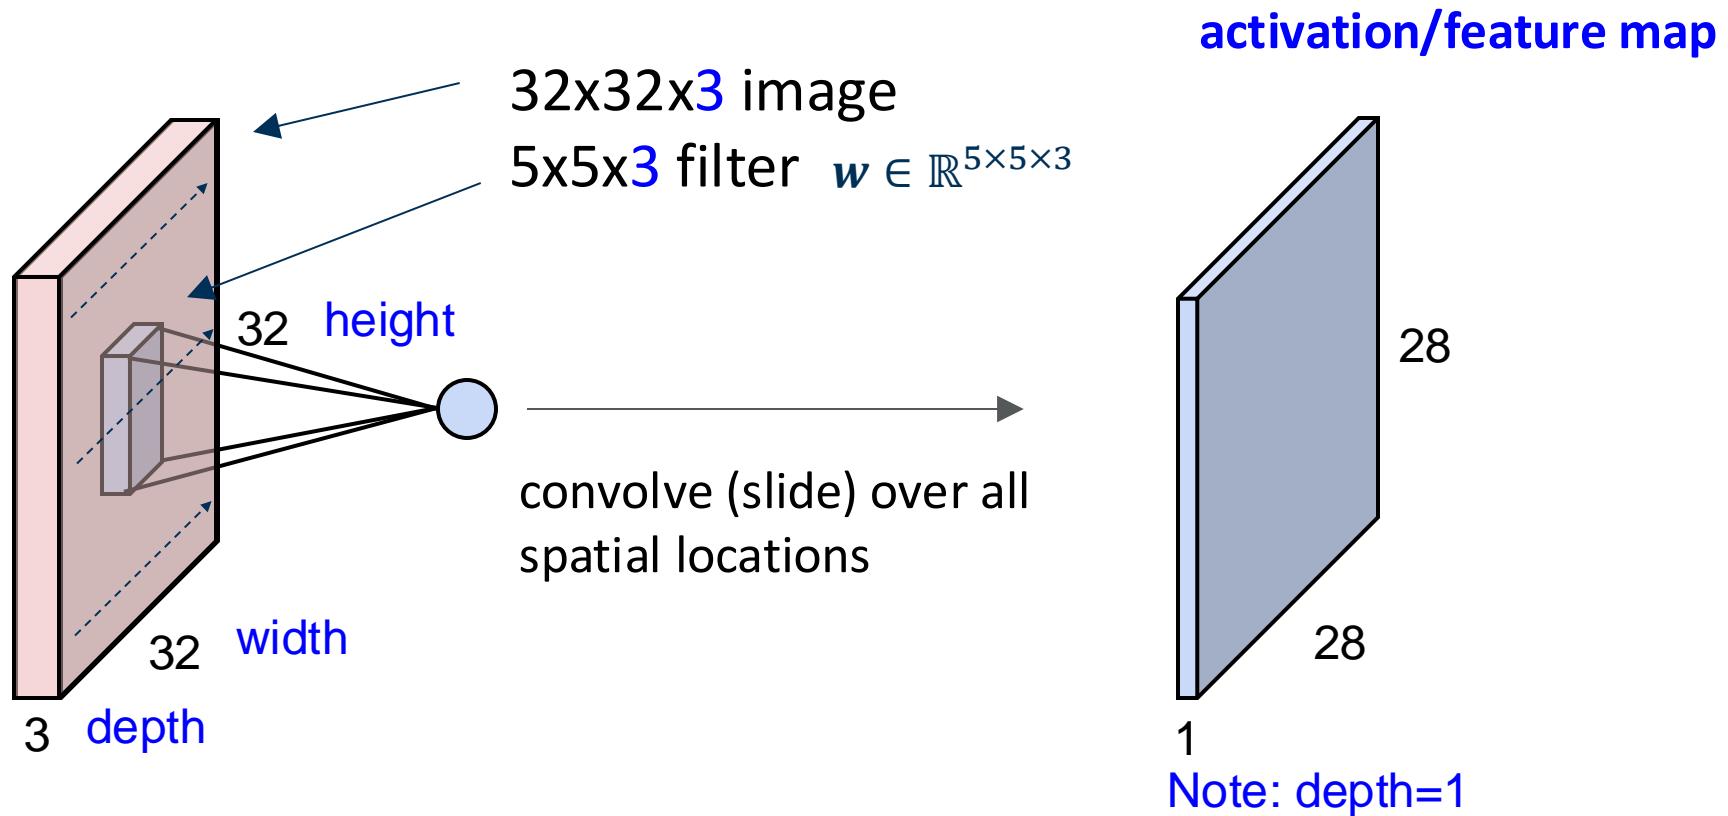
the result of taking a dot product between filter
and a small 5x5x3 patch of the image
(i.e. $5 \times 5 \times 3 = 75$ -dimensional dot product + bias)

$$w^T x + b \text{ (vectorized)}$$

where $x \in \mathbb{R}^{5 \times 5 \times 3}$

Convolutional Layer

Spatial Arrangement

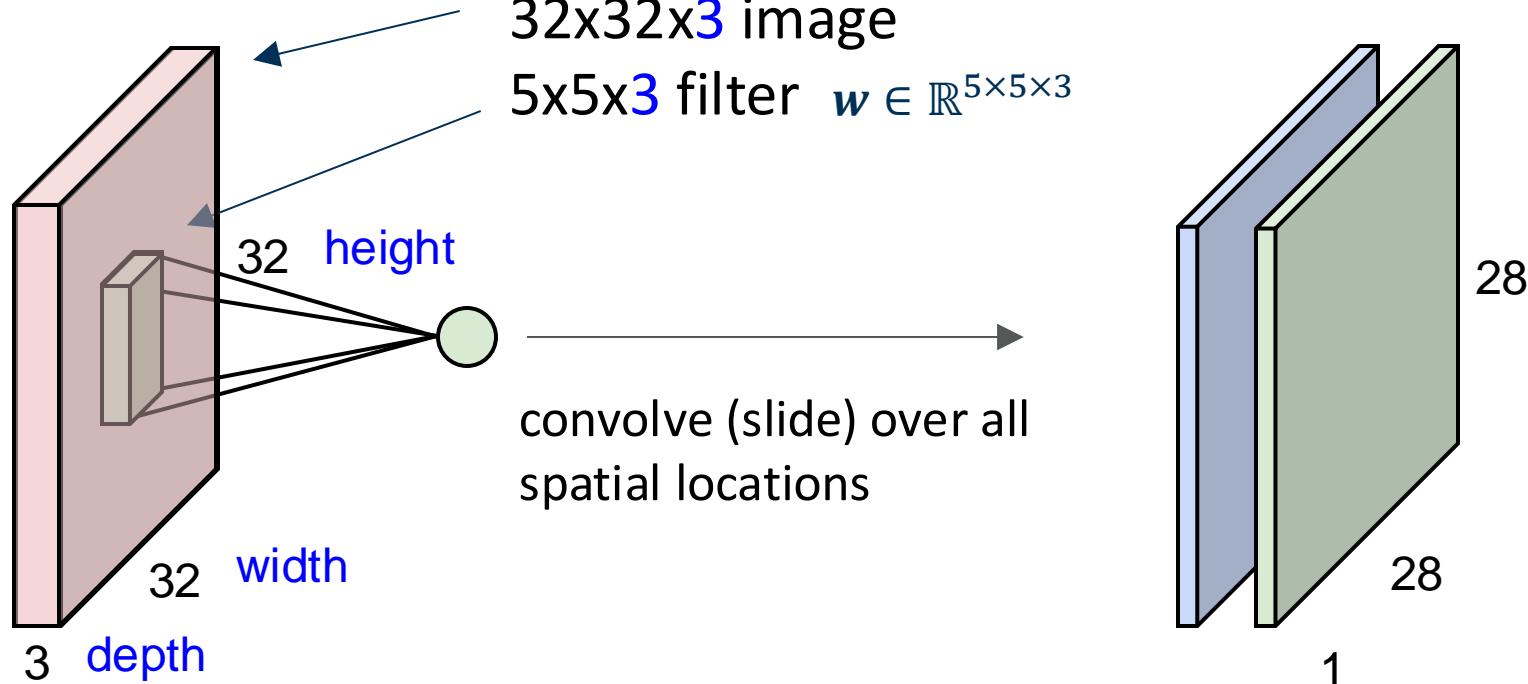


Convolutional Layer

Spatial Arrangement

consider a second filter

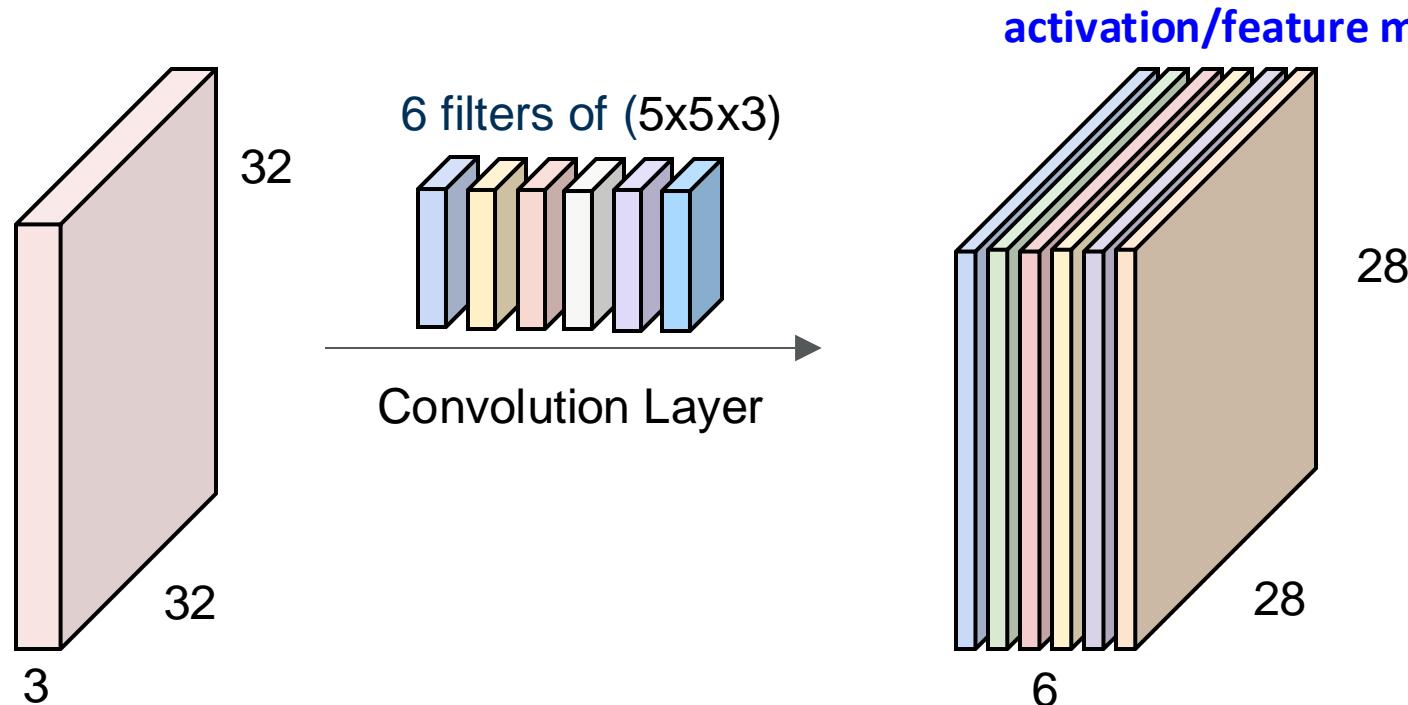
activation/feature maps



Convolutional Layer

Spatial Arrangement

If we had 6 filters of size $5 \times 5 \times 3$, the outputs are 6 separate activation maps:



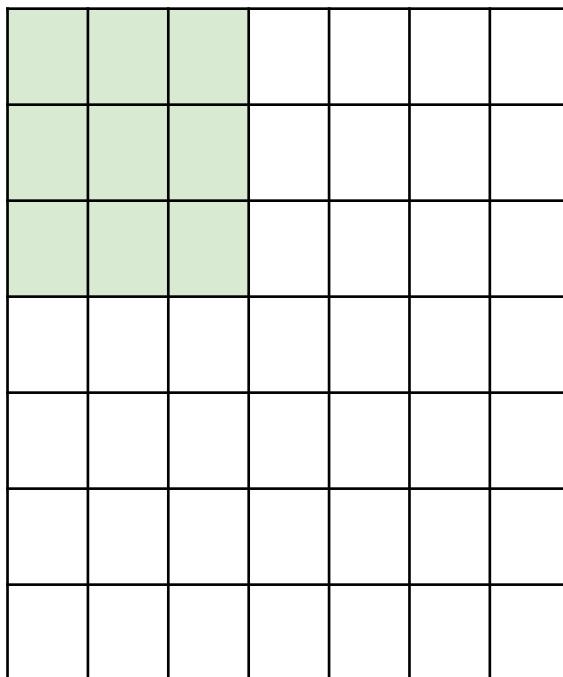
We stack these up to get a set of output feature maps of size $28 \times 28 \times 6$!

Convolutional Layer

Spatial Arrangement

A closer look at spatial dimensions:

7



input

7x7 input (spatially)
assume 3x3 filter



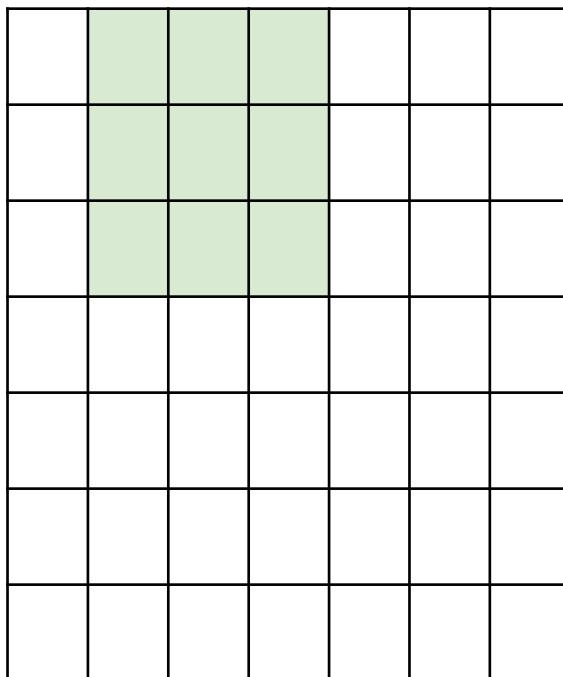
output

Convolutional Layer

Spatial Arrangement

A closer look at spatial dimensions:

7



input

7x7 input (spatially)
assume 3x3 filter



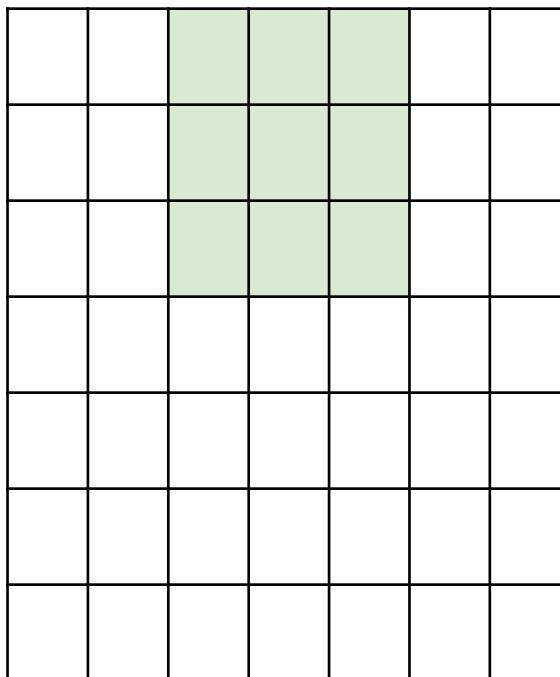
output

Convolutional Layer

Spatial Arrangement

A closer look at spatial dimensions:

7



input

7x7 input (spatially)
assume 3x3 filter

7



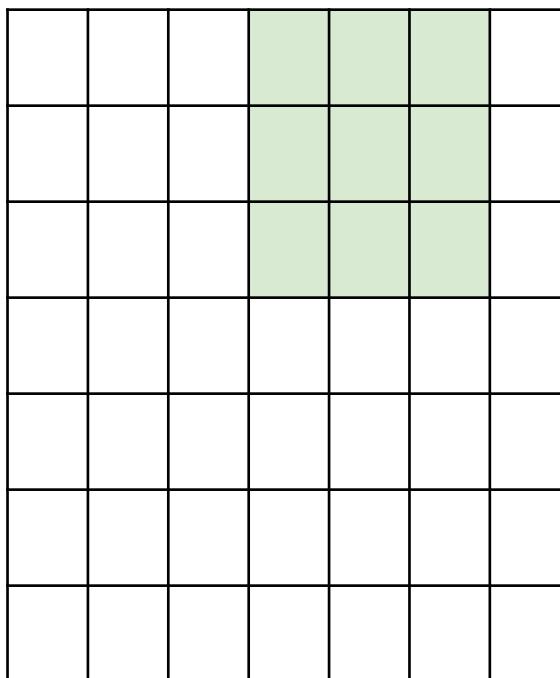
output

Convolutional Layer

Spatial Arrangement

A closer look at spatial dimensions:

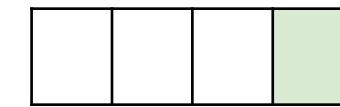
7



input

7x7 input (spatially)
assume 3x3 filter

7



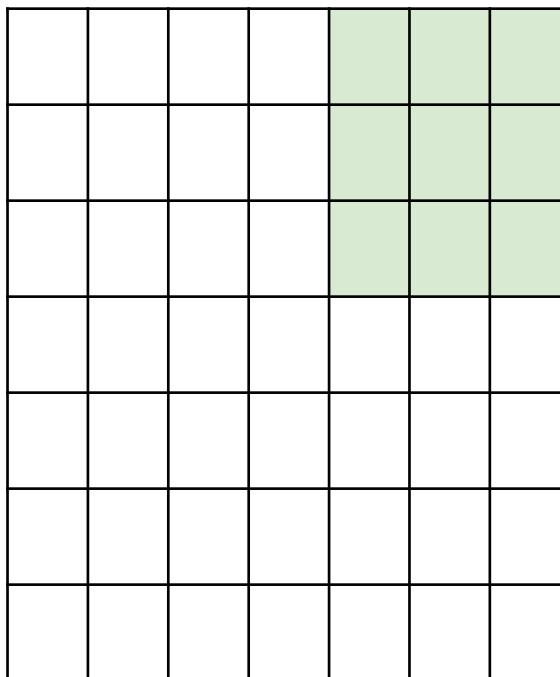
output

Convolutional Layer

Spatial Arrangement

A closer look at spatial dimensions:

7



input

7x7 input (spatially)
assume 3x3 filter

7

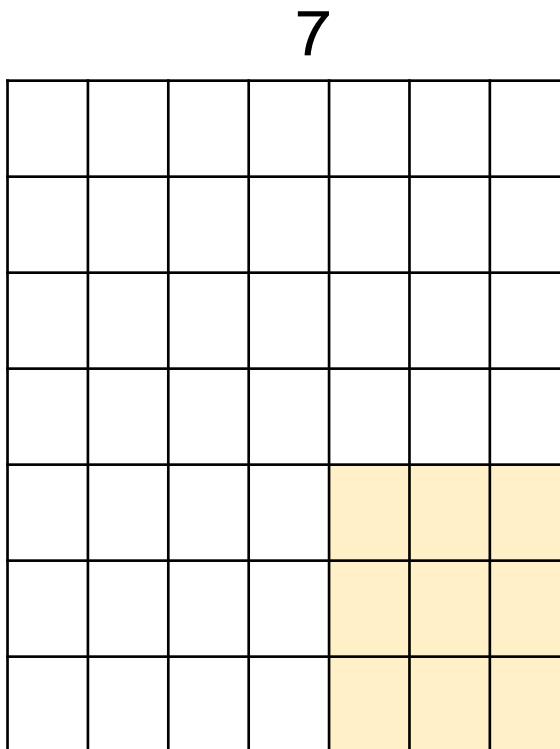


output

Convolutional Layer

Spatial Arrangement

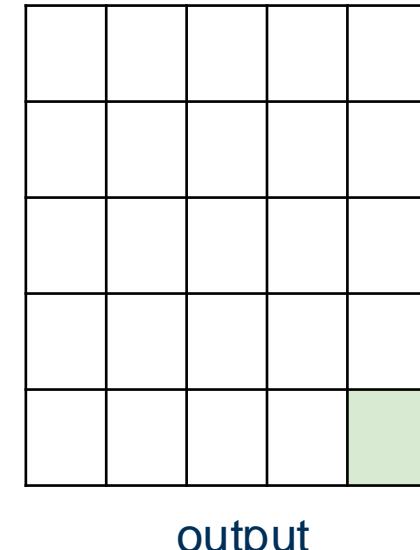
A closer look at spatial dimensions:



input

7x7 input (spatially)
assume 3x3 filter

=> 5x5 output

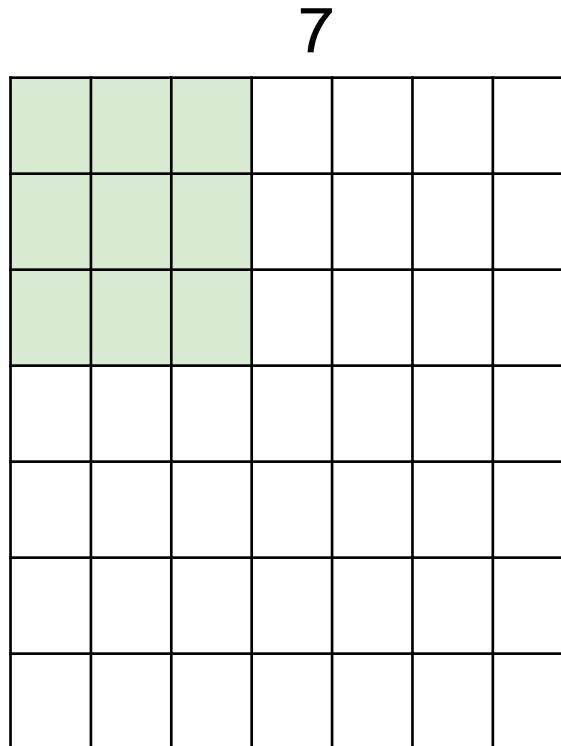


output

Convolutional Layer

Spatial Arrangement

A closer look at spatial dimensions:



input

7x7 input (spatially)
assume 3x3 filter
applied **with stride 2**

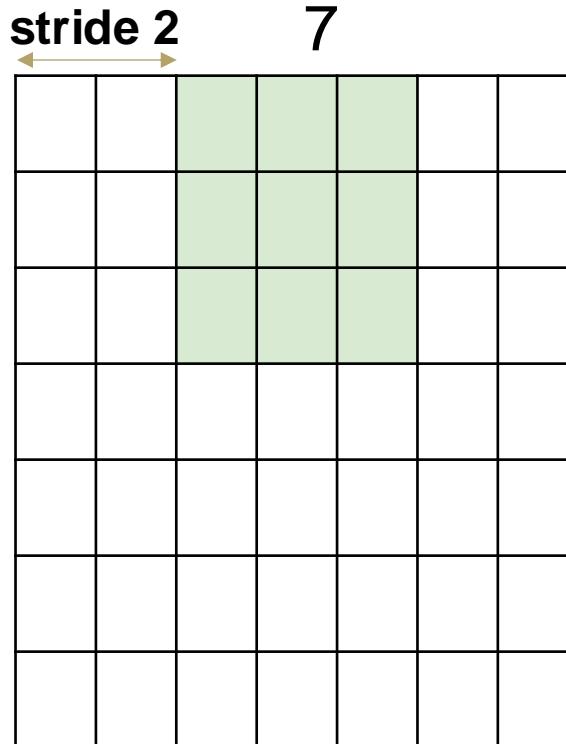


output

Convolutional Layer

Spatial Arrangement

A closer look at spatial dimensions:



input

7x7 input (spatially)
assume 3x3 filter
applied **with stride 2**

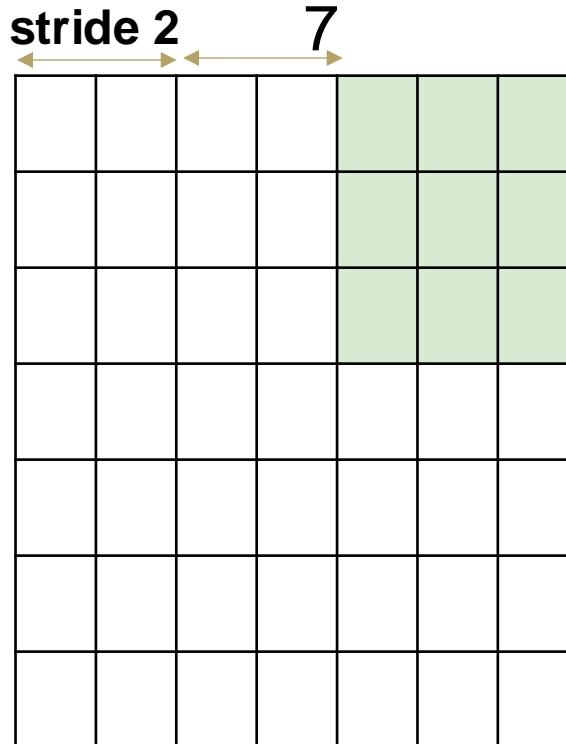


output

Convolutional Layer

Spatial Arrangement

A closer look at spatial dimensions:



input

7x7 input (spatially)
assume 3x3 filter
applied **with stride 2**

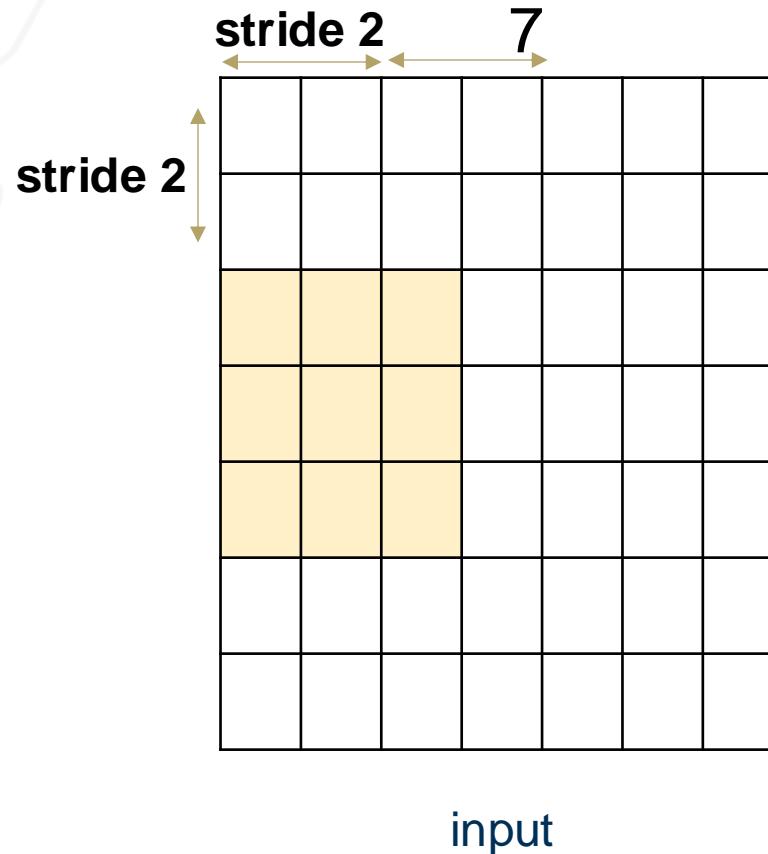


output

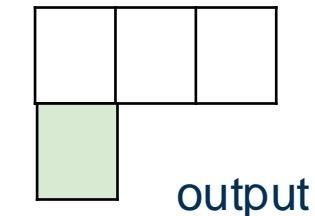
Convolutional Layer

Spatial Arrangement

A closer look at spatial dimensions:



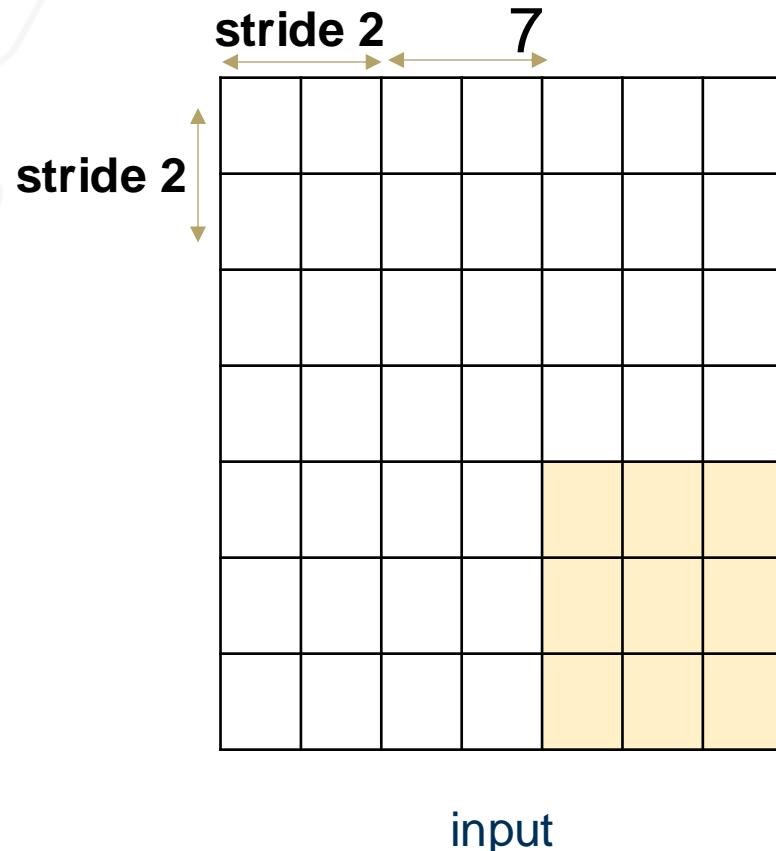
7x7 input (spatially)
assume 3x3 filter
applied **with stride 2**



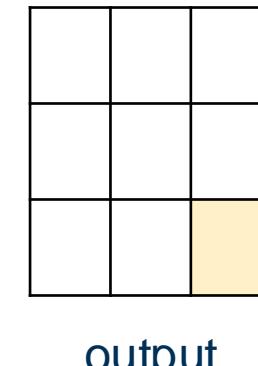
Convolutional Layer

Spatial Arrangement

A closer look at spatial dimensions:



7x7 input (spatially)
assume 3x3 filter
applied **with stride 2**
7
=> 3x3 output!

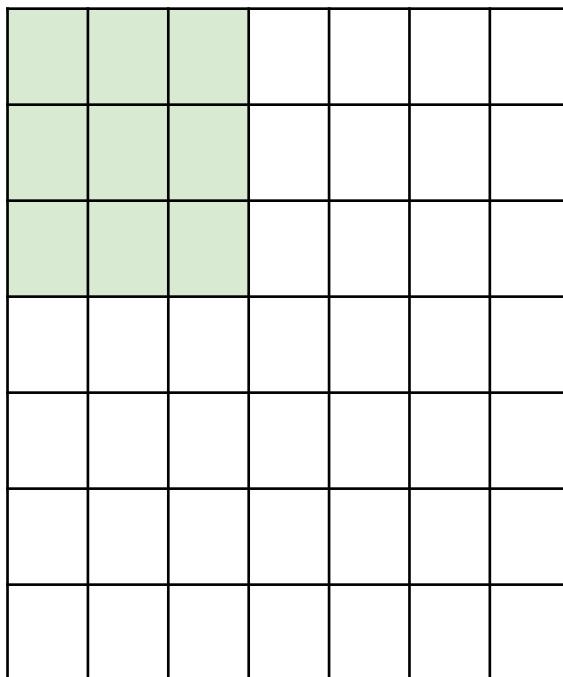


Convolutional Layer

Spatial Arrangement

A closer look at spatial dimensions:

7



input

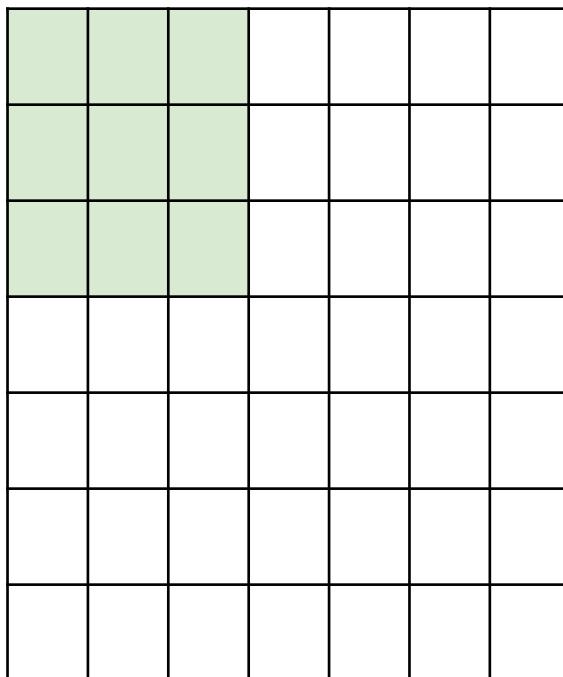
7x7 input (spatially)
assume 3x3 filter
applied **with stride 3?**

Convolutional Layer

Spatial Arrangement

A closer look at spatial dimensions:

7



input

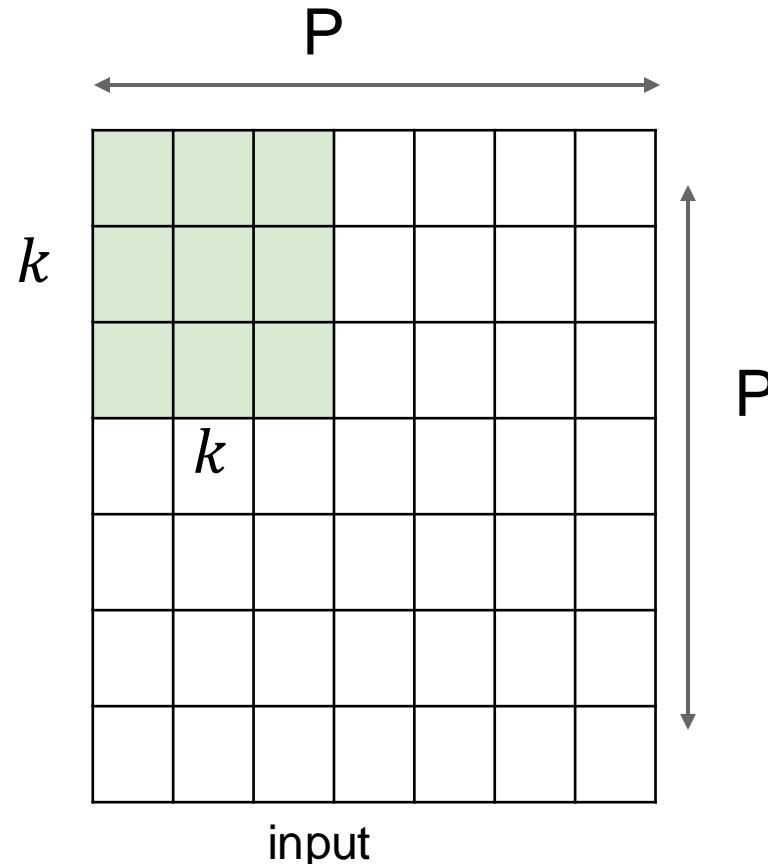
7x7 input (spatially)
assume 3x3 filter
applied **with stride 3?**

7

doesn't fit!
cannot apply 3x3 filter on
7x7 input with stride 3.

Convolutional Layer

Spatial Arrangement



- Input size: $P \times P$
- Filter (kernel) size: $k \times k$
- Output size: **$(P - k) / \text{stride} + 1$**

e.g., $P = 7$, $k = 3$:

$$\text{stride } 1 \Rightarrow (7 - 3)/1 + 1 = 5$$

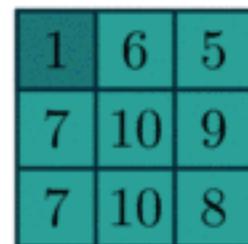
$$\text{stride } 2 \Rightarrow (7 - 3)/2 + 1 = 3$$

$$\text{stride } 3 \Rightarrow (7 - 3)/3 + 1 = \mathbf{2.33}$$

Convolutional Layer

Spatial Arrangement

zero-padding & strides



kernel: $\begin{pmatrix} 2 & 0 & 1 \\ 1 & 0 & 0 \\ 0 & 1 & 1 \end{pmatrix}$

Padded with 1x1 border of zeros

stride 2: kernel moves 2 pixels at the input

Output size: $(P - k) / \text{stride} + 1 = (7 - 3) / 2 + 1 = 3$

Convolutional Layer

Spatial Arrangement

0	0	0	0	0	0			
0								
0								
0								
0								

e.g. input 7x7

3x3 filter, applied with **stride 1**

pad with 1 pixel border => what is the output?

7x7 output!

in general, if convolutional layers with stride 1,
filters of size $k \times k$, and zero-padding with $(k-1)/2$.
(will preserve size spatially)

e.g. $k = 3 \Rightarrow$ zero pad with 1

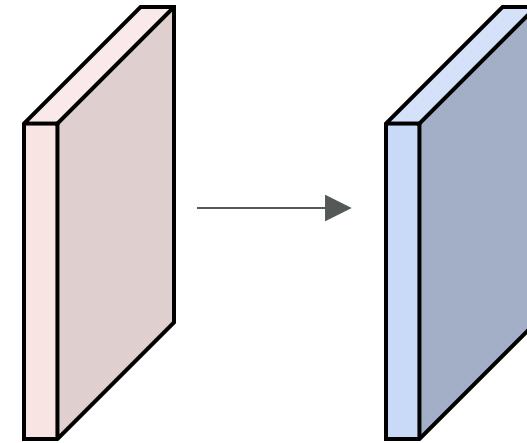
$k = 5 \Rightarrow$ zero pad with 2

$k = 7 \Rightarrow$ zero pad with 3

Convolutional Layer

Output Dimensions

- Input volume: **32x32x3**
- **10 5x5 filters with stride 1, pad 2**



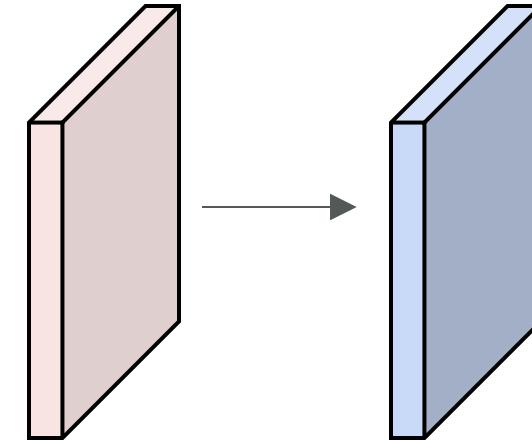
- Output volume size:
 $(32+2*2-5)/1+1 = 32$ spatially, so
32x32x10

$$\frac{\text{Input Size} + 2 \times \text{Padding} - \text{Filter size}}{\text{Stride}} + 1$$

Convolutional Layer

Number of Parameters

- Input volume: **32x32x3**
- **10 5x5 filters with stride 1, pad 2**



- Number of parameters in this layer?
each filter has $5*5*3 + 1 = 76$ params (+1 for bias)
 $\Rightarrow 76*10 = 760$

Convolutional Layer

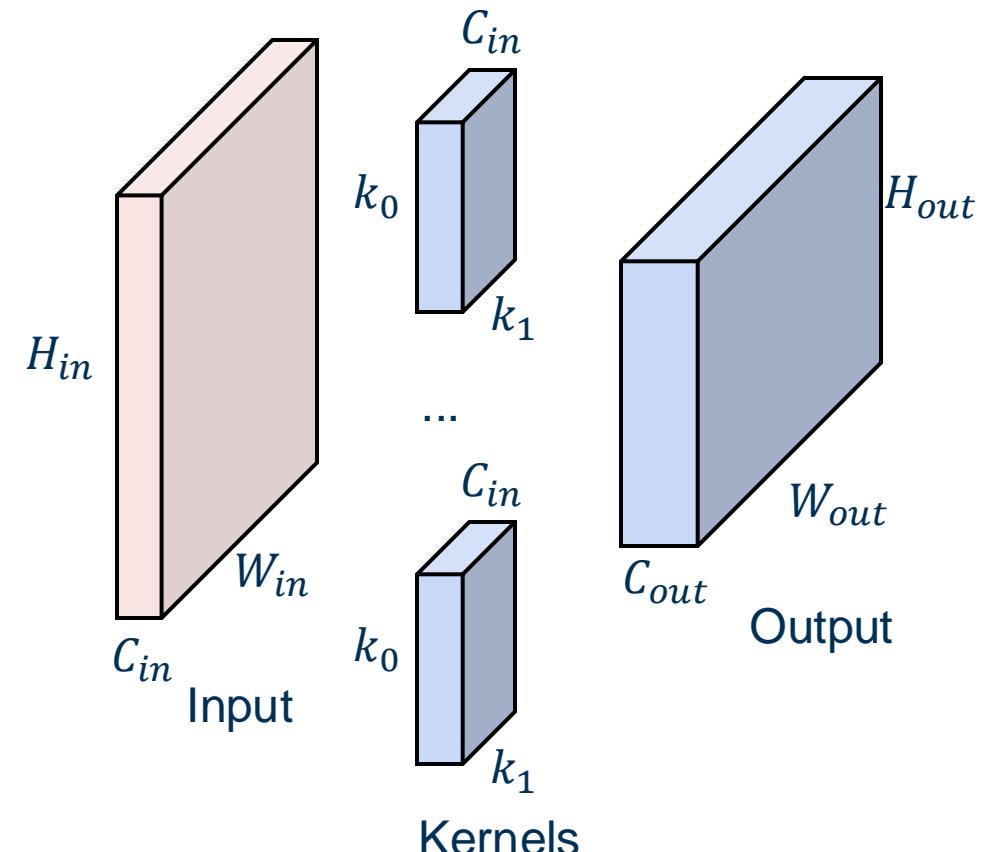
Dimensions

In general, for a convolutional layer:

- Input: $(N, C_{in}, H_{in}, W_{in})$
- Kernel: $(C_{out}, C_{in}, k_0, k_1)$
- Output: $(N, C_{out}, H_{out}, W_{out})$

where

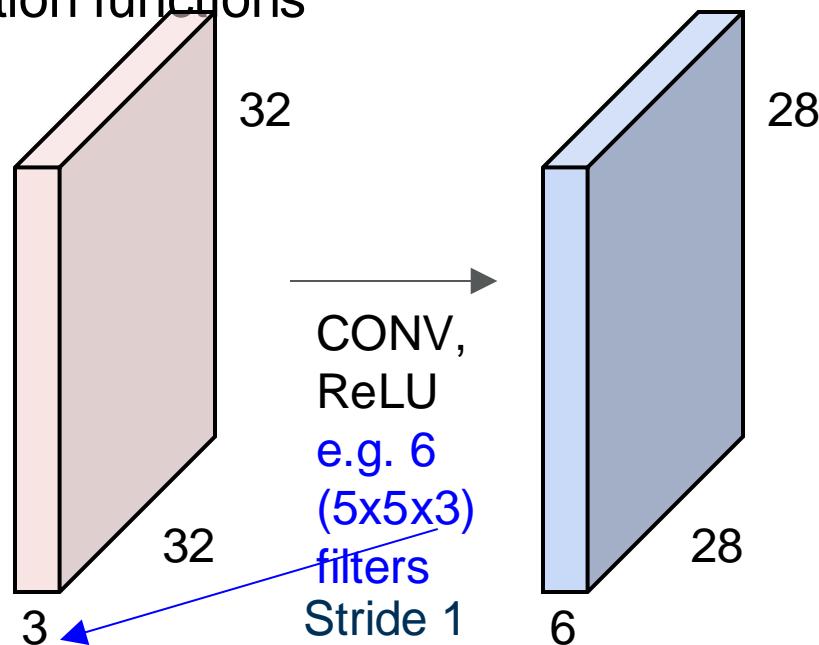
- $H_{out} = \left\lfloor \frac{H_{in} + 2 \times Padding_H - k_0}{Stride_H} + 1 \right\rfloor$
- $W_{out} = \left\lfloor \frac{W_{in} + 2 \times Padding_W - k_1}{Stride_W} + 1 \right\rfloor$
- N : number of samples
- $\lfloor \cdot \rfloor$: floor operation, $\lfloor x \rfloor$ is the largest integer i , such that $i \leq x$.



Convolutional Neural networks

Introduction

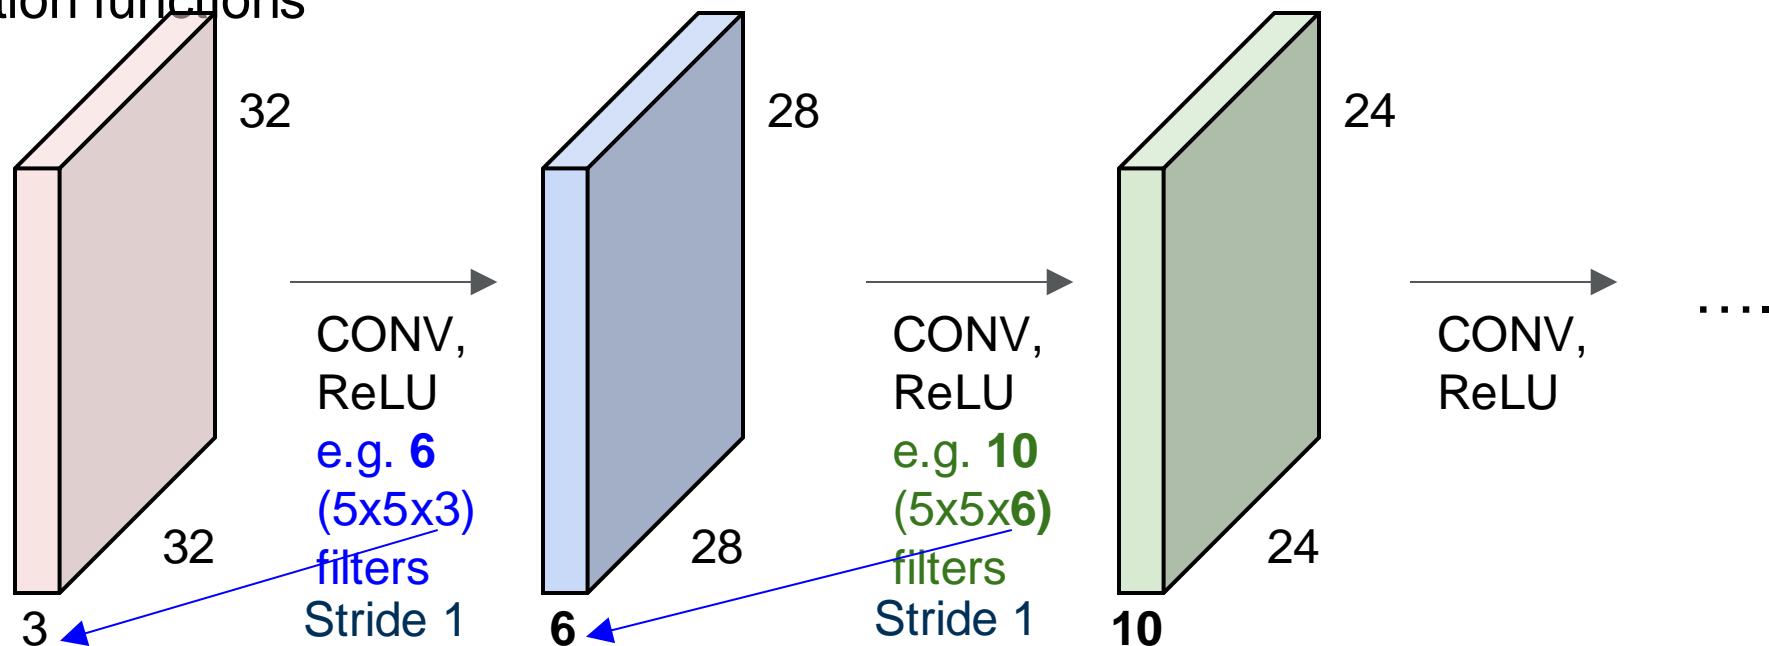
Preview: ConvNet is (initially) a sequence of Convolutional Layers, interspersed with activation functions



Convolutional Neural networks

Introduction

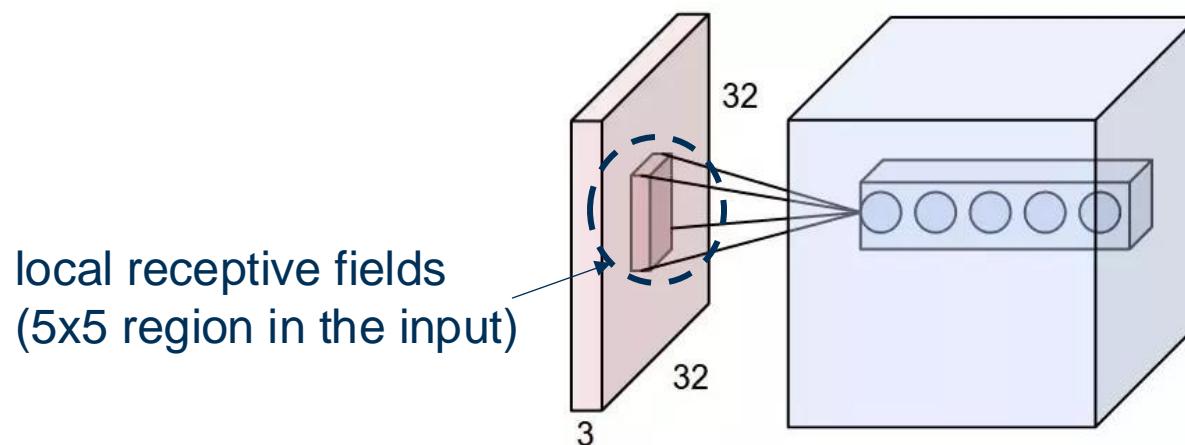
Preview: ConvNet is (initially) a sequence of Convolutional Layers, interspersed with activation functions



Convolutional Neural networks

Introduction

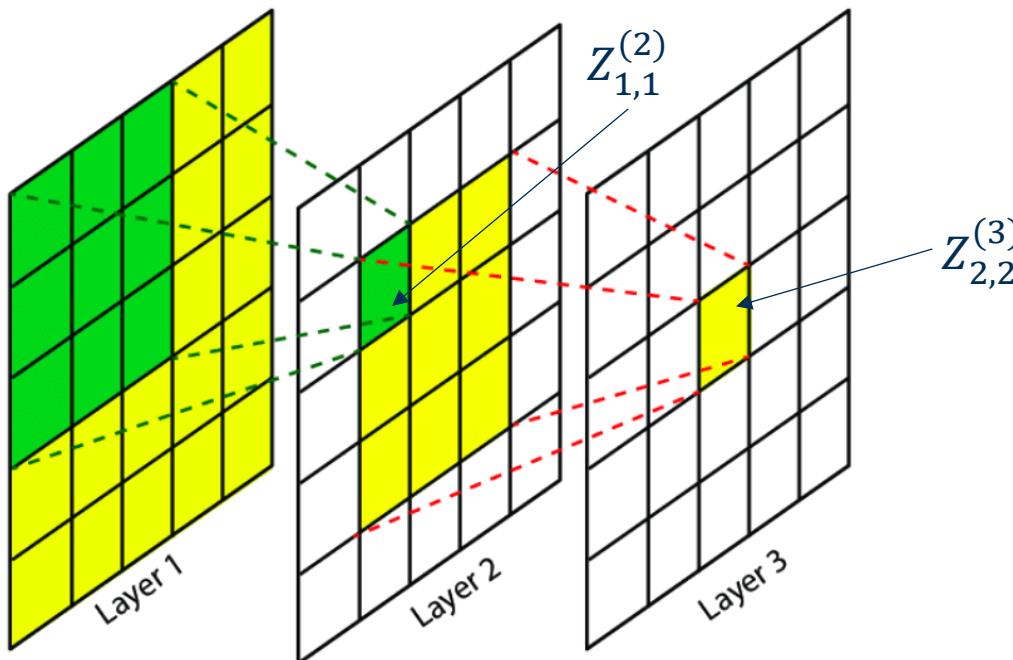
- Fully-connected networks:
 - Output of each neuron depends on the entire input
- Convolutional networks :
 - Output of a neuron only depends on a **local region** of the input. This region in the input is the **receptive field** for that neuron.



Convolutional Neural networks

Introduction

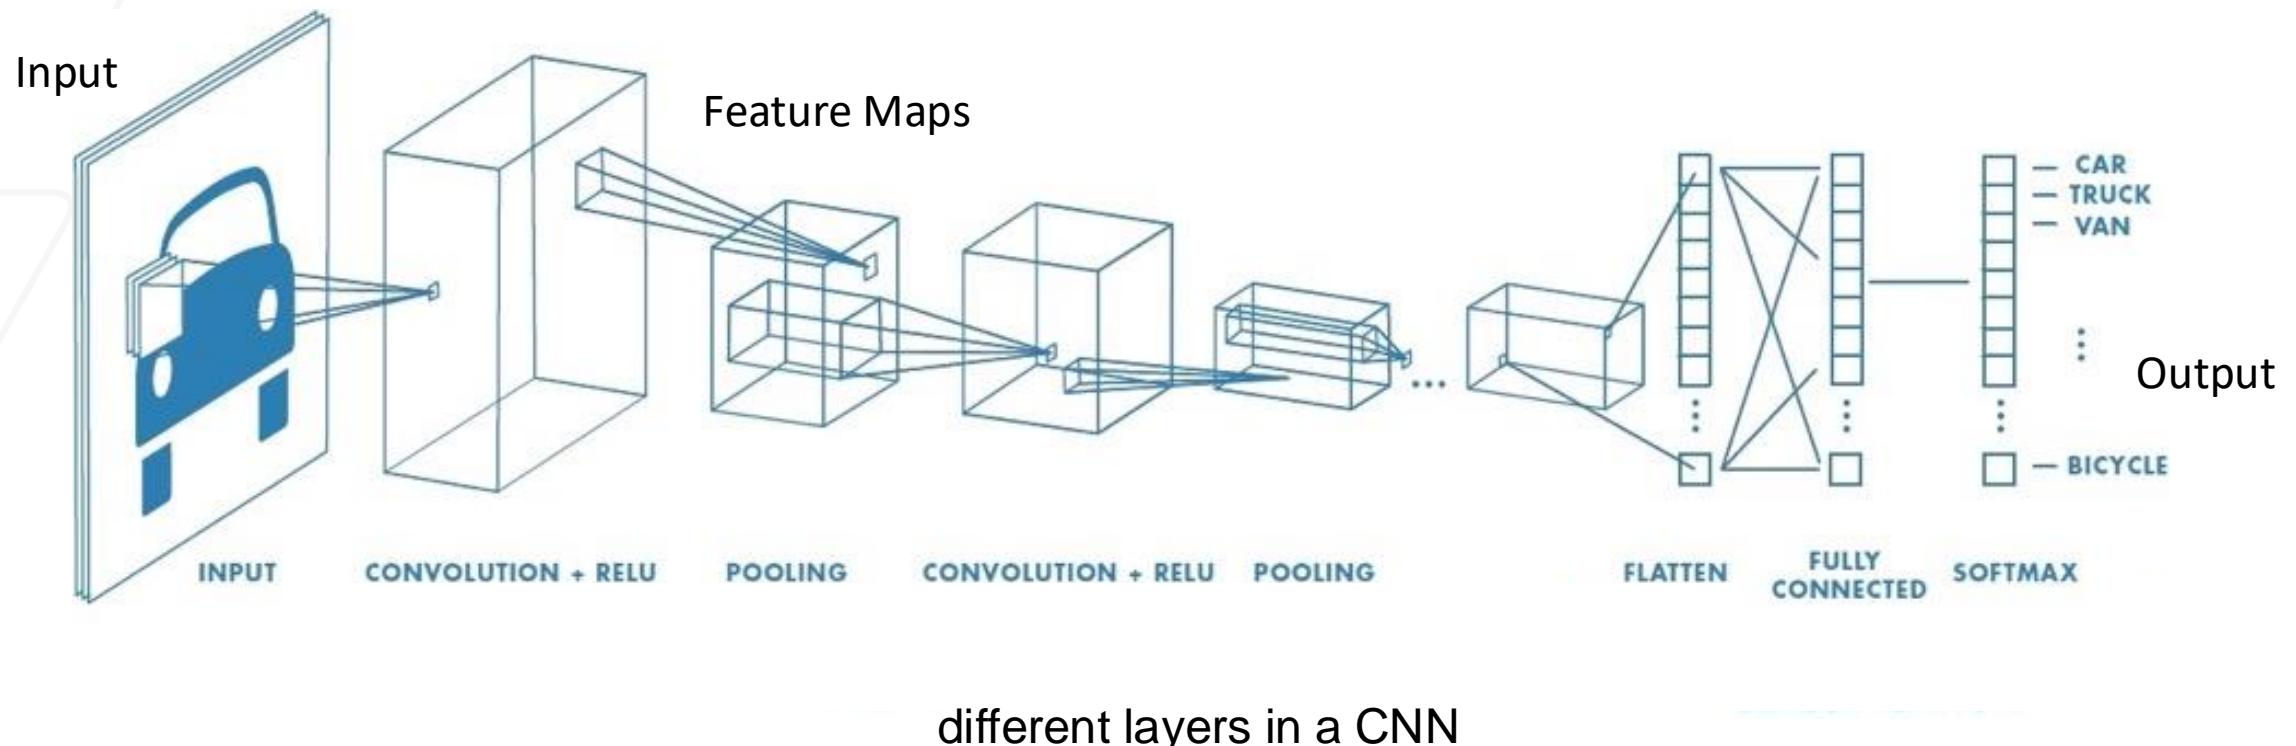
- The receptive field of $Z_{1,1}^{(2)}$ in the layer 1 has dimension of 3x3 (green region)
- The receptive field of $Z_{2,2}^{(3)}$ in the layer 1 has dimension of 5x5 (yellow region)



- Stacking more convolutional layers increases the receptive field size

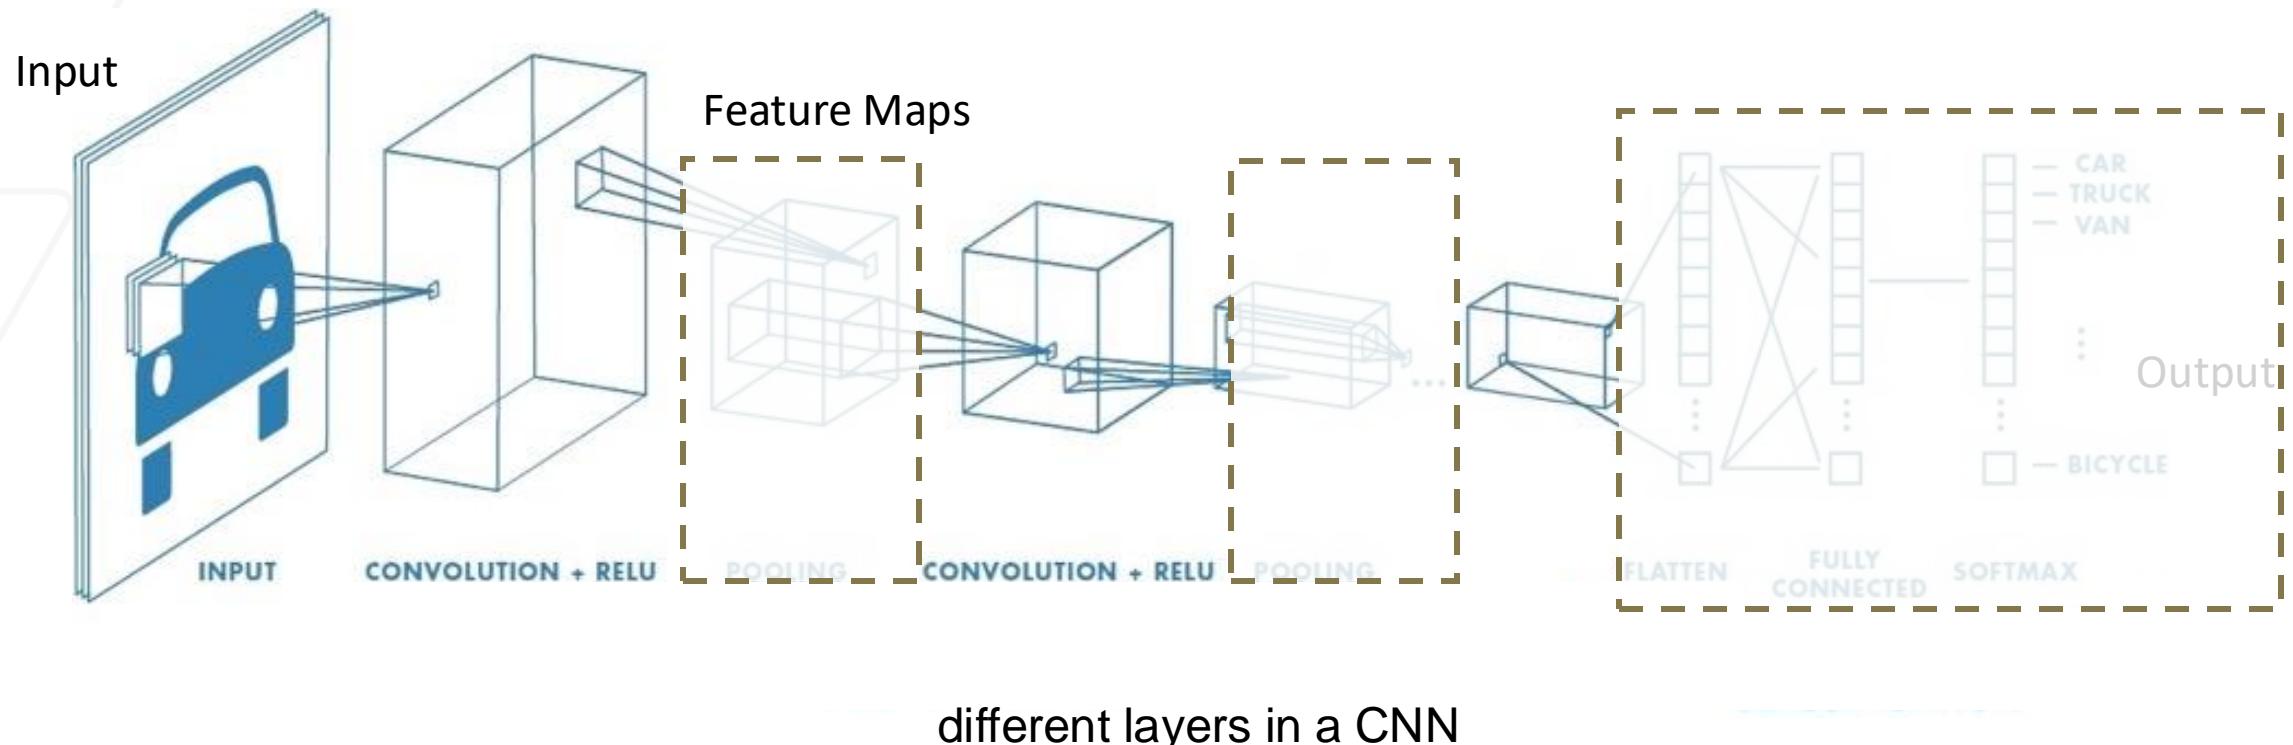
Convolutional Neural networks

Introduction



Convolutional Neural networks

Introduction



Overview

In this Lecture..

Locally Connected Layer

- Limitations of ANNNs
- Convolutional Layers

Introduction to Convolutional Neural Networks

Layers used to build Convolutional Neural Networks

- Convolutional layer
- Pooling Layer
- Fully-connected layers

Examples of Deep CNNs Architectures

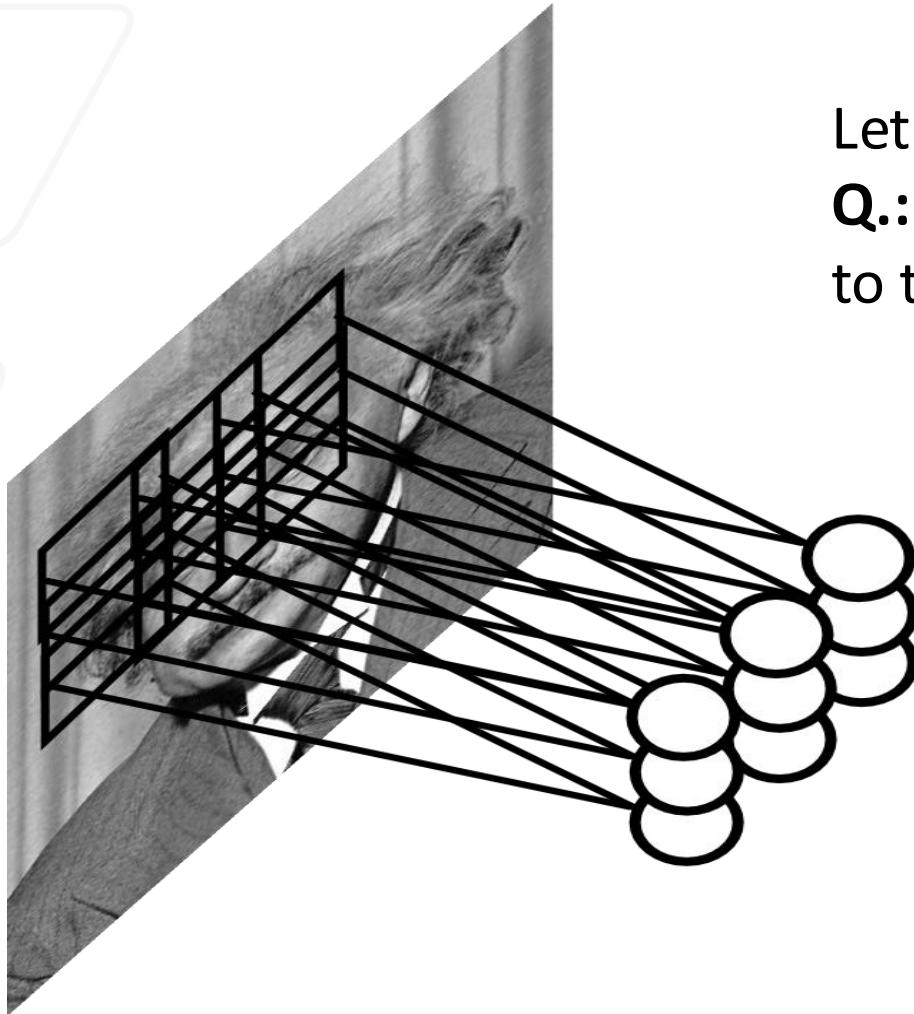
Training CNNs

- Optimization

Visualization of Convolutional Neural Networks

Convolutional Neural networks

Pooling Layer

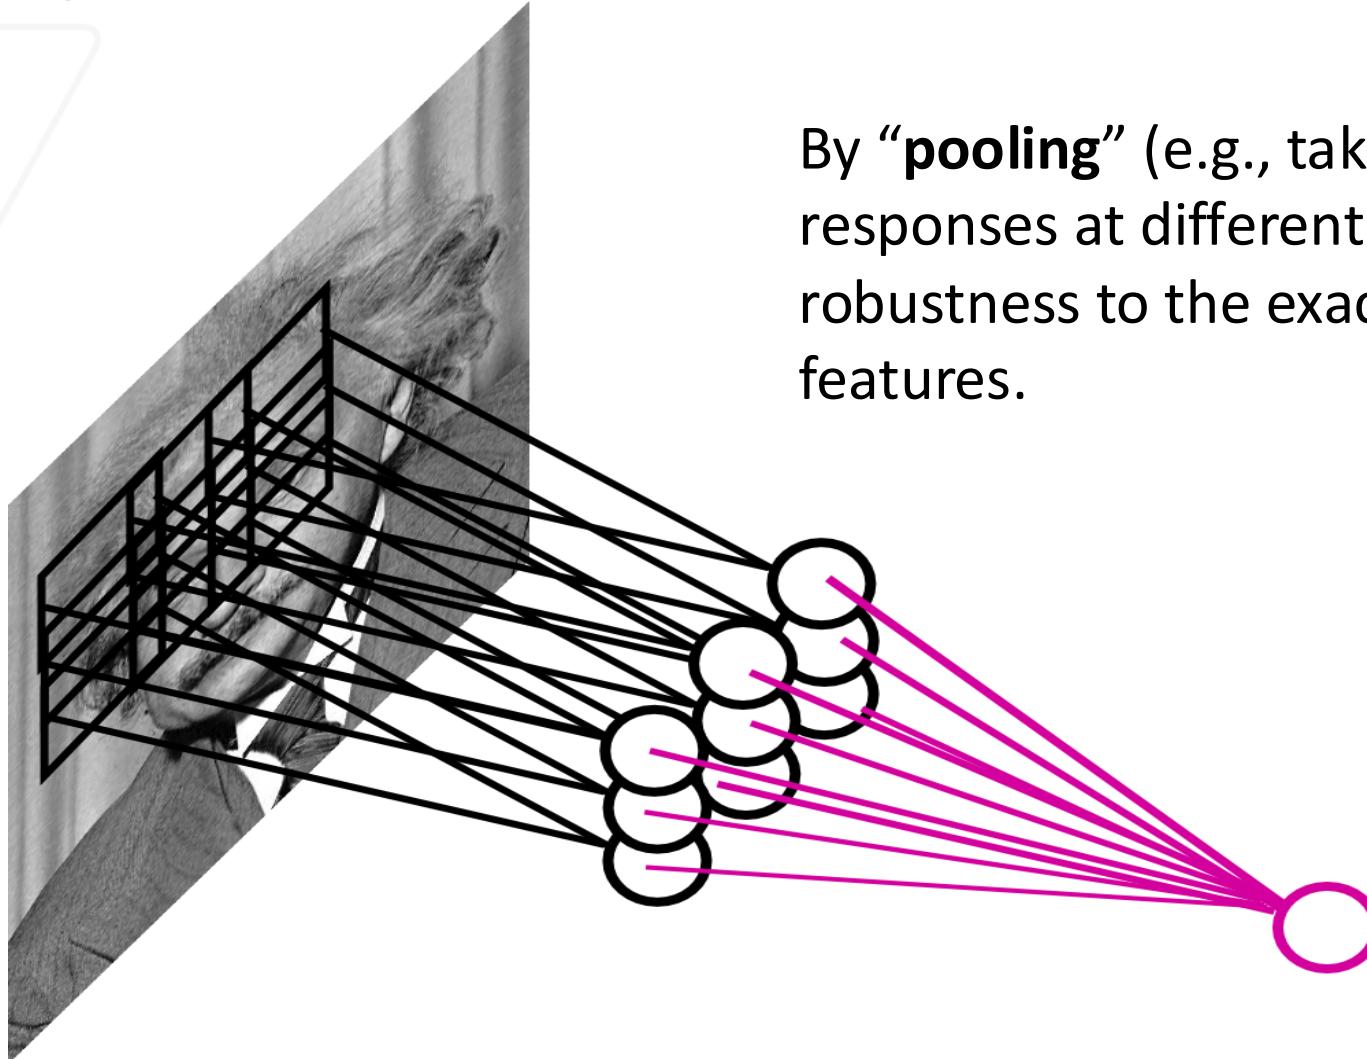


Let us assume the filter is an “eye” detector.

Q.: how can we make the detection robust
to the exact location of the eye?

Convolutional Neural networks

Pooling Layer

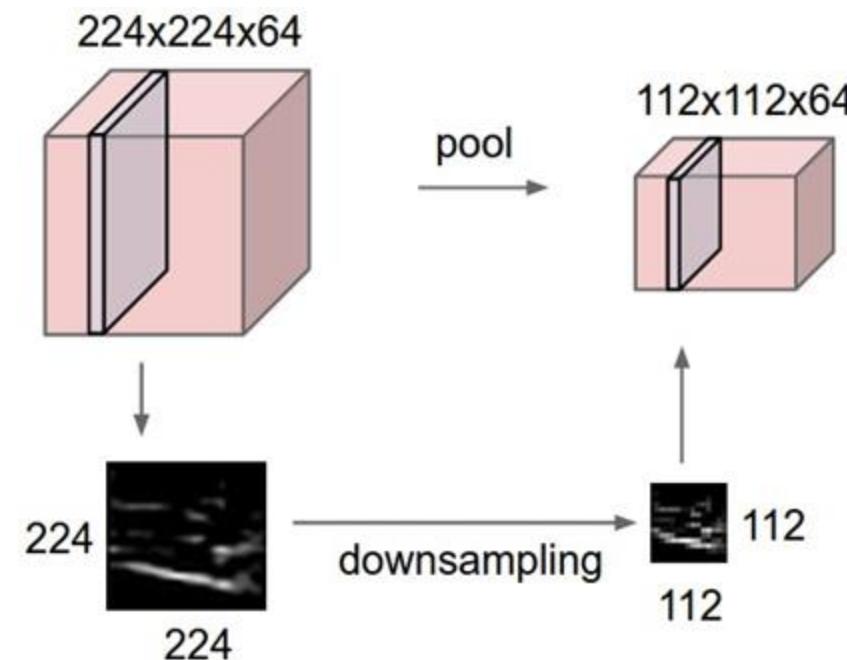


By “**pooling**” (e.g., taking max) filter responses at different locations we gain robustness to the exact spatial location of features.

Convolutional Neural networks

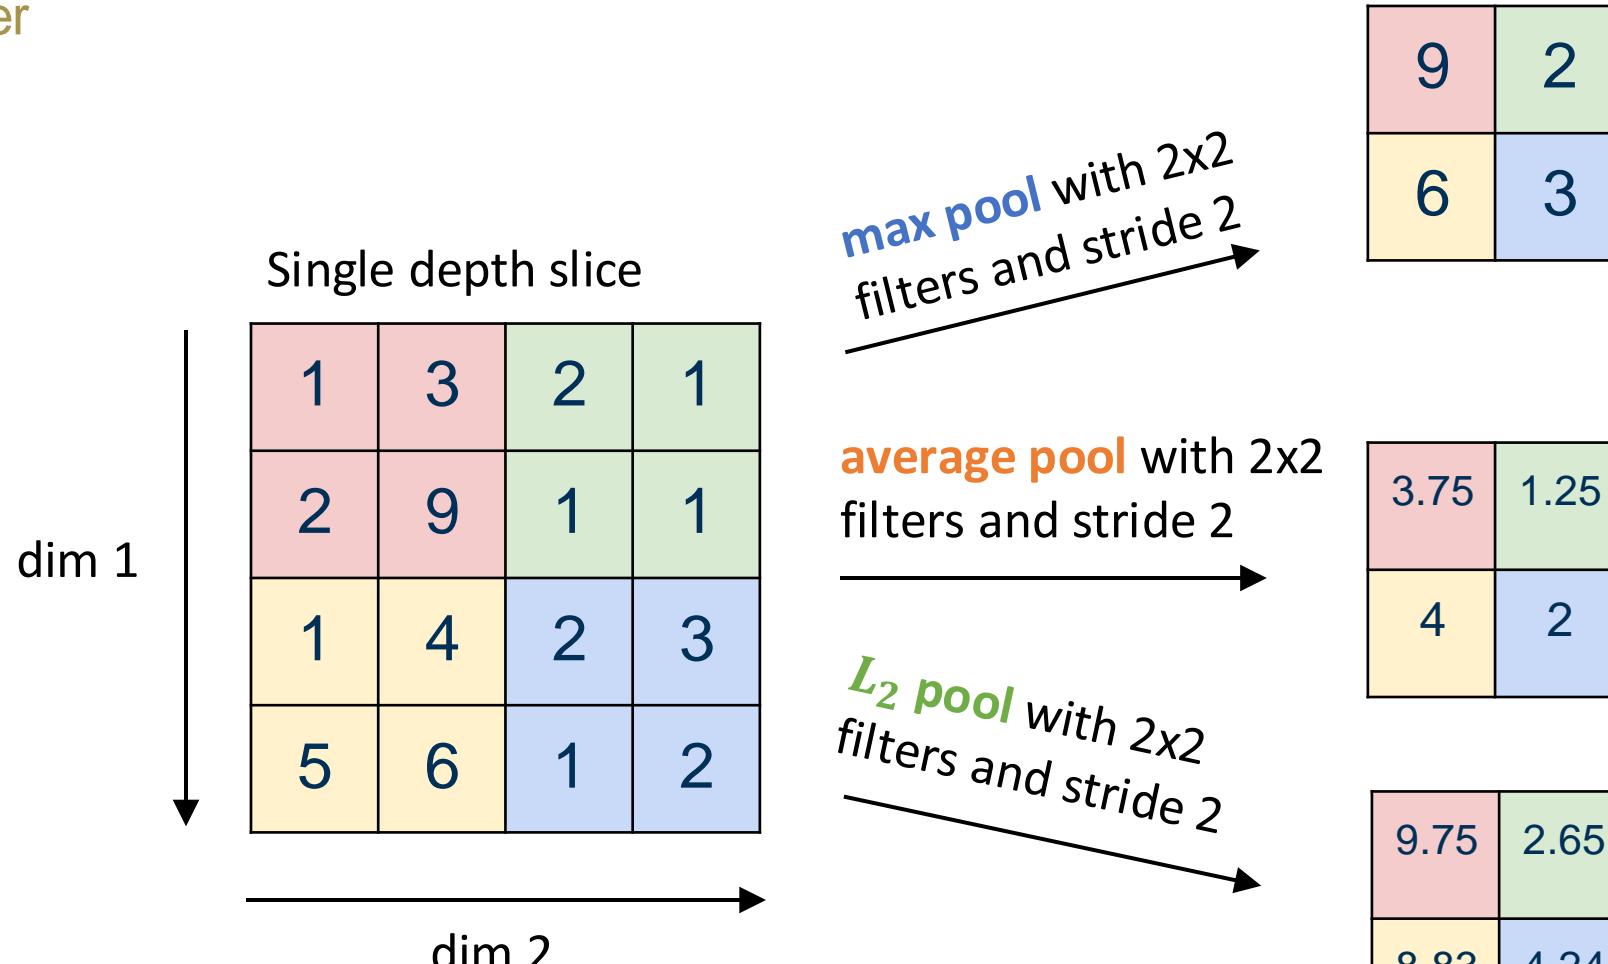
Pooling Layer

- makes the activation maps smaller and more manageable
- operates over each activation map independently



Convolutional Neural networks

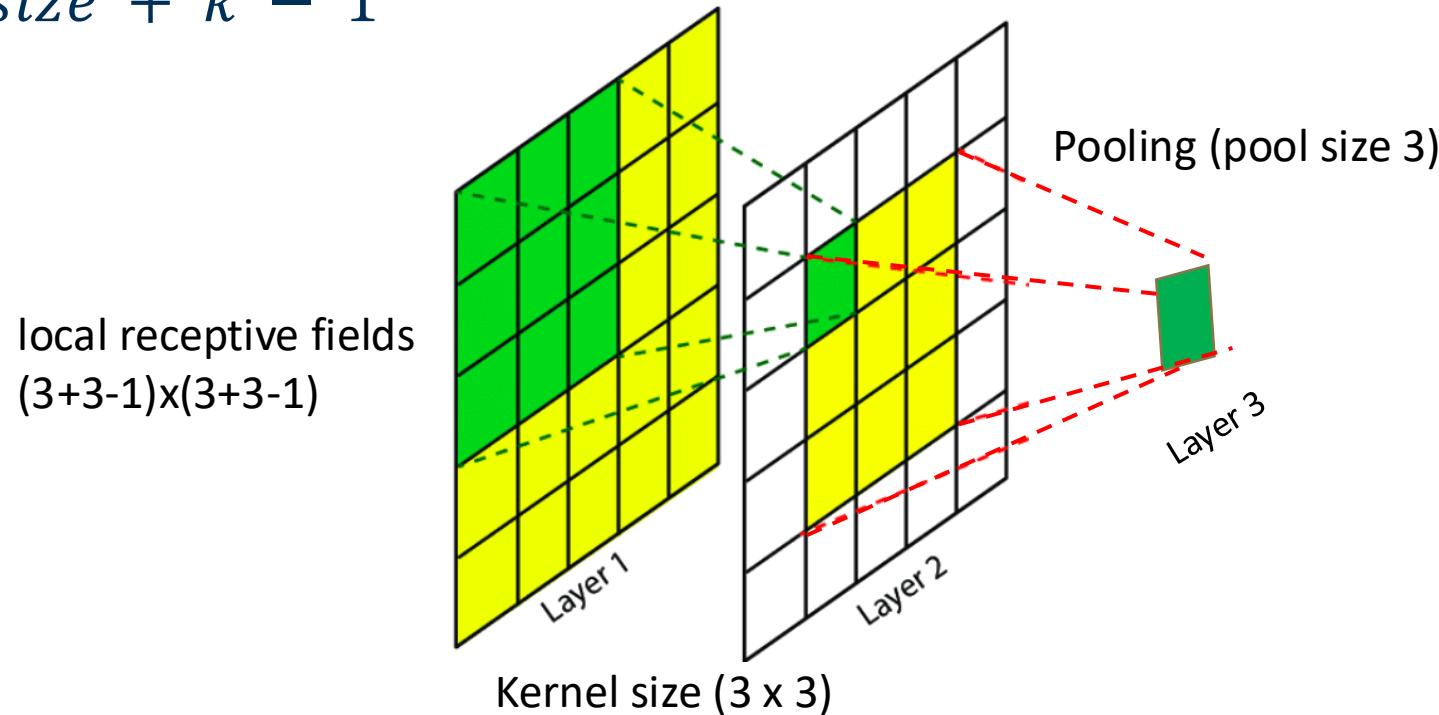
Pooling Layer



Convolutional Neural networks

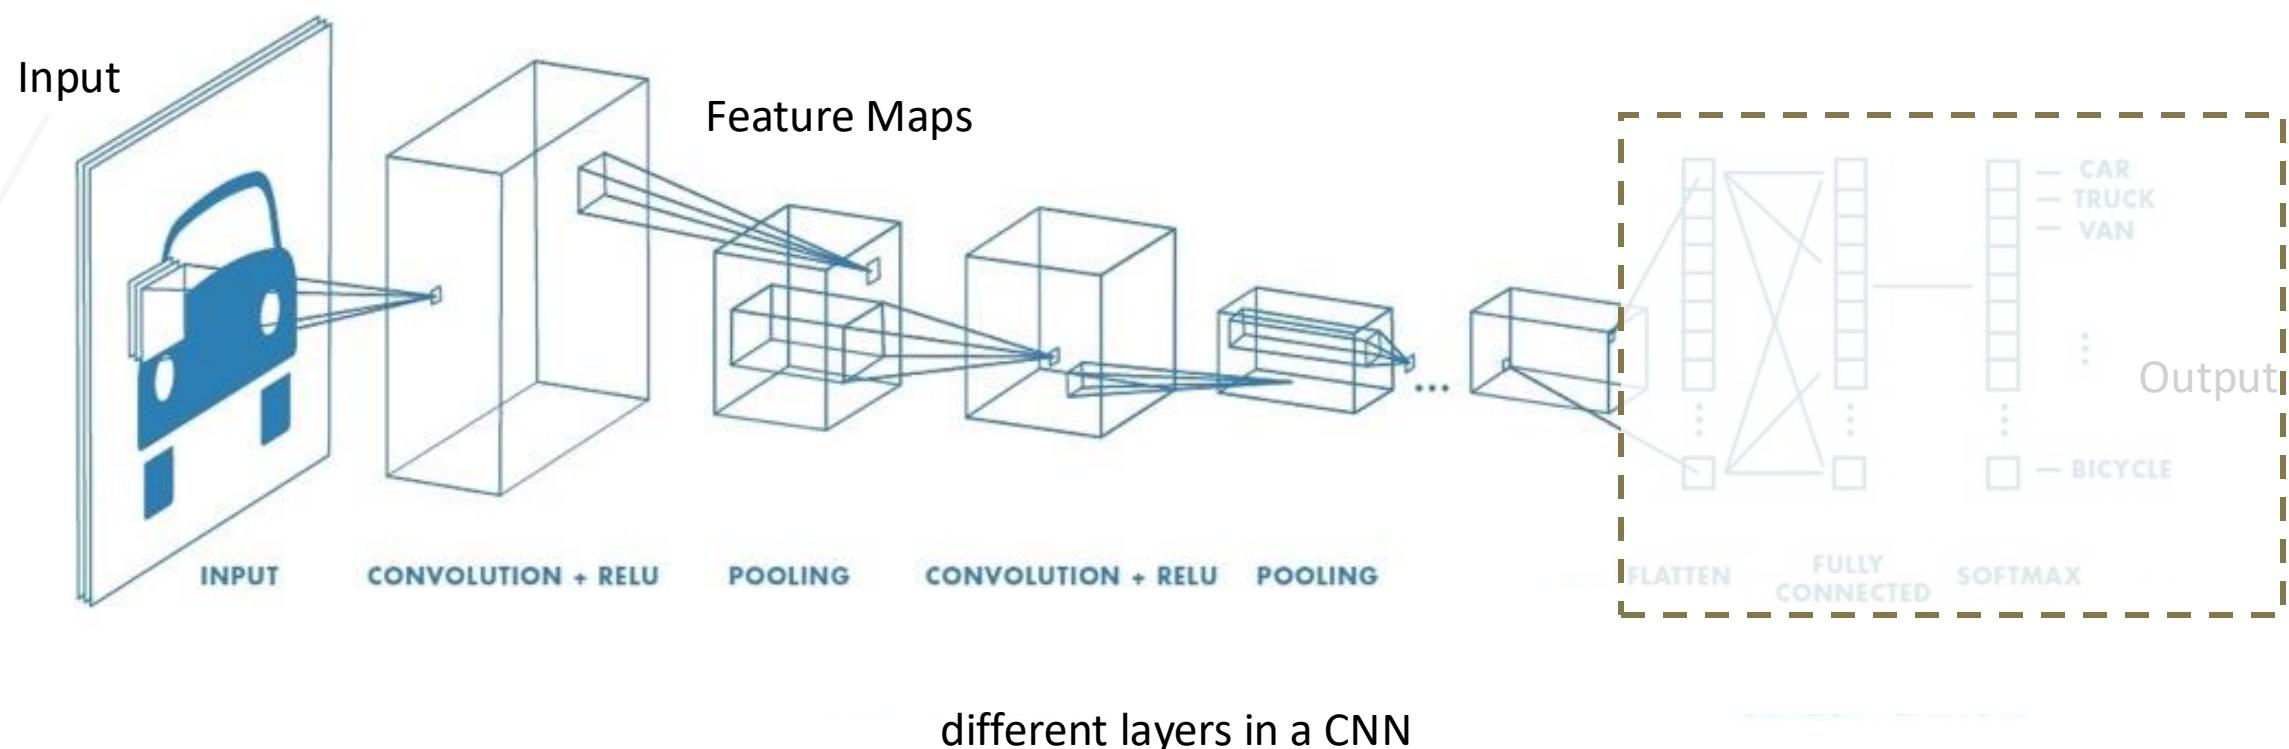
Pooling Layer: Receptive Field Size

- If filters are $k \times k$ and stride 1
- then each neuron in pooling layer 3 depends upon a patch in layer 1 of size: $pool_size + k - 1$



Convolutional Neural networks

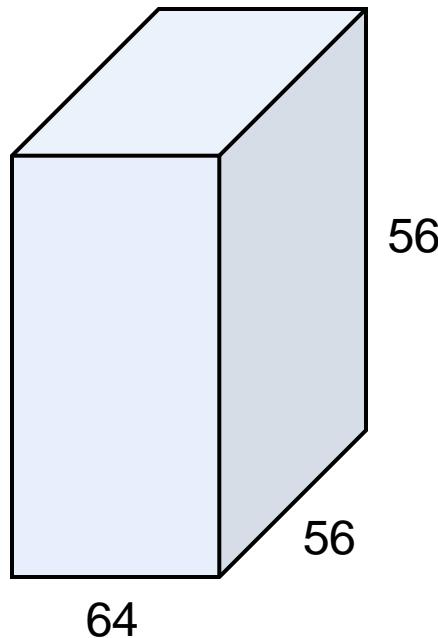
Introduction



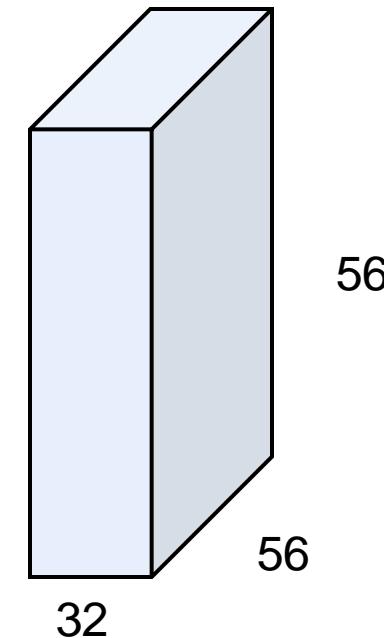
Convolutional Neural networks

Fully-connected Layer: 1x1 Convolution

1x1 convolution layers make perfect sense



1x1 CONV
with 32 filters
→
(each filter has size
1x1x64, and performs a
64-dimensional dot
product)

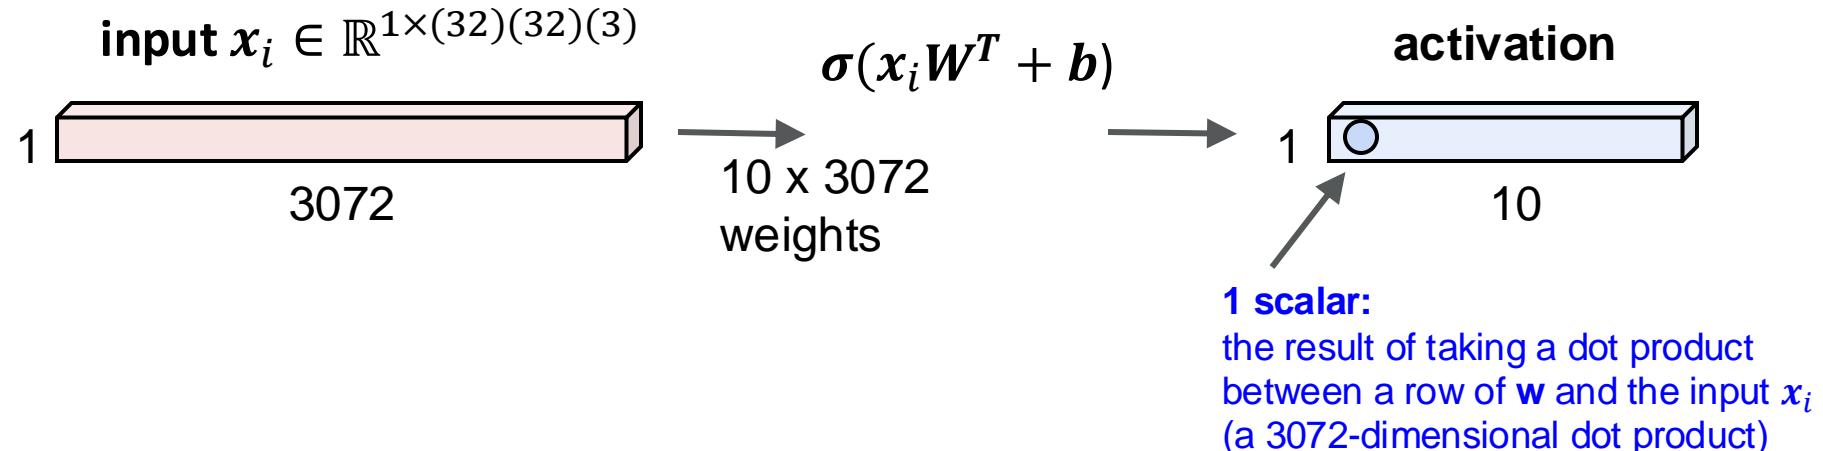


Convolutional Neural networks

Fully-connected Layer: 1x1 Convolution

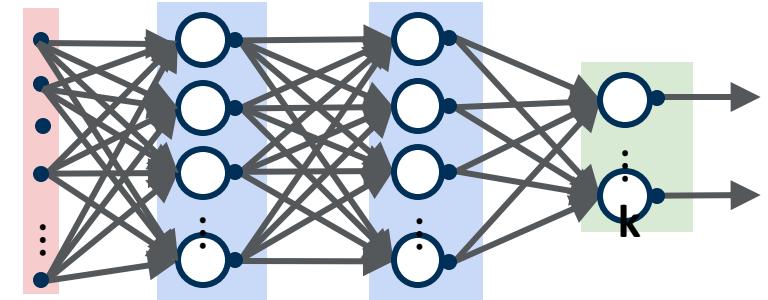
Fully-Connected Layer as 1x1 Convolution

32x32x3 image -> stretch to 3072 x 1

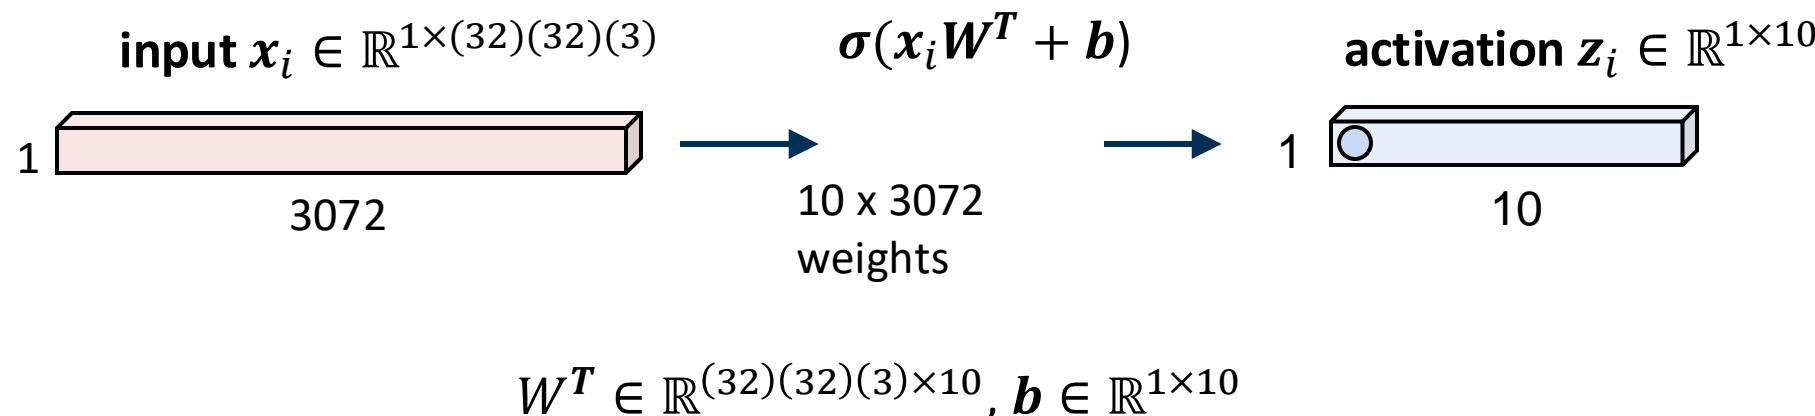


Convolutional Neural networks

Fully-connected Layer: 1x1 Convolution



Neurons in a fully connected layer have full connections to all activations in the previous layer, as seen in Artificial Neural Networks. Their activations can hence be computed with a matrix multiplication followed by a bias offset.

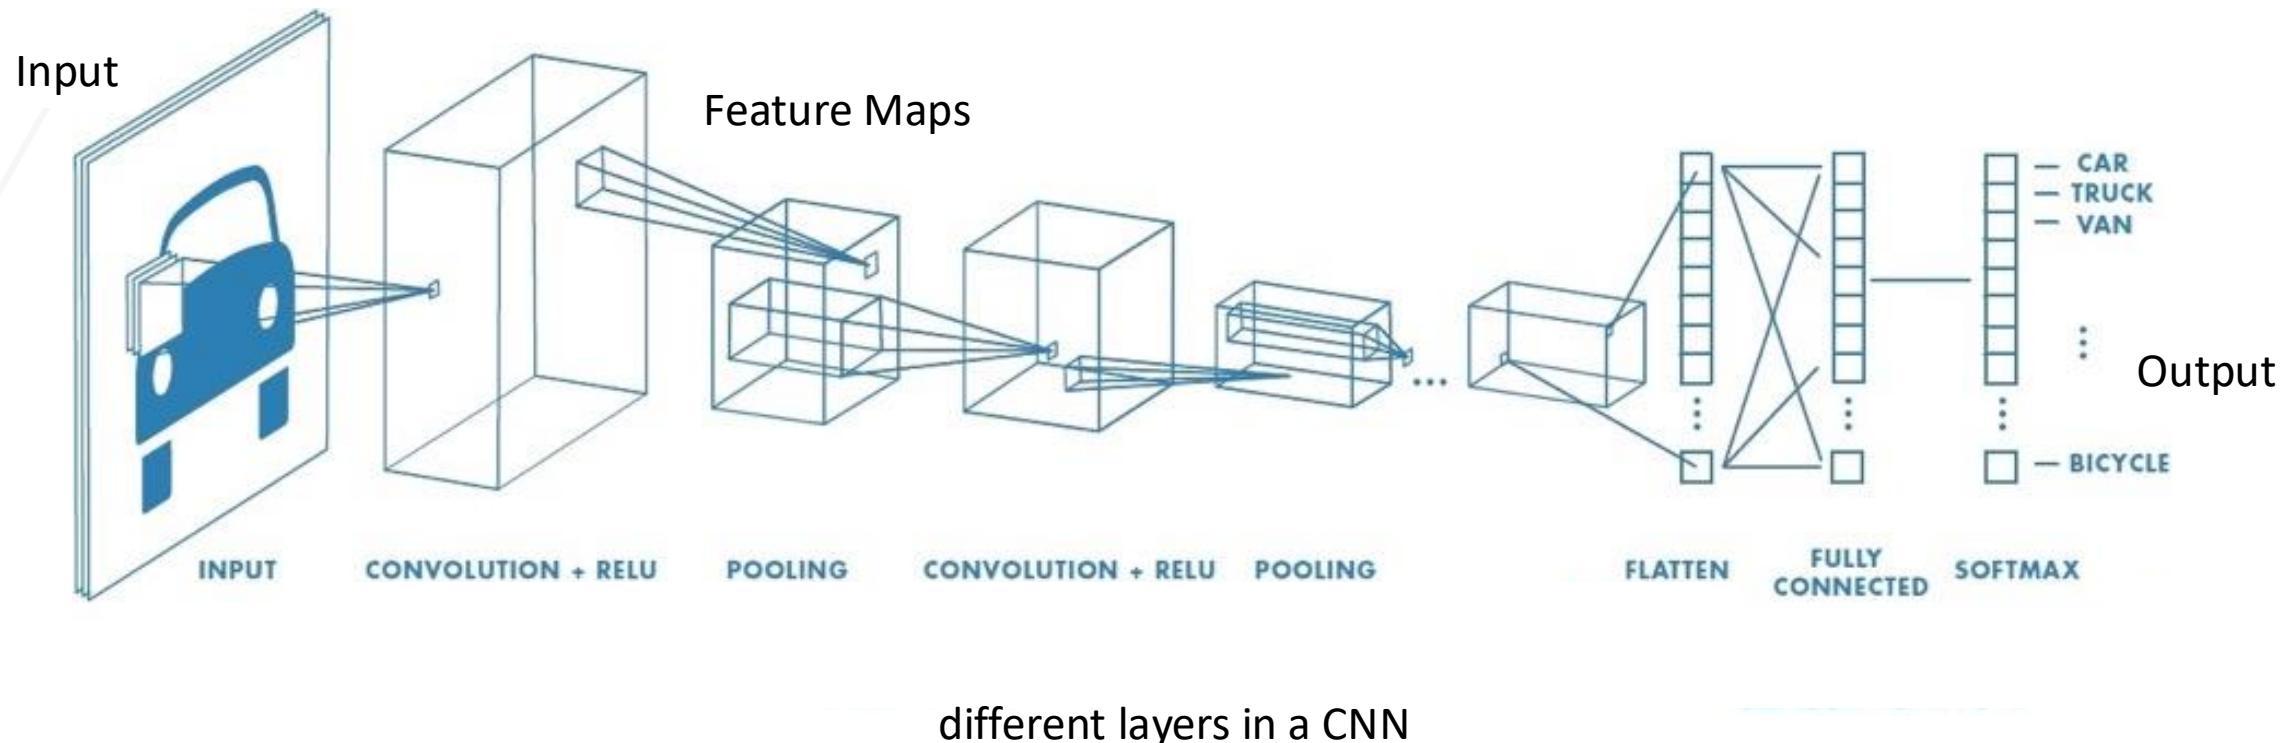


Convolutional Neural networks

Fully-connected Layer: 1x1 Convolution

Three main types of layers to build ConvNet architectures:

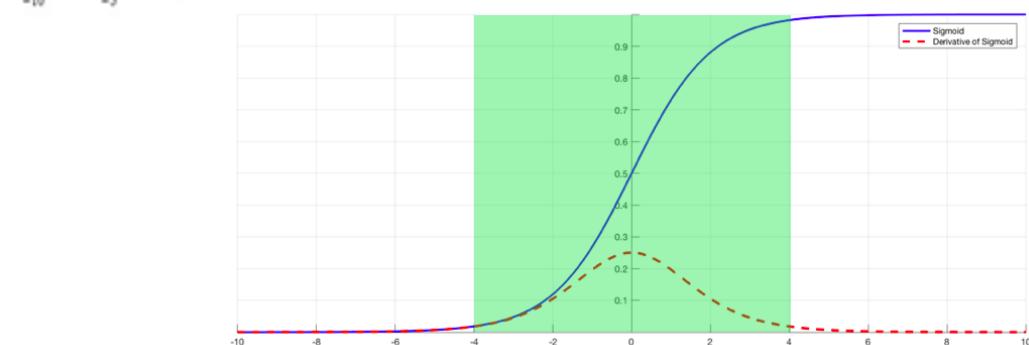
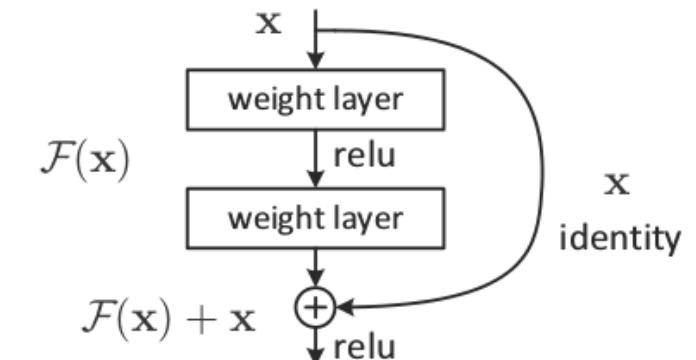
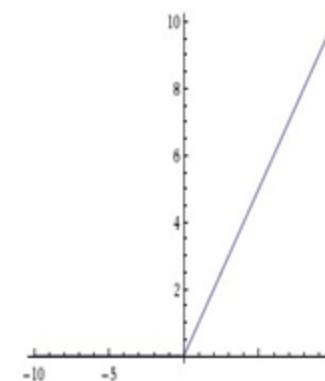
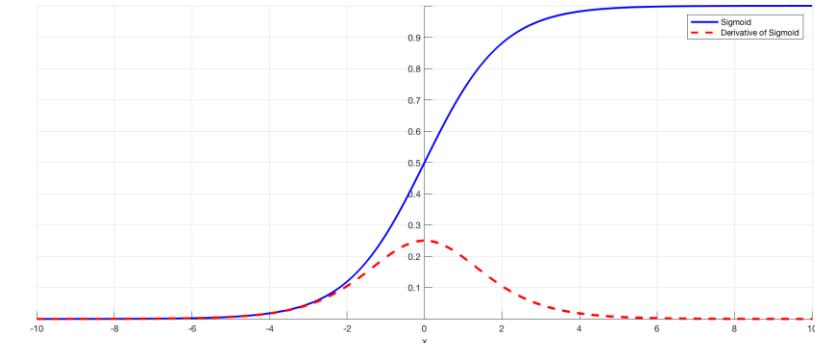
- Convolutional Layer, Pooling Layer, and Fully-Connected Layer (ANN)



Convolutional Neural networks

Vanishing Gradient Problem

- Using the sigmoid, the gradient could become smaller and smaller and in deep networks, that may cause the training to be ineffective as no change could take place [check the derivative of the sigmoid in the plot]
- Possible solutions:
 - Use ReLU
 - Use Residual networks
 - Normalization



Terminology

- *Distribution*: (sample space) the set of all possible samples
- *Dataset*: a set of samples drawn from a distribution
- *Batch*: a subset of samples drawn from the dataset
- *Sample*: a single data object represented as a set of features
- *Feature*: value of a single attribute, property, in a sample. Could be numeric or categorical.

Appendix A: Notations

- x_i : a single feature
- \boldsymbol{x}_i : feature vector (a data sample)
- $\boldsymbol{x}_{:,i}$: feature vector of all data samples
- \boldsymbol{X} : matrix of feature vectors (dataset)
- N : number of data samples
- P : number of features in a feature vector
- α : learning rate
- Bold letter/symbol: vector
- Bold capital letters/symbol: matrix